



in **f**ra

SVE: Distributed Video Processing at Facebook Scale

Qi Huang

Petchean Ang, Peter Knowles, Tomasz Nykiel, Iaroslav Tverdokhlib,
Amit Yajurvedi, Paul Dapolito IV, Xifan Yan, Maxim Bykov, Chuen Liang, Mohit
Talwar, Abhishek Mathur, Sachin Kulkarni, Matthew Burke, Wyatt Lloyd

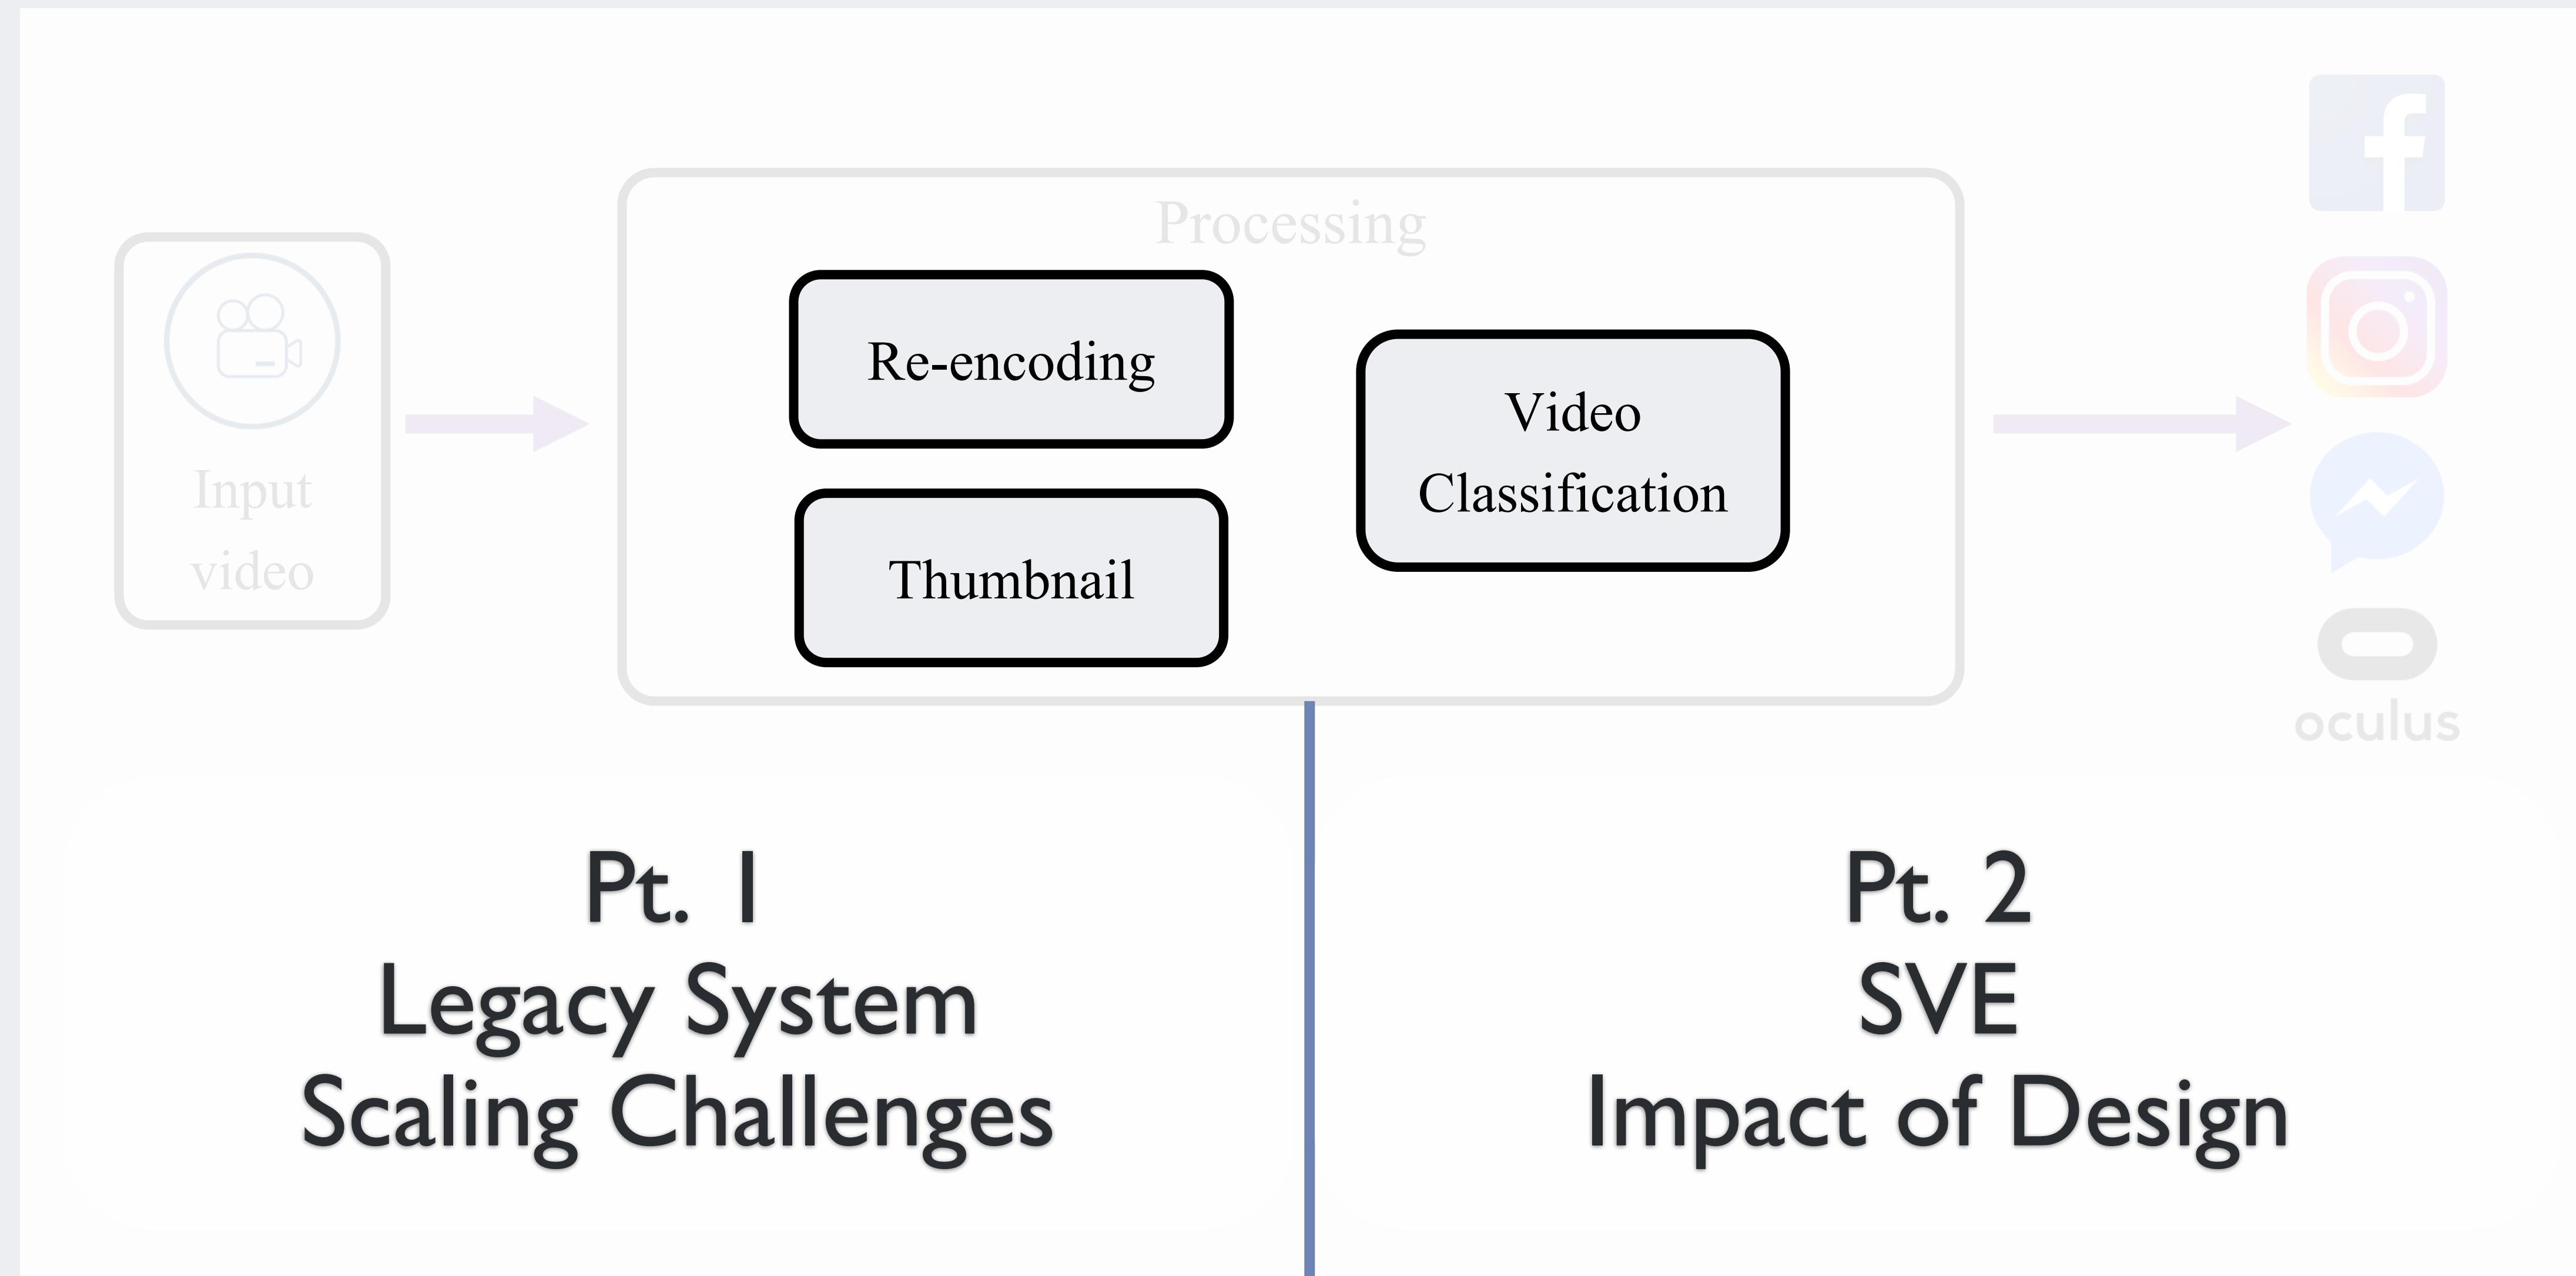
Facebook, University of Southern California, Cornell, Princeton

Video is growing across Facebook

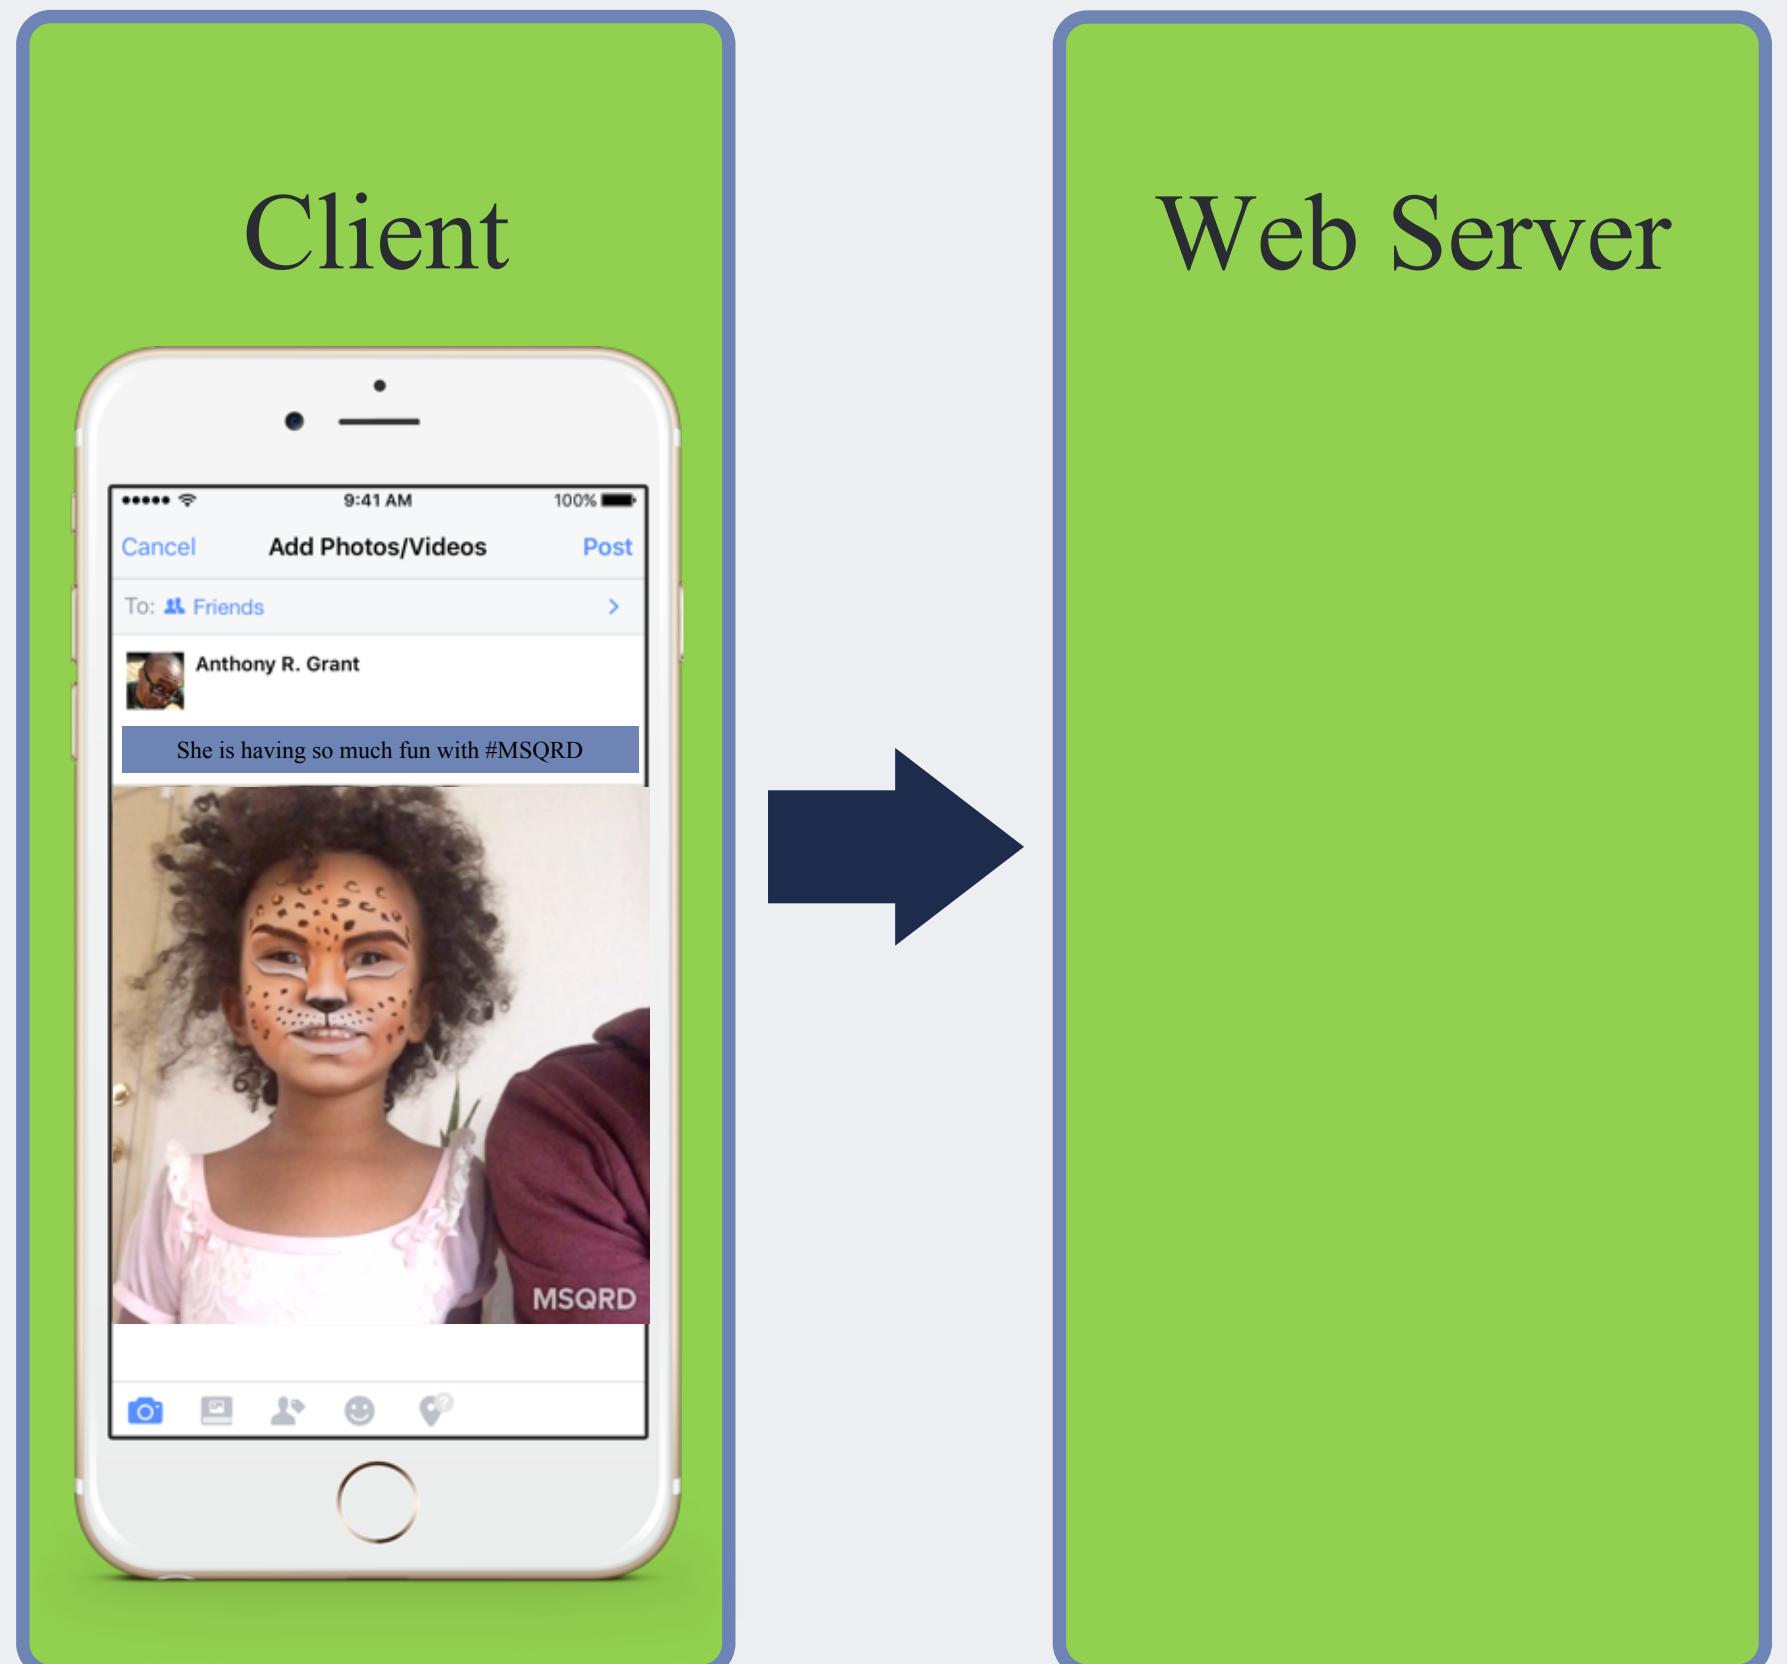


- FB: **500M** users watch **100M hours** video daily (Mar. 16)
- Instagram: **250M** daily active users for stories (Jun. 17)
- All: **many tens of millions** of daily uploads, **3X** NYE spike

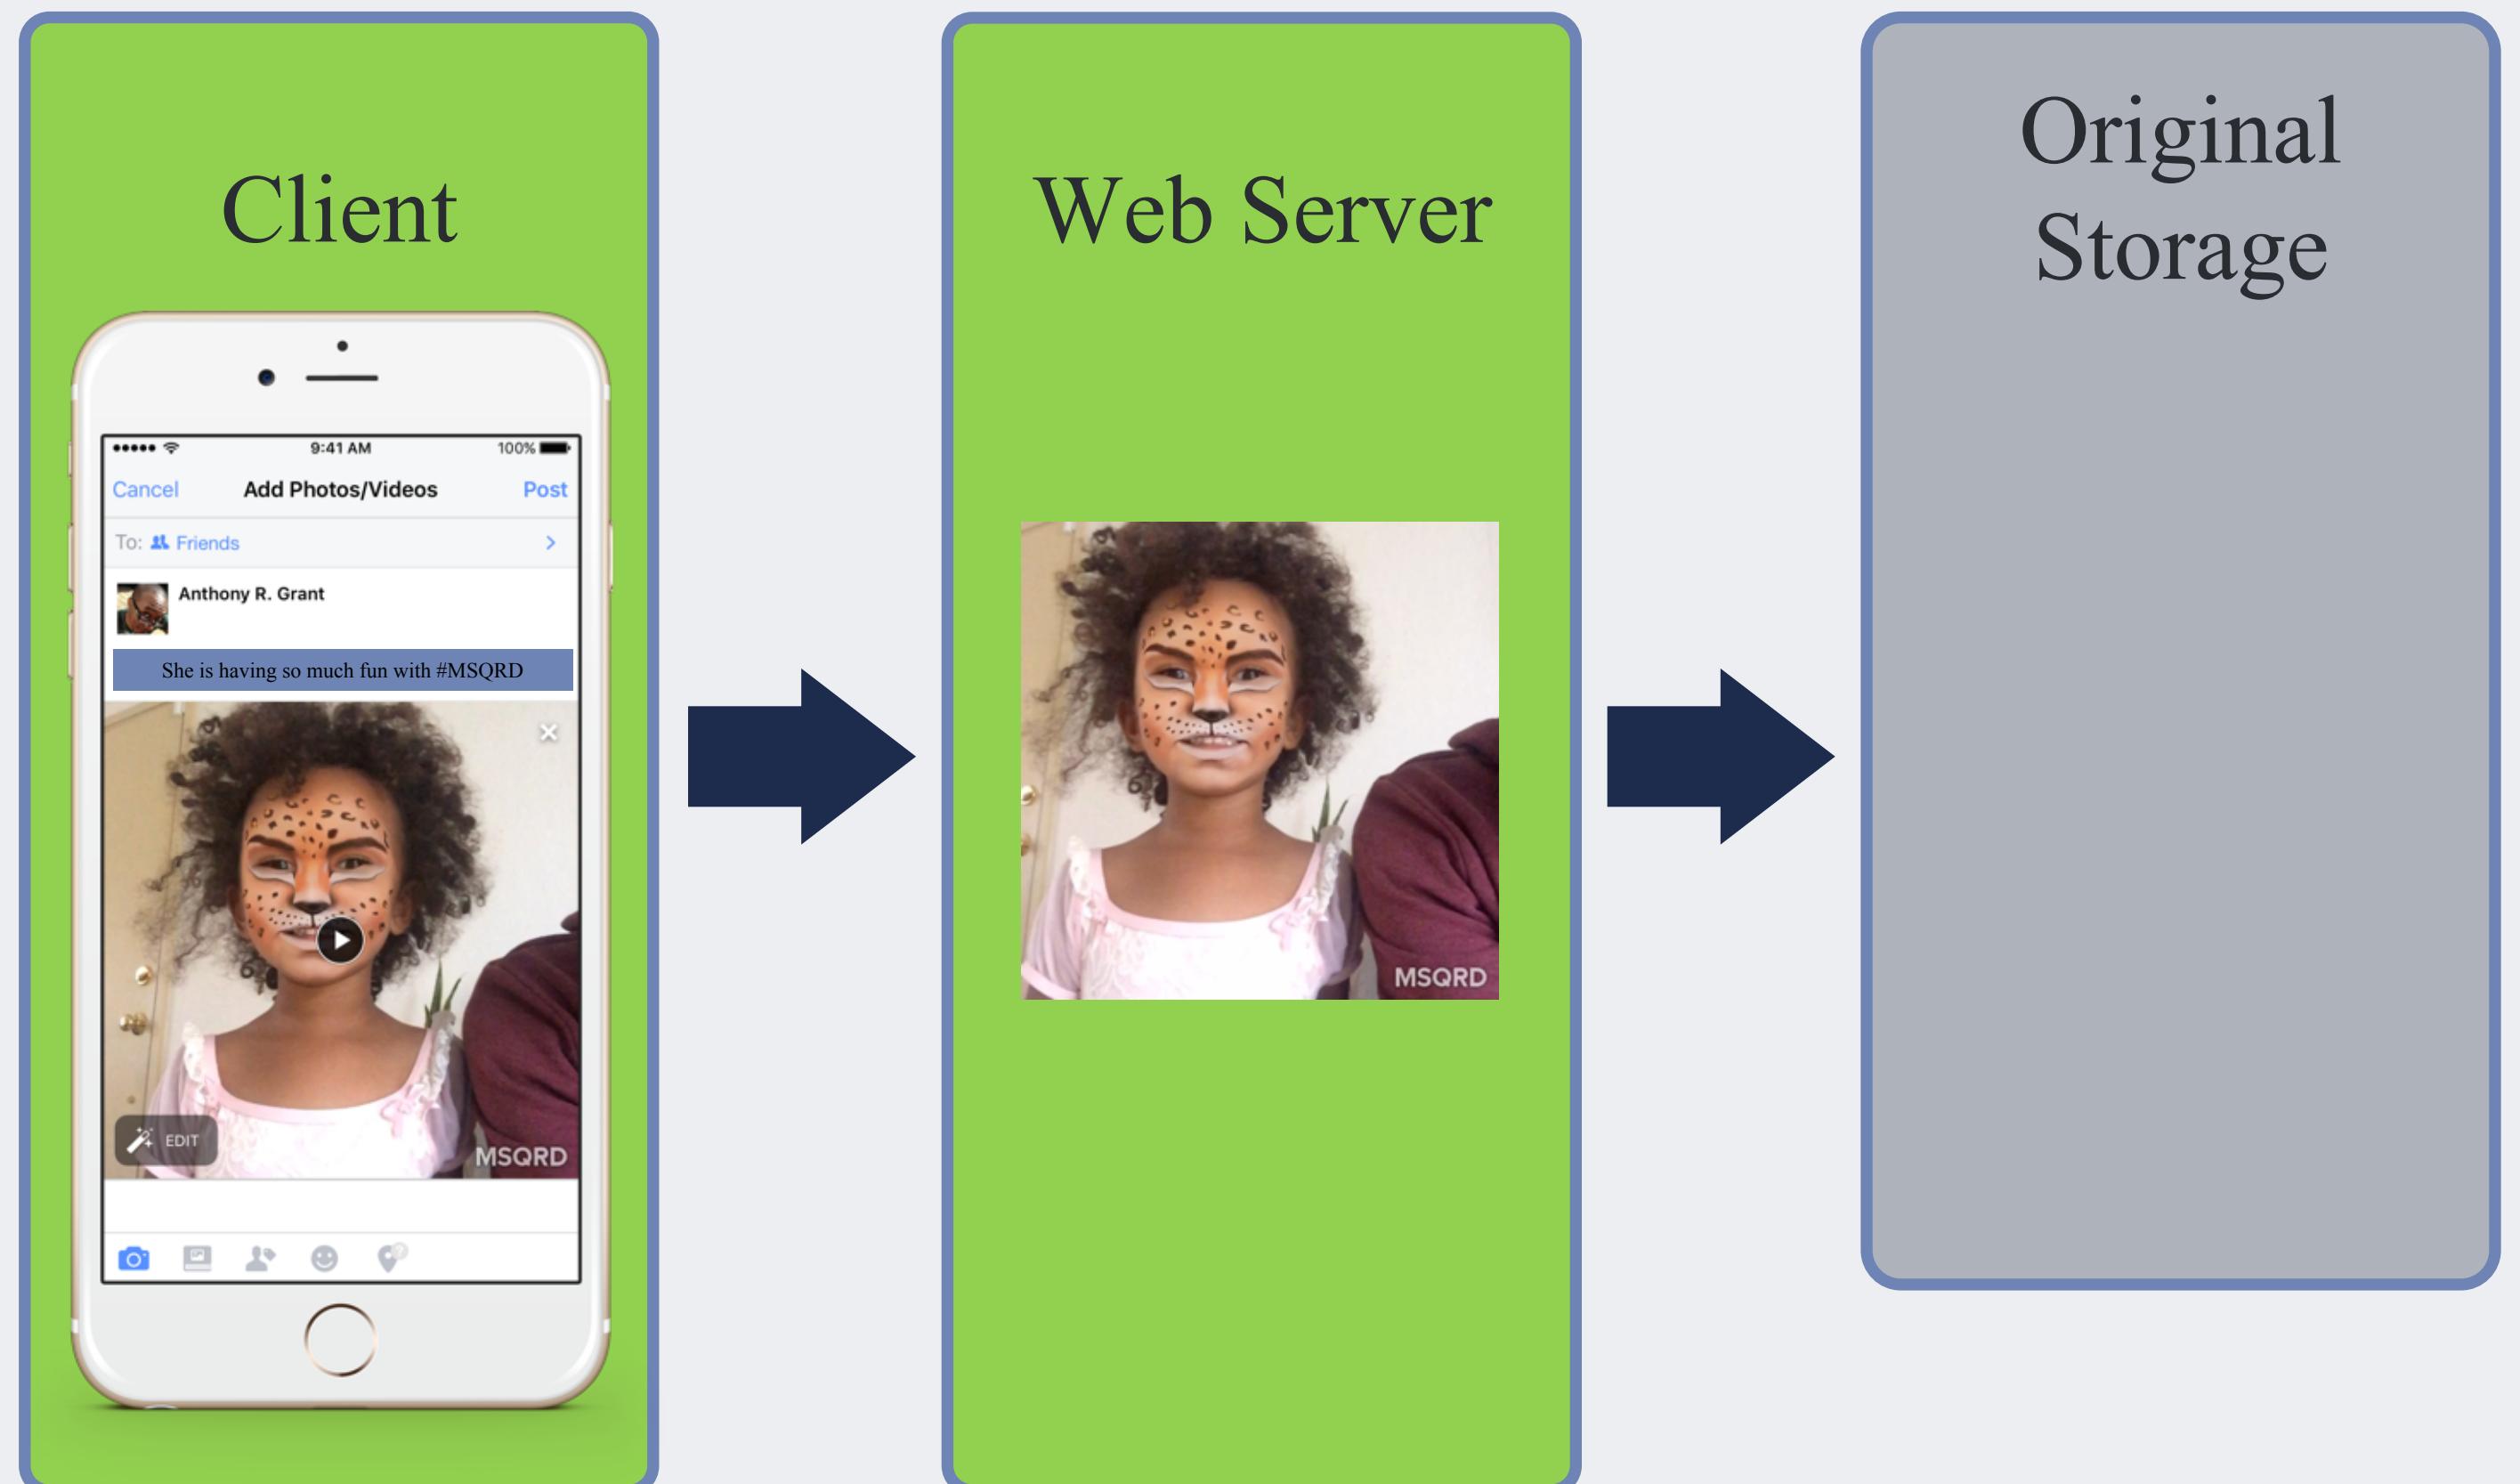
Processing is diverse and demanding



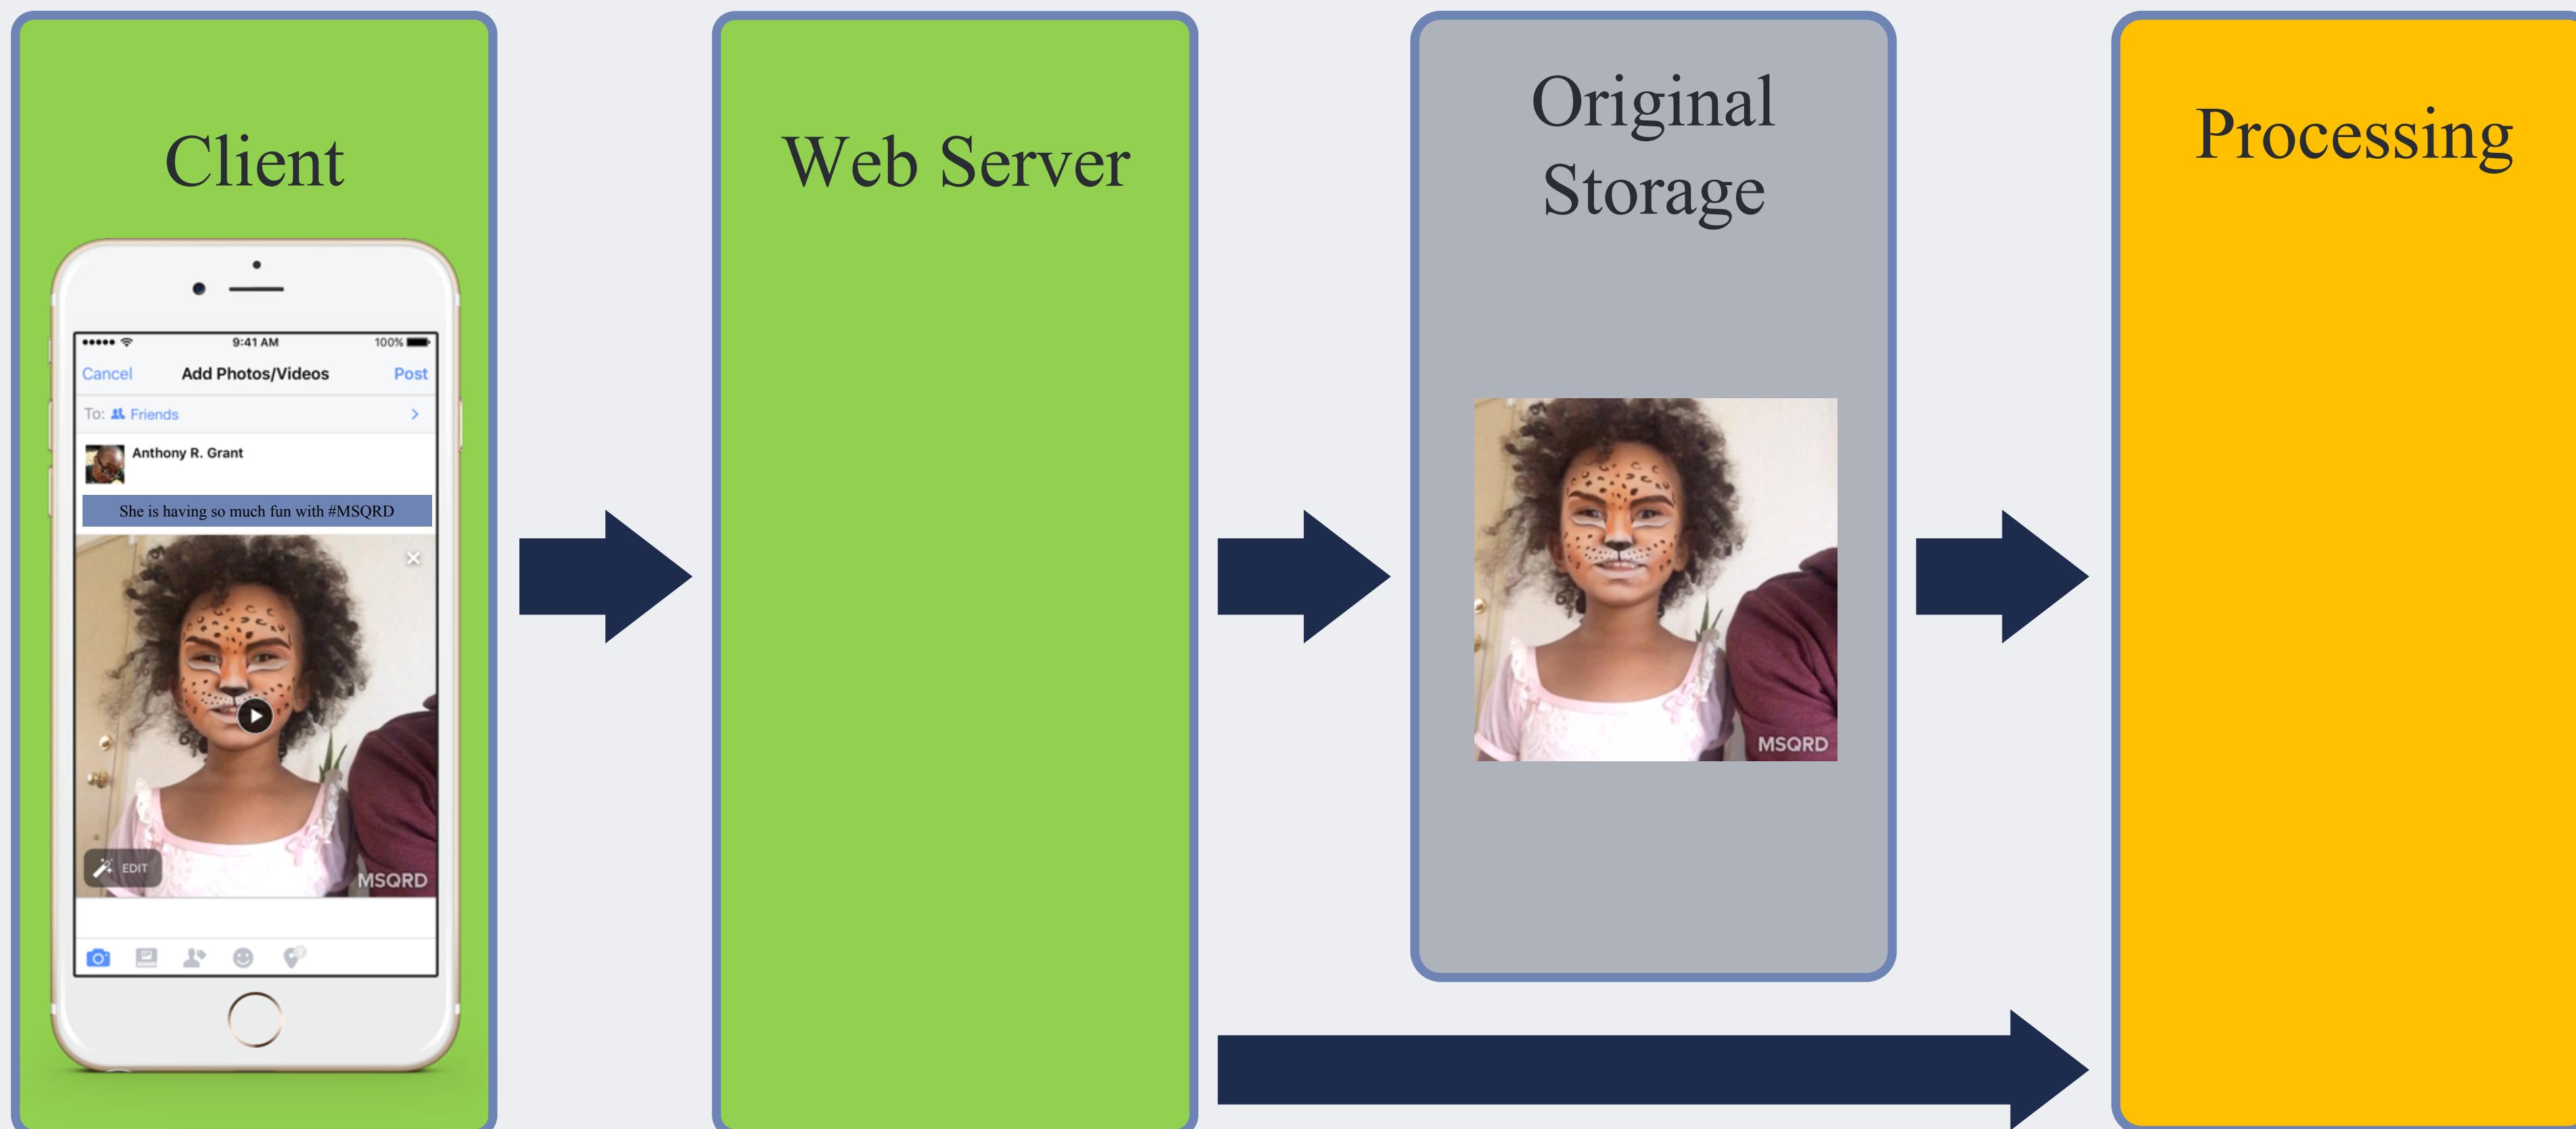
Legacy: upload video file to web server



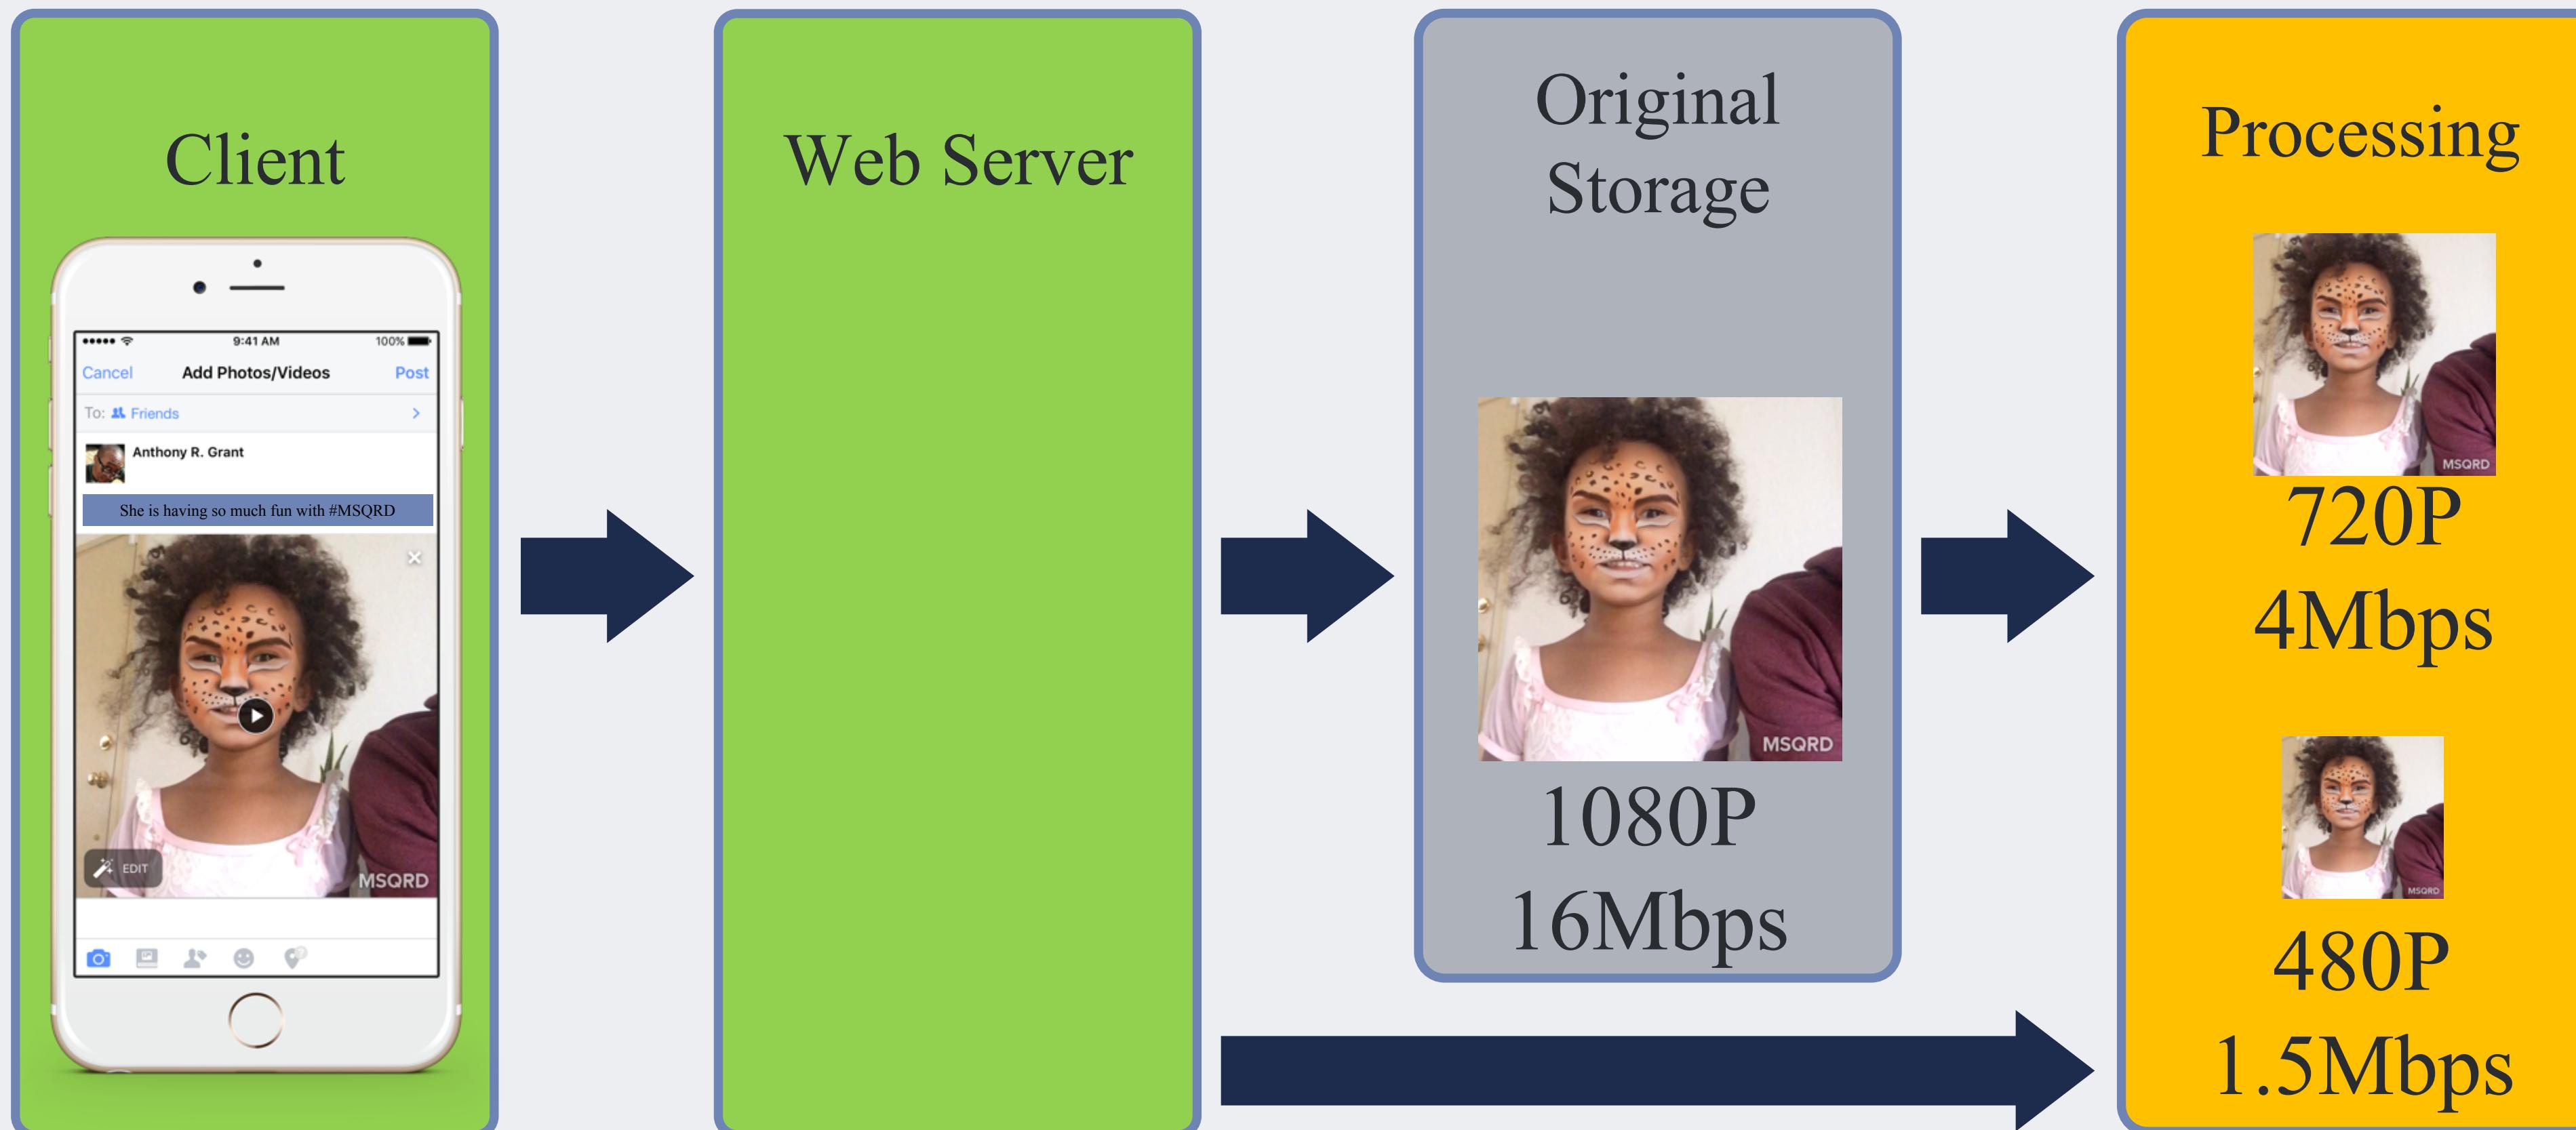
Legacy: preserve original for reliability



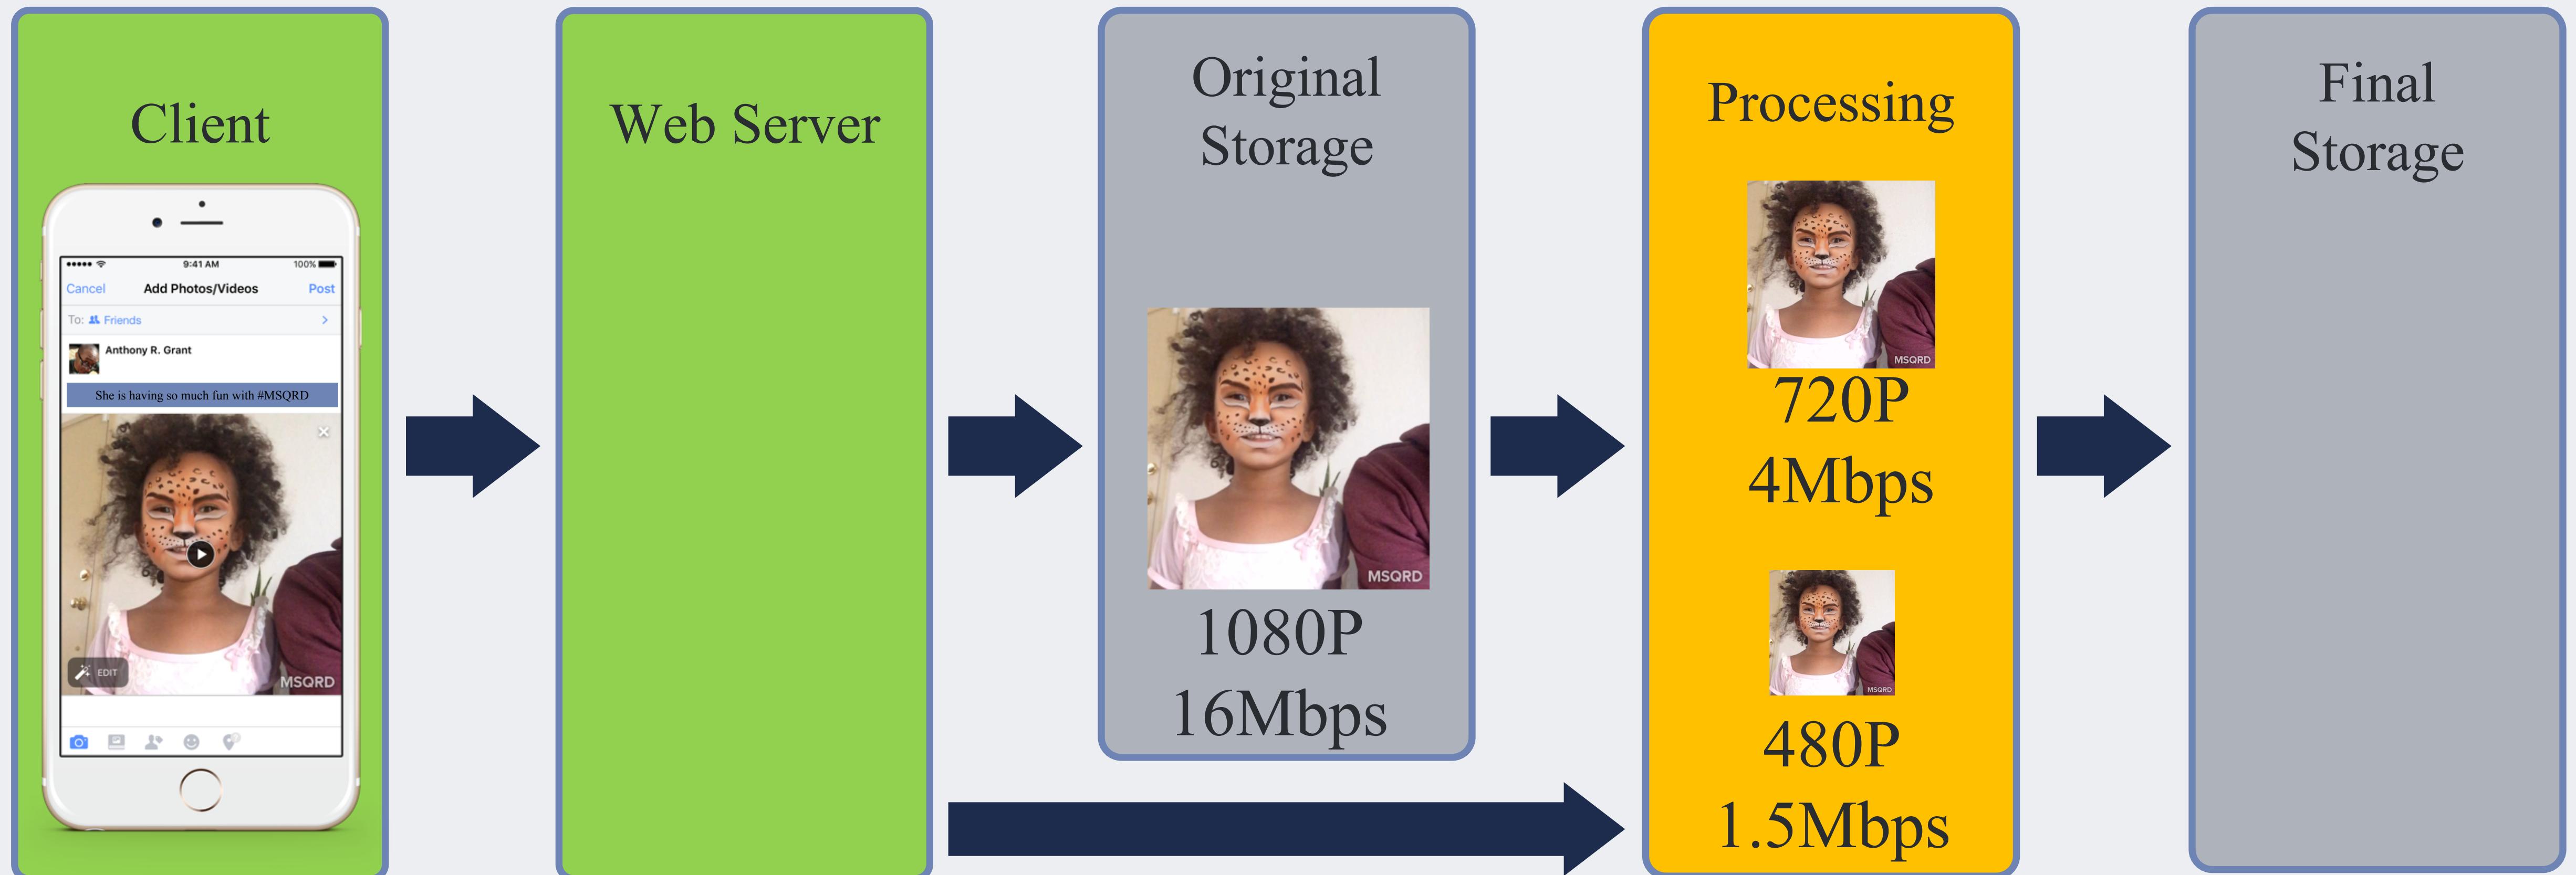
Legacy: process after upload completes



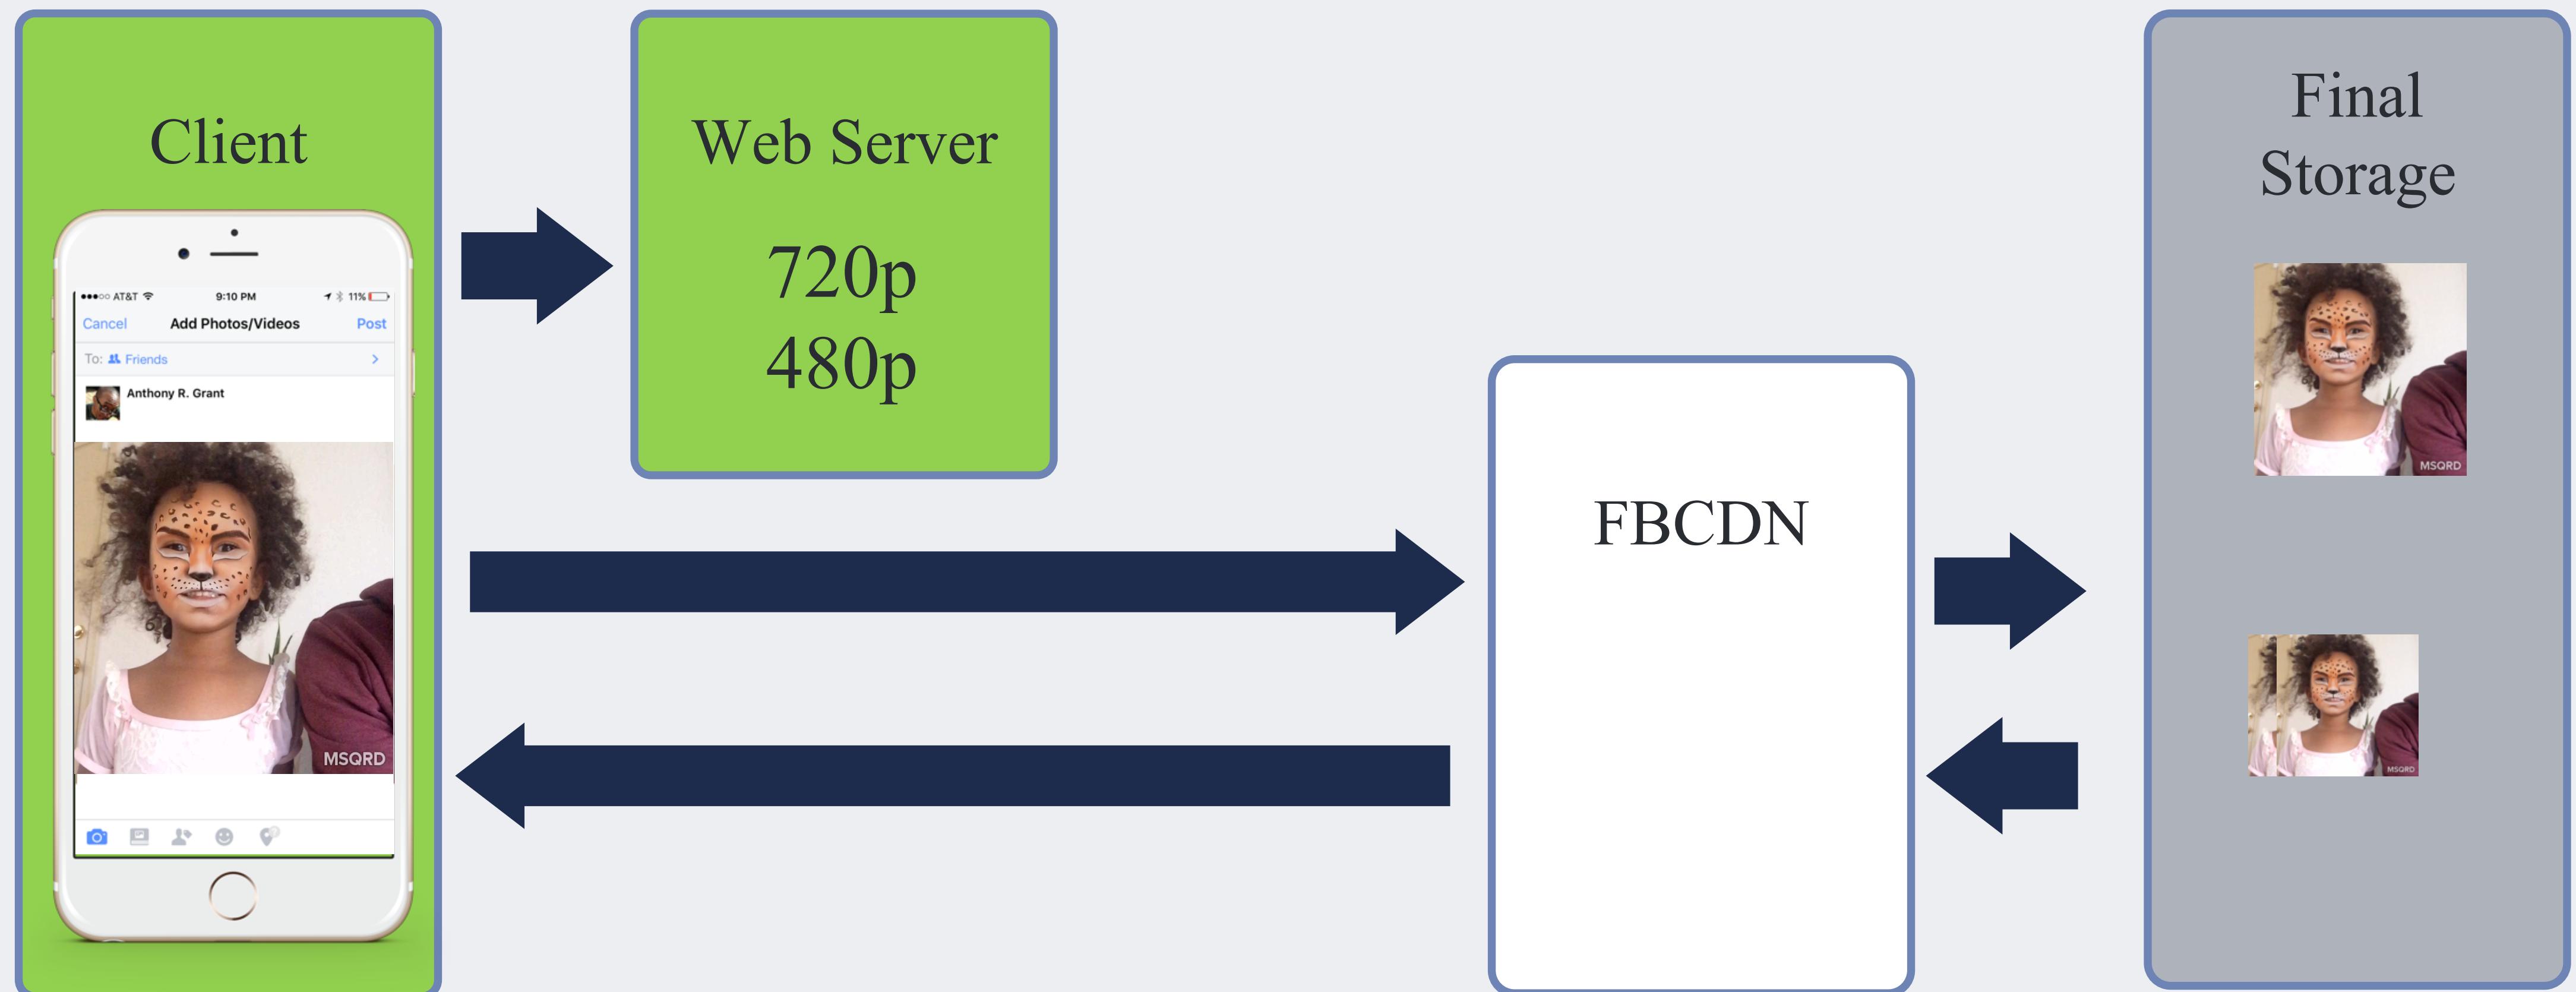
Legacy: encode w/ varying bitrates



Legacy: store encodings before sharing



Sharing with adaptive streaming

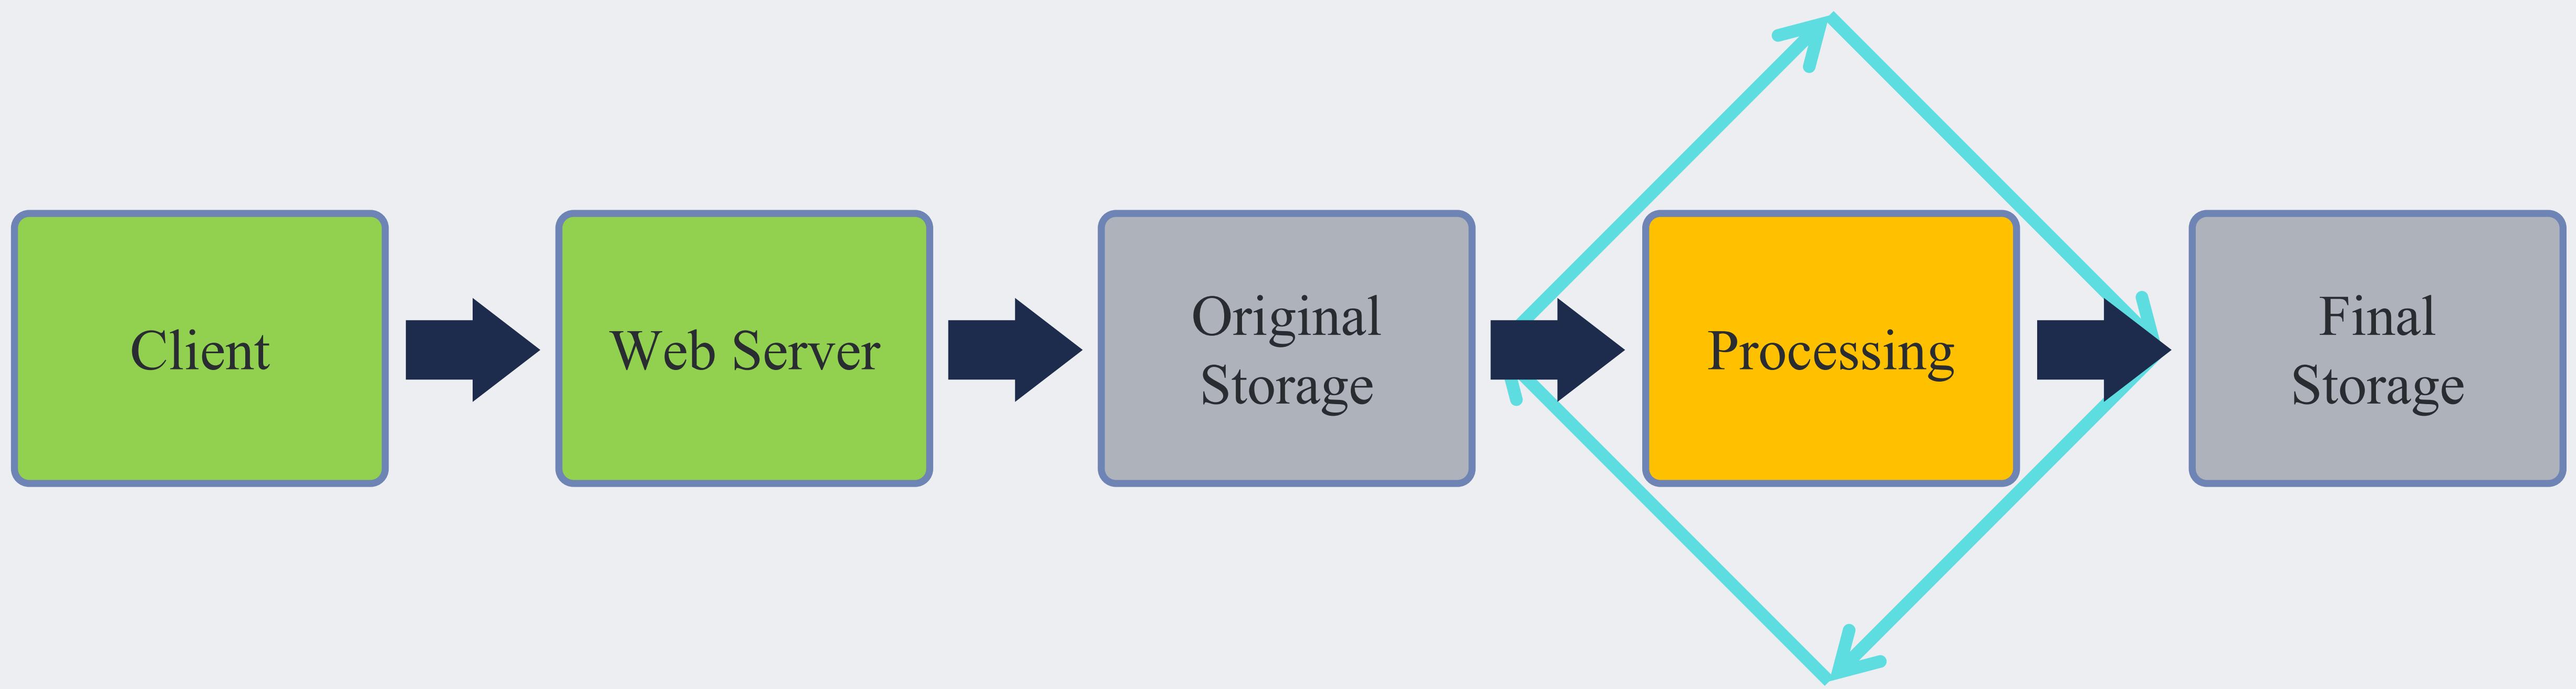


Focus: pre-sharing pipeline



All steps from when a user starts an upload until a video is ready to be shared

Serial pipeline leads to slow processing



Monolithic script slows development

“Let’s experiment speech recognition,
add a logic to extract audio and analysis”

“We need to change the thumbnail
generation logic for videos > x
minutes to create scene-based
scrubber preview”

“Pass-through for
small and well-
formatted videos”

“We want to experiment AI-based
encodings to spend 10x CPU for 30%
compression improvement on
popular videos”

“Change color
coding at
different time”

Processing

Final Blob
Storage

Challenges for video processing @ FB

Speedy

Users can share videos quickly

Flexible

Thousands of engineers can write pipelines for tens of apps

Robust

Handle faults and overload that is inevitable at scale

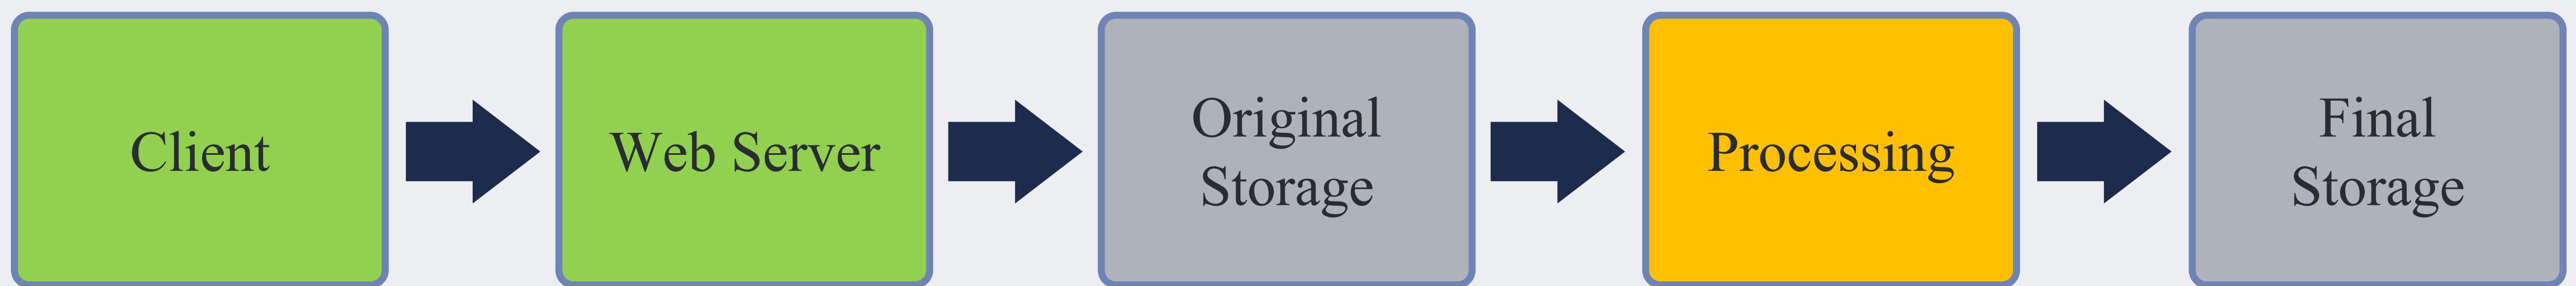
Our Streaming Video Engine (SVE)
is speedy, flexible, and robust

Speedy: harness parallelism

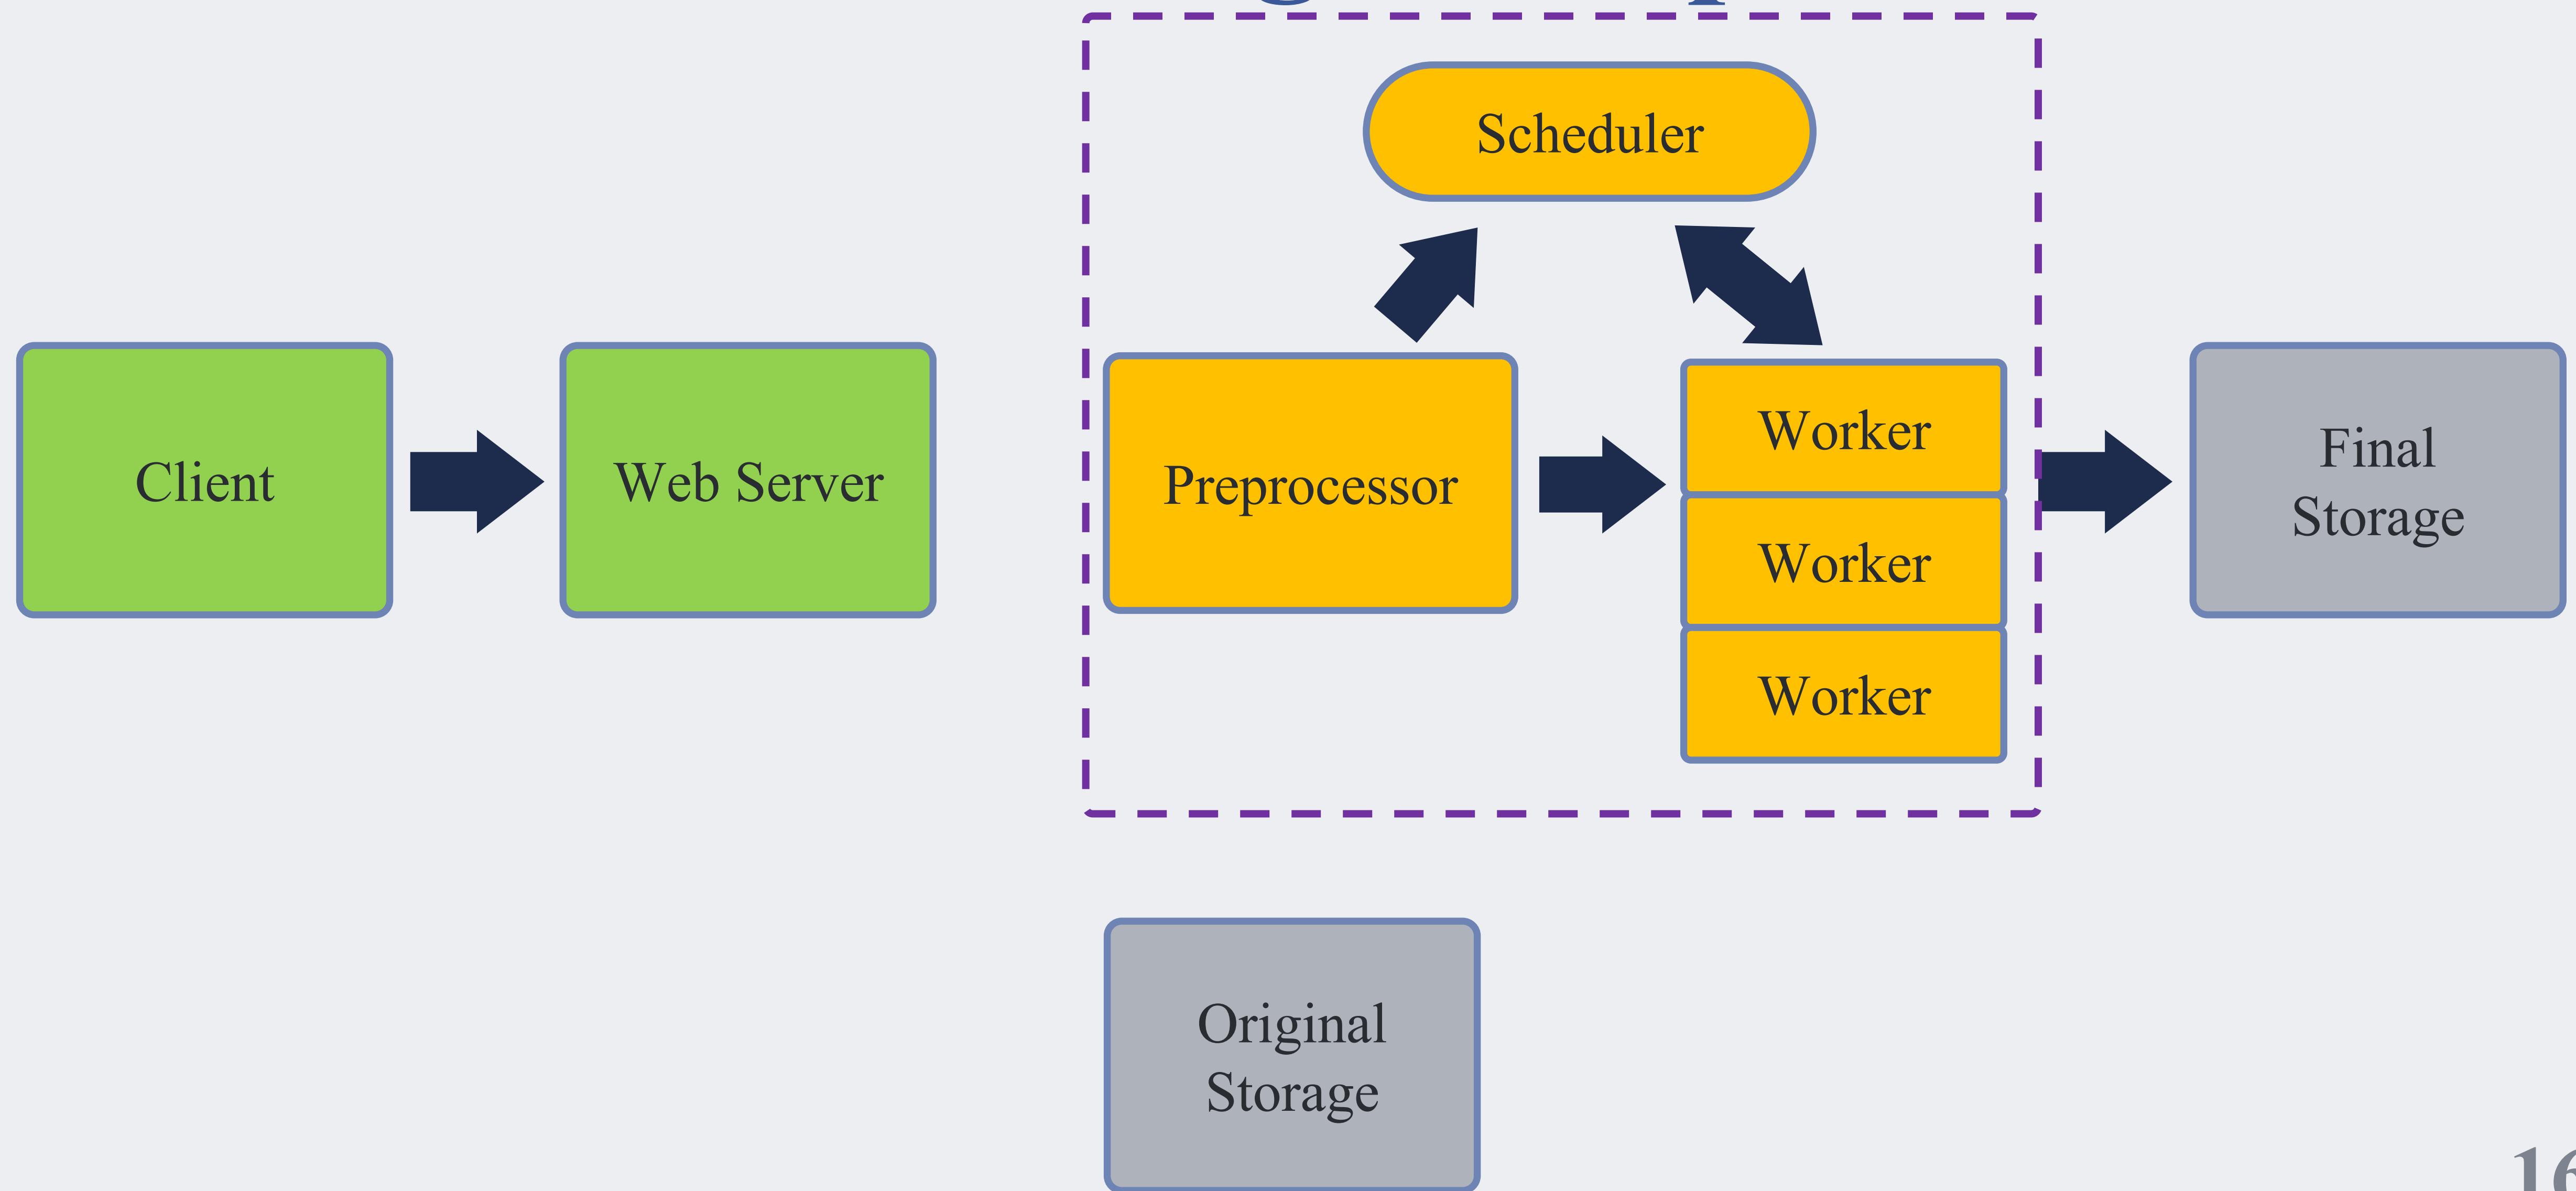
Users can share videos quickly

- Overlap fault tolerance and processing
- Overlap upload and processing
- Parallel processing

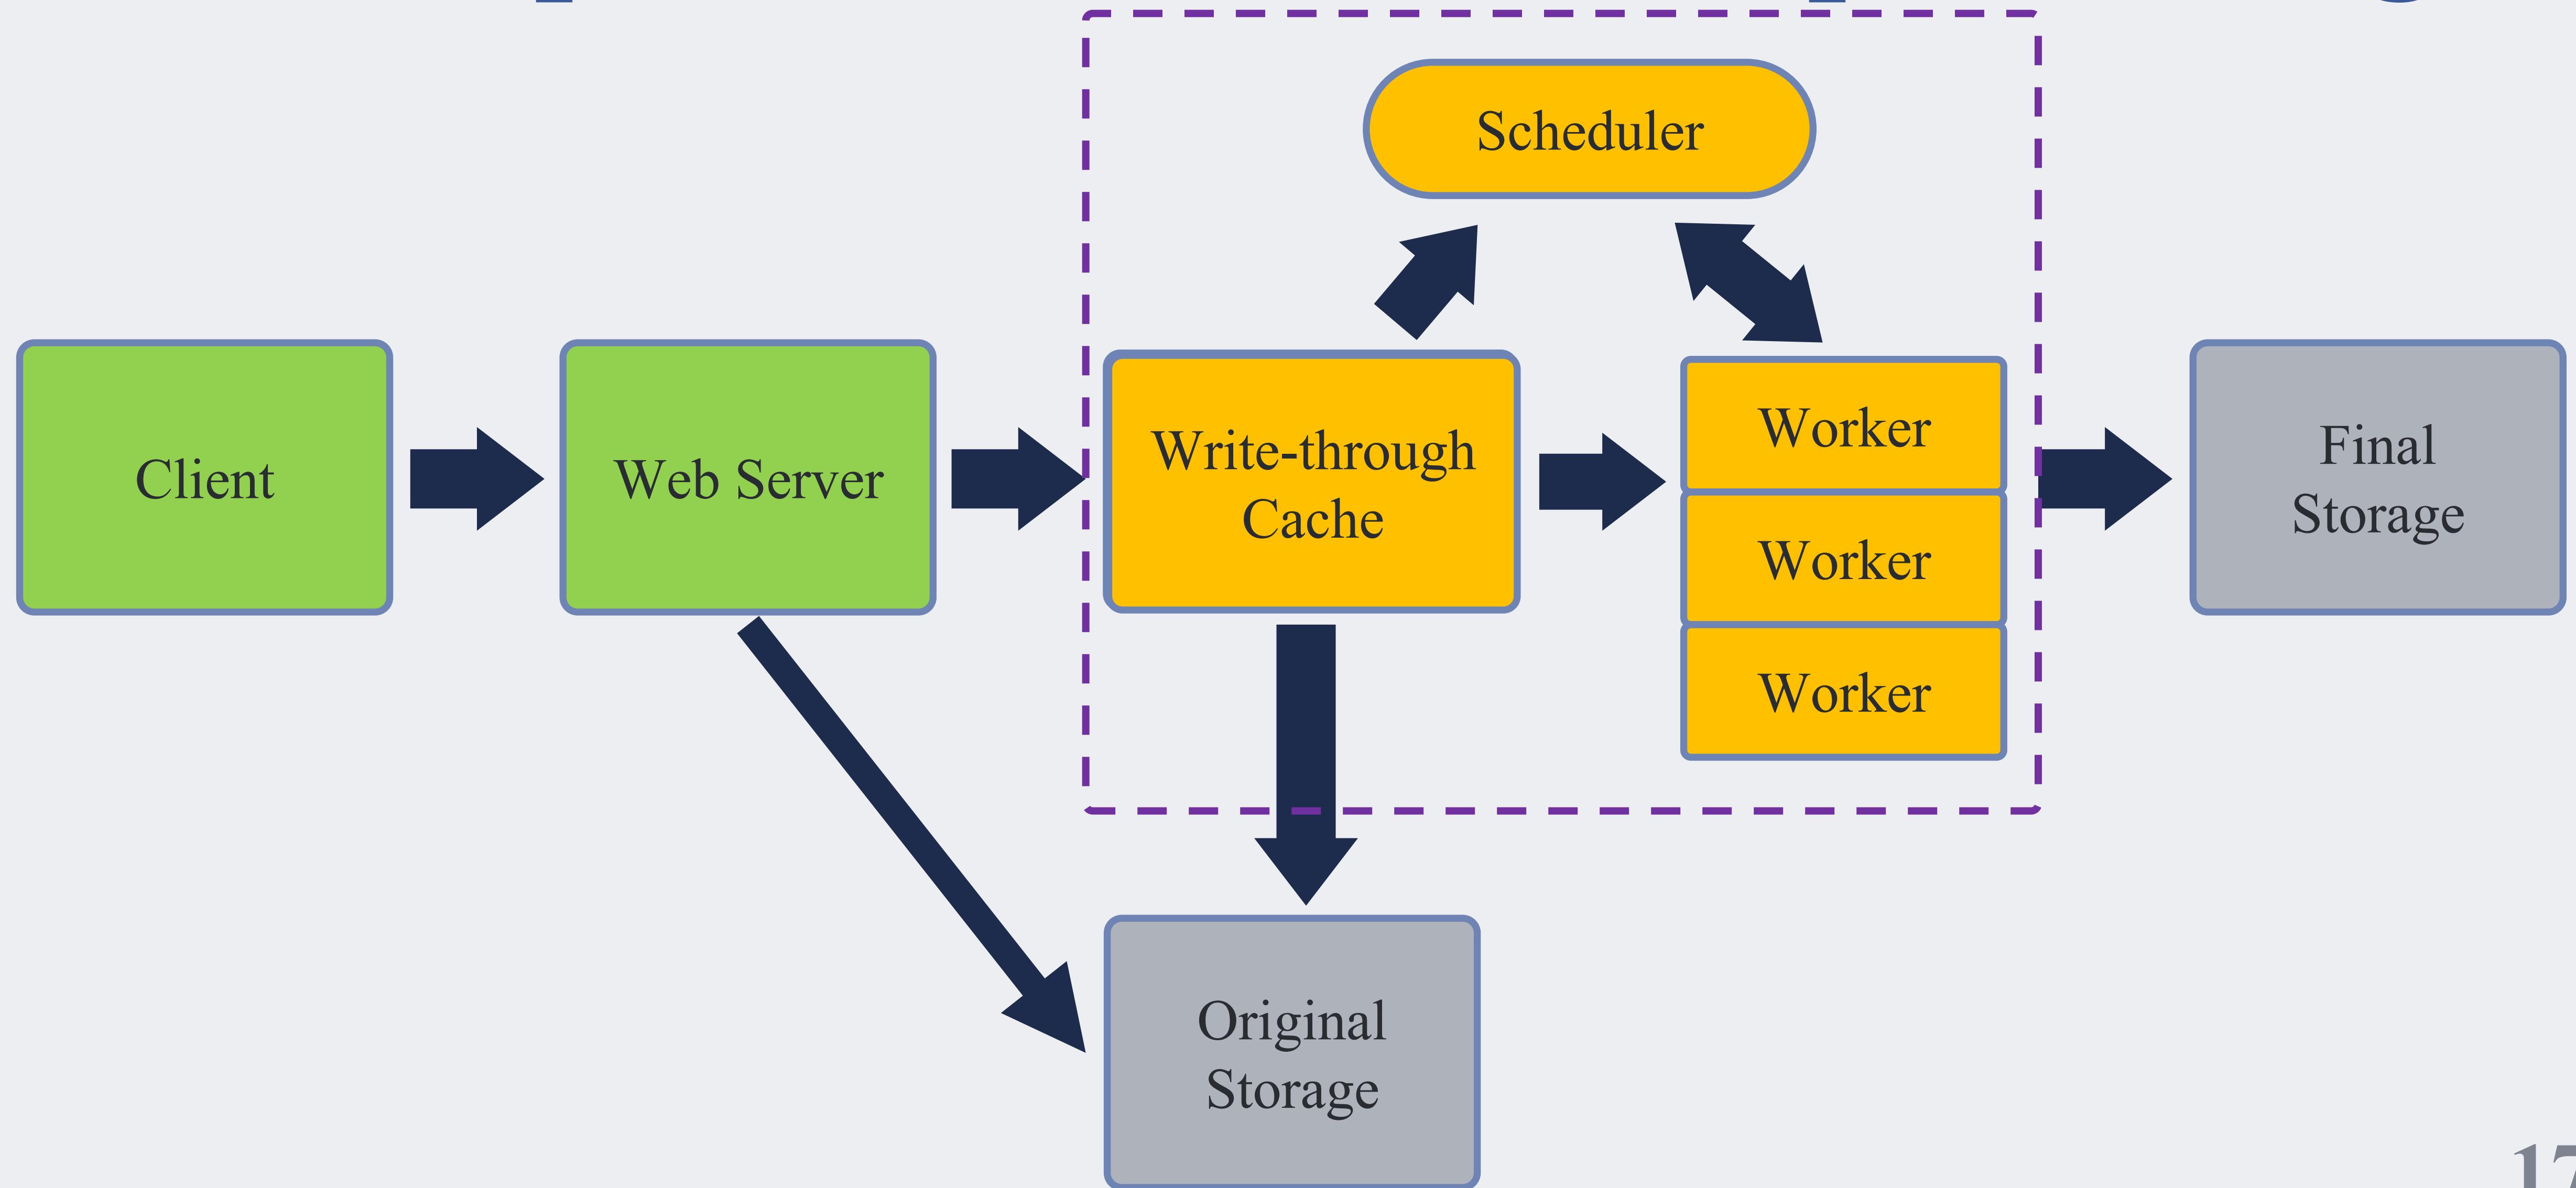
Architectural changes for parallelism



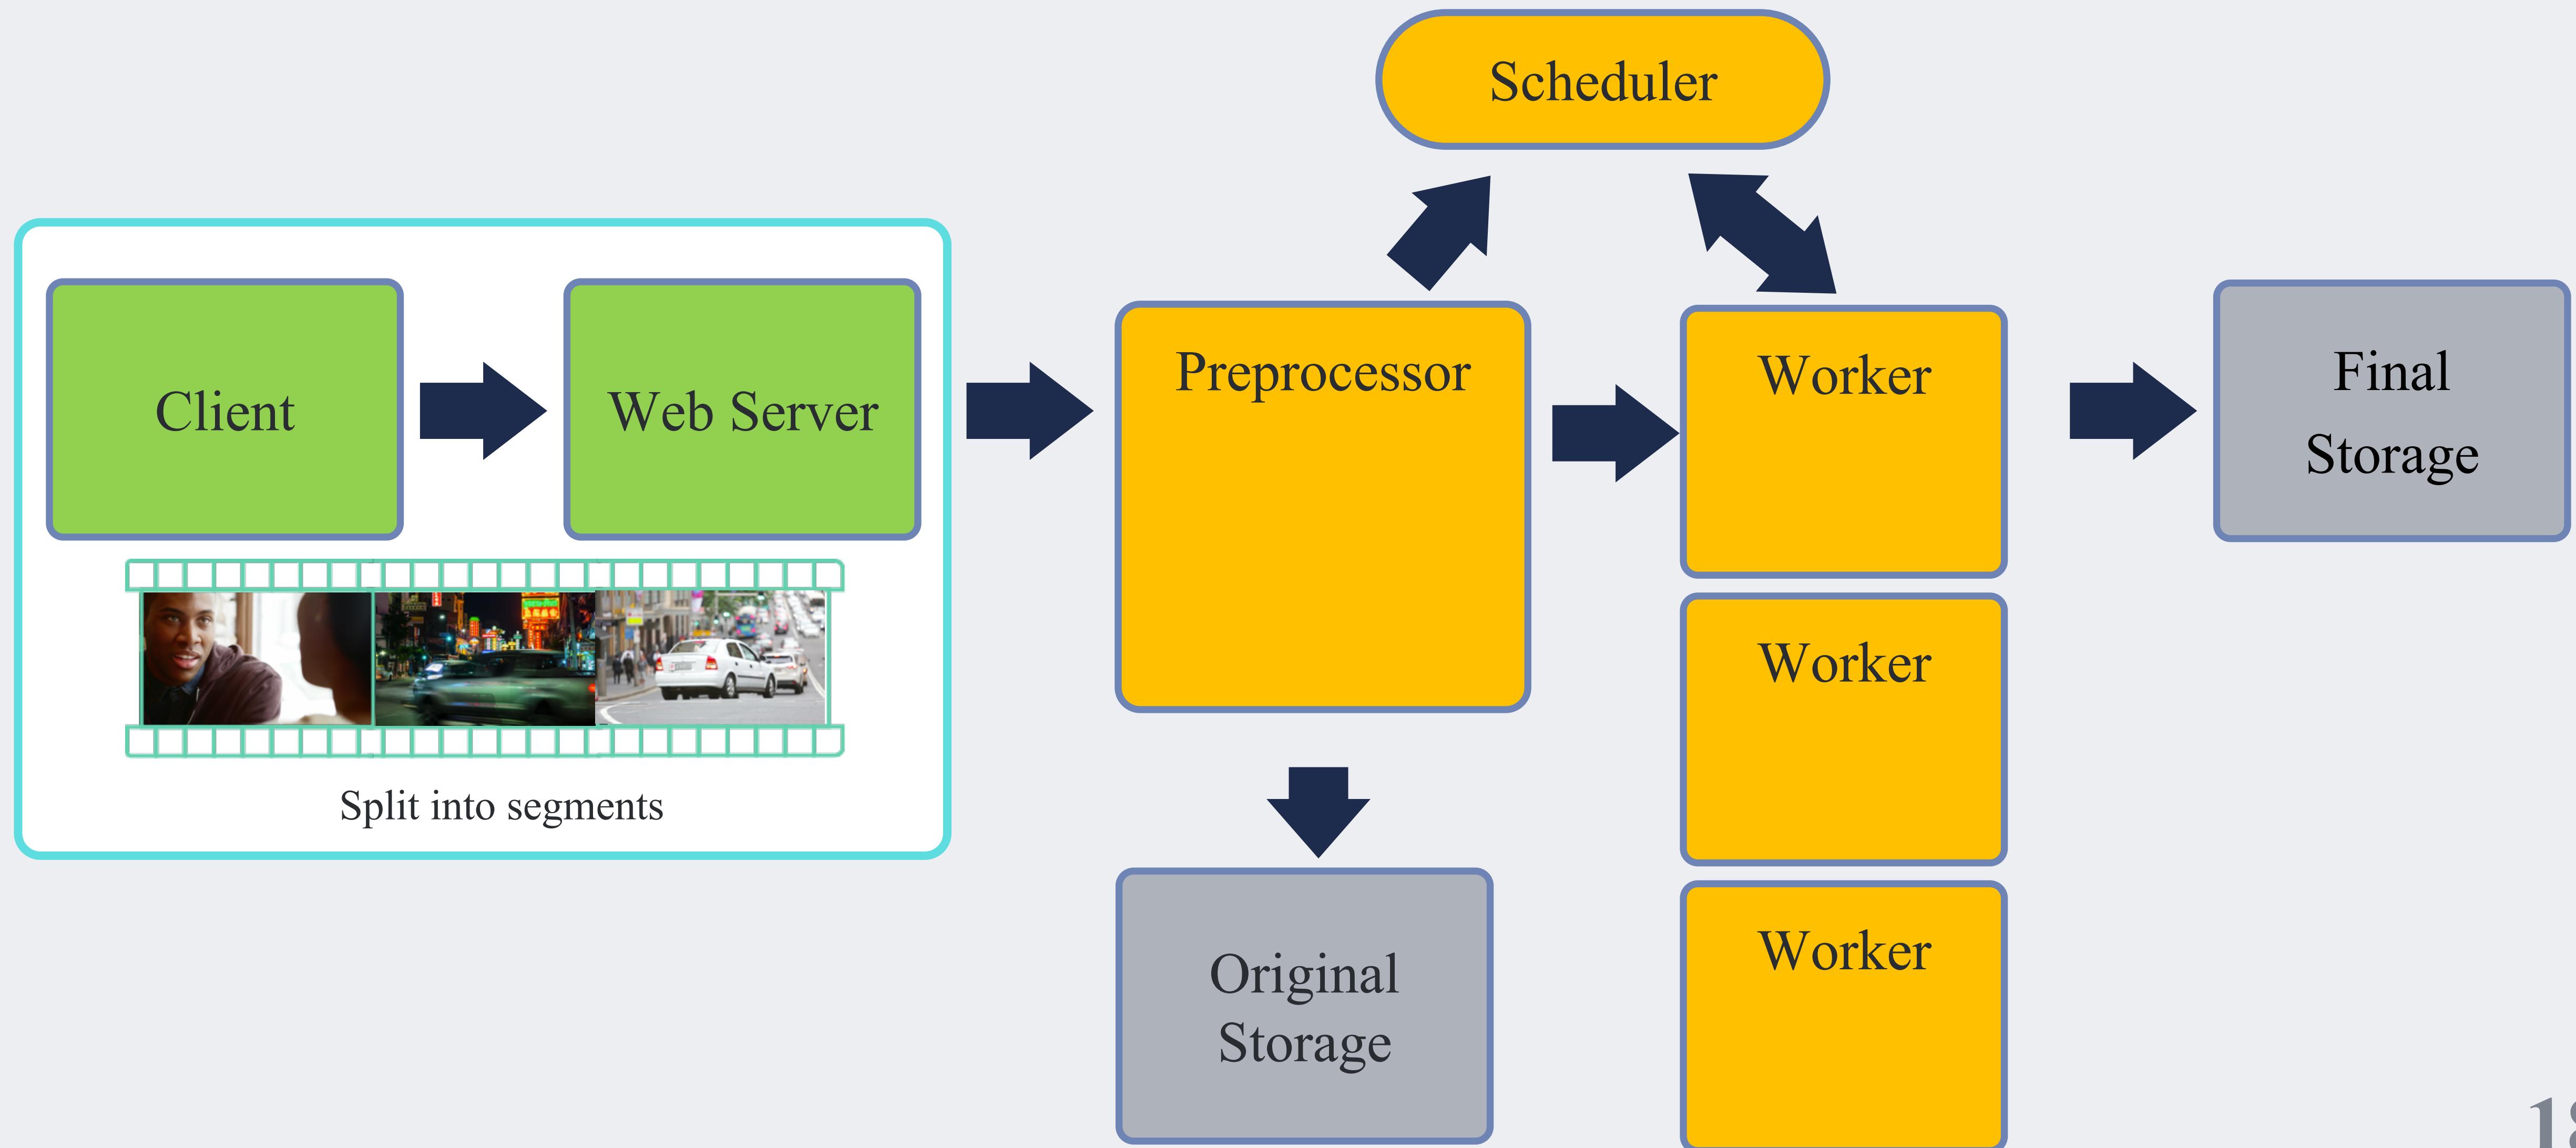
Architectural changes for parallelism



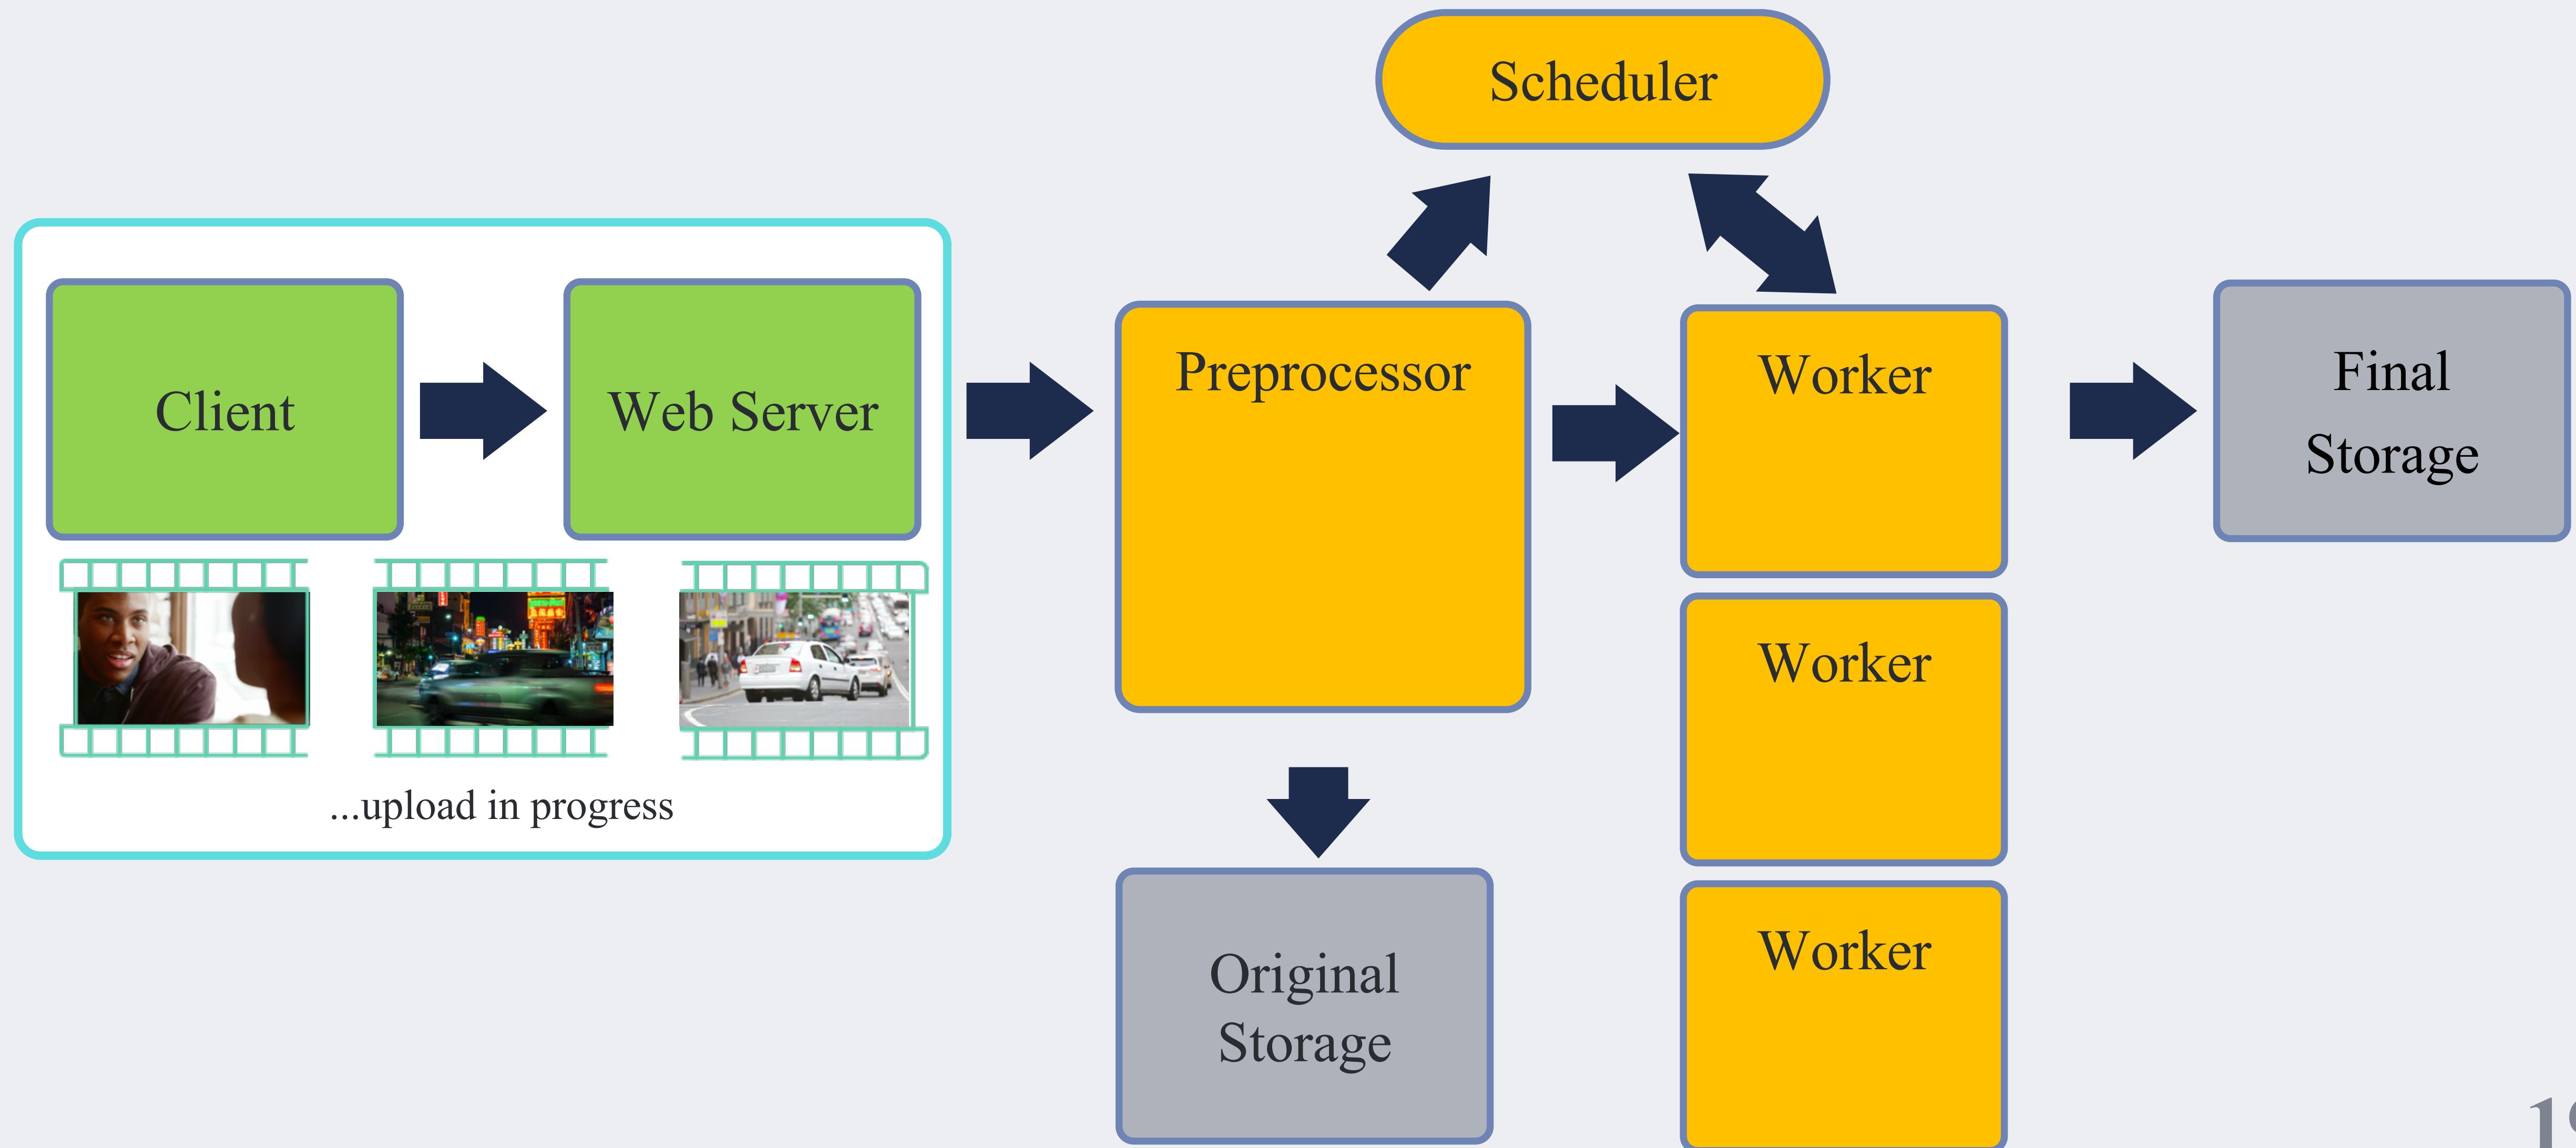
Overlap fault tolerance and processing



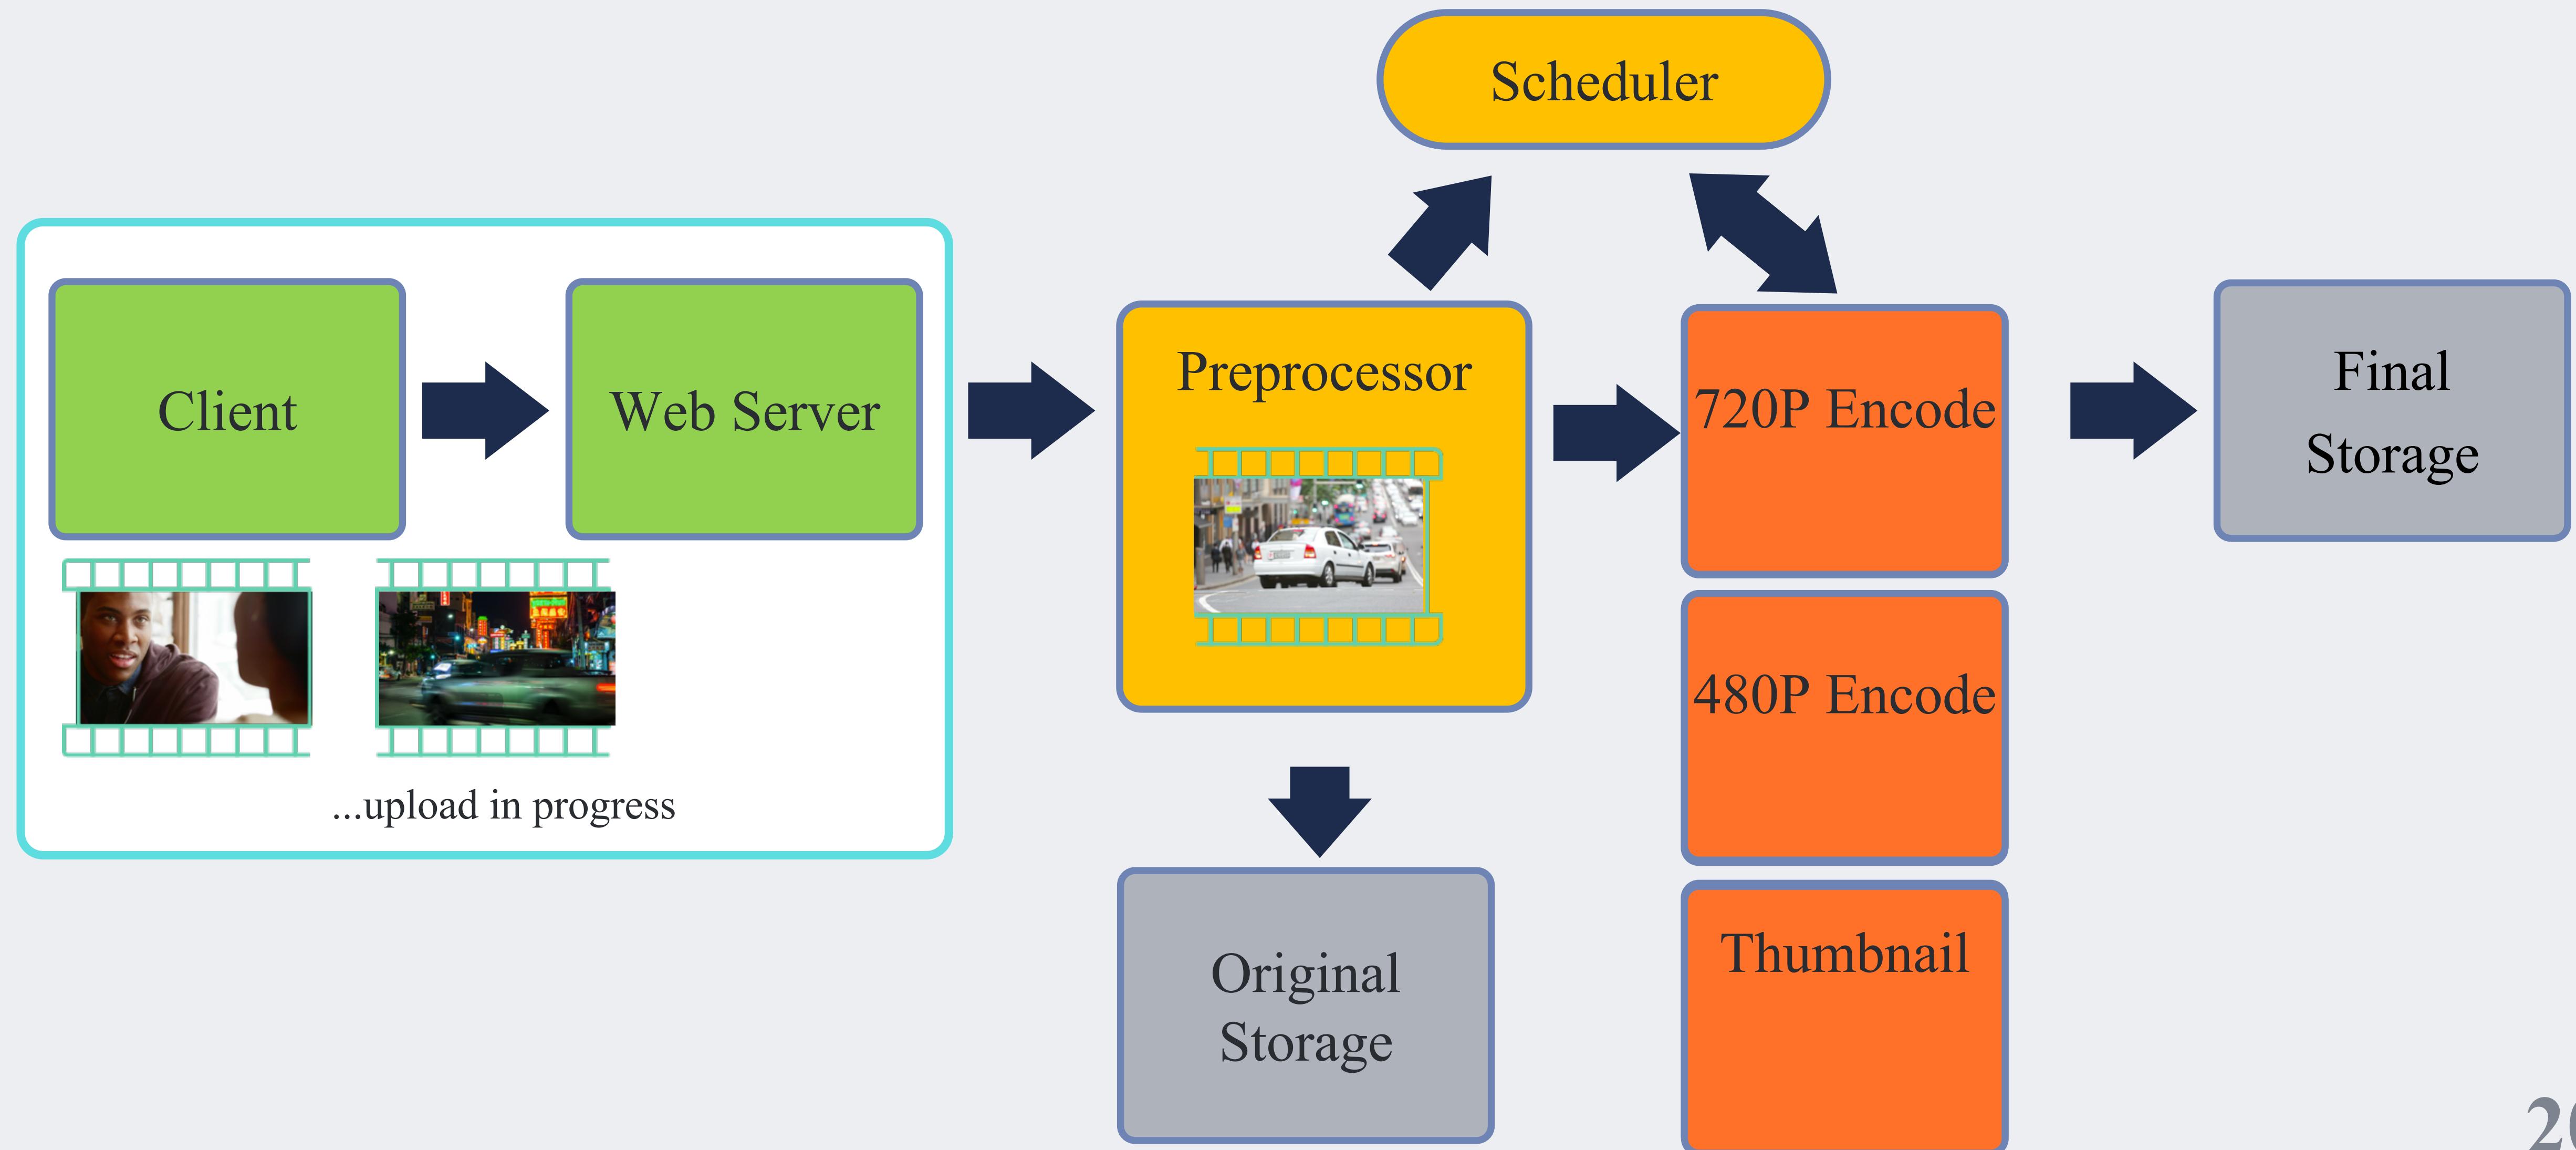
Overlap upload and processing



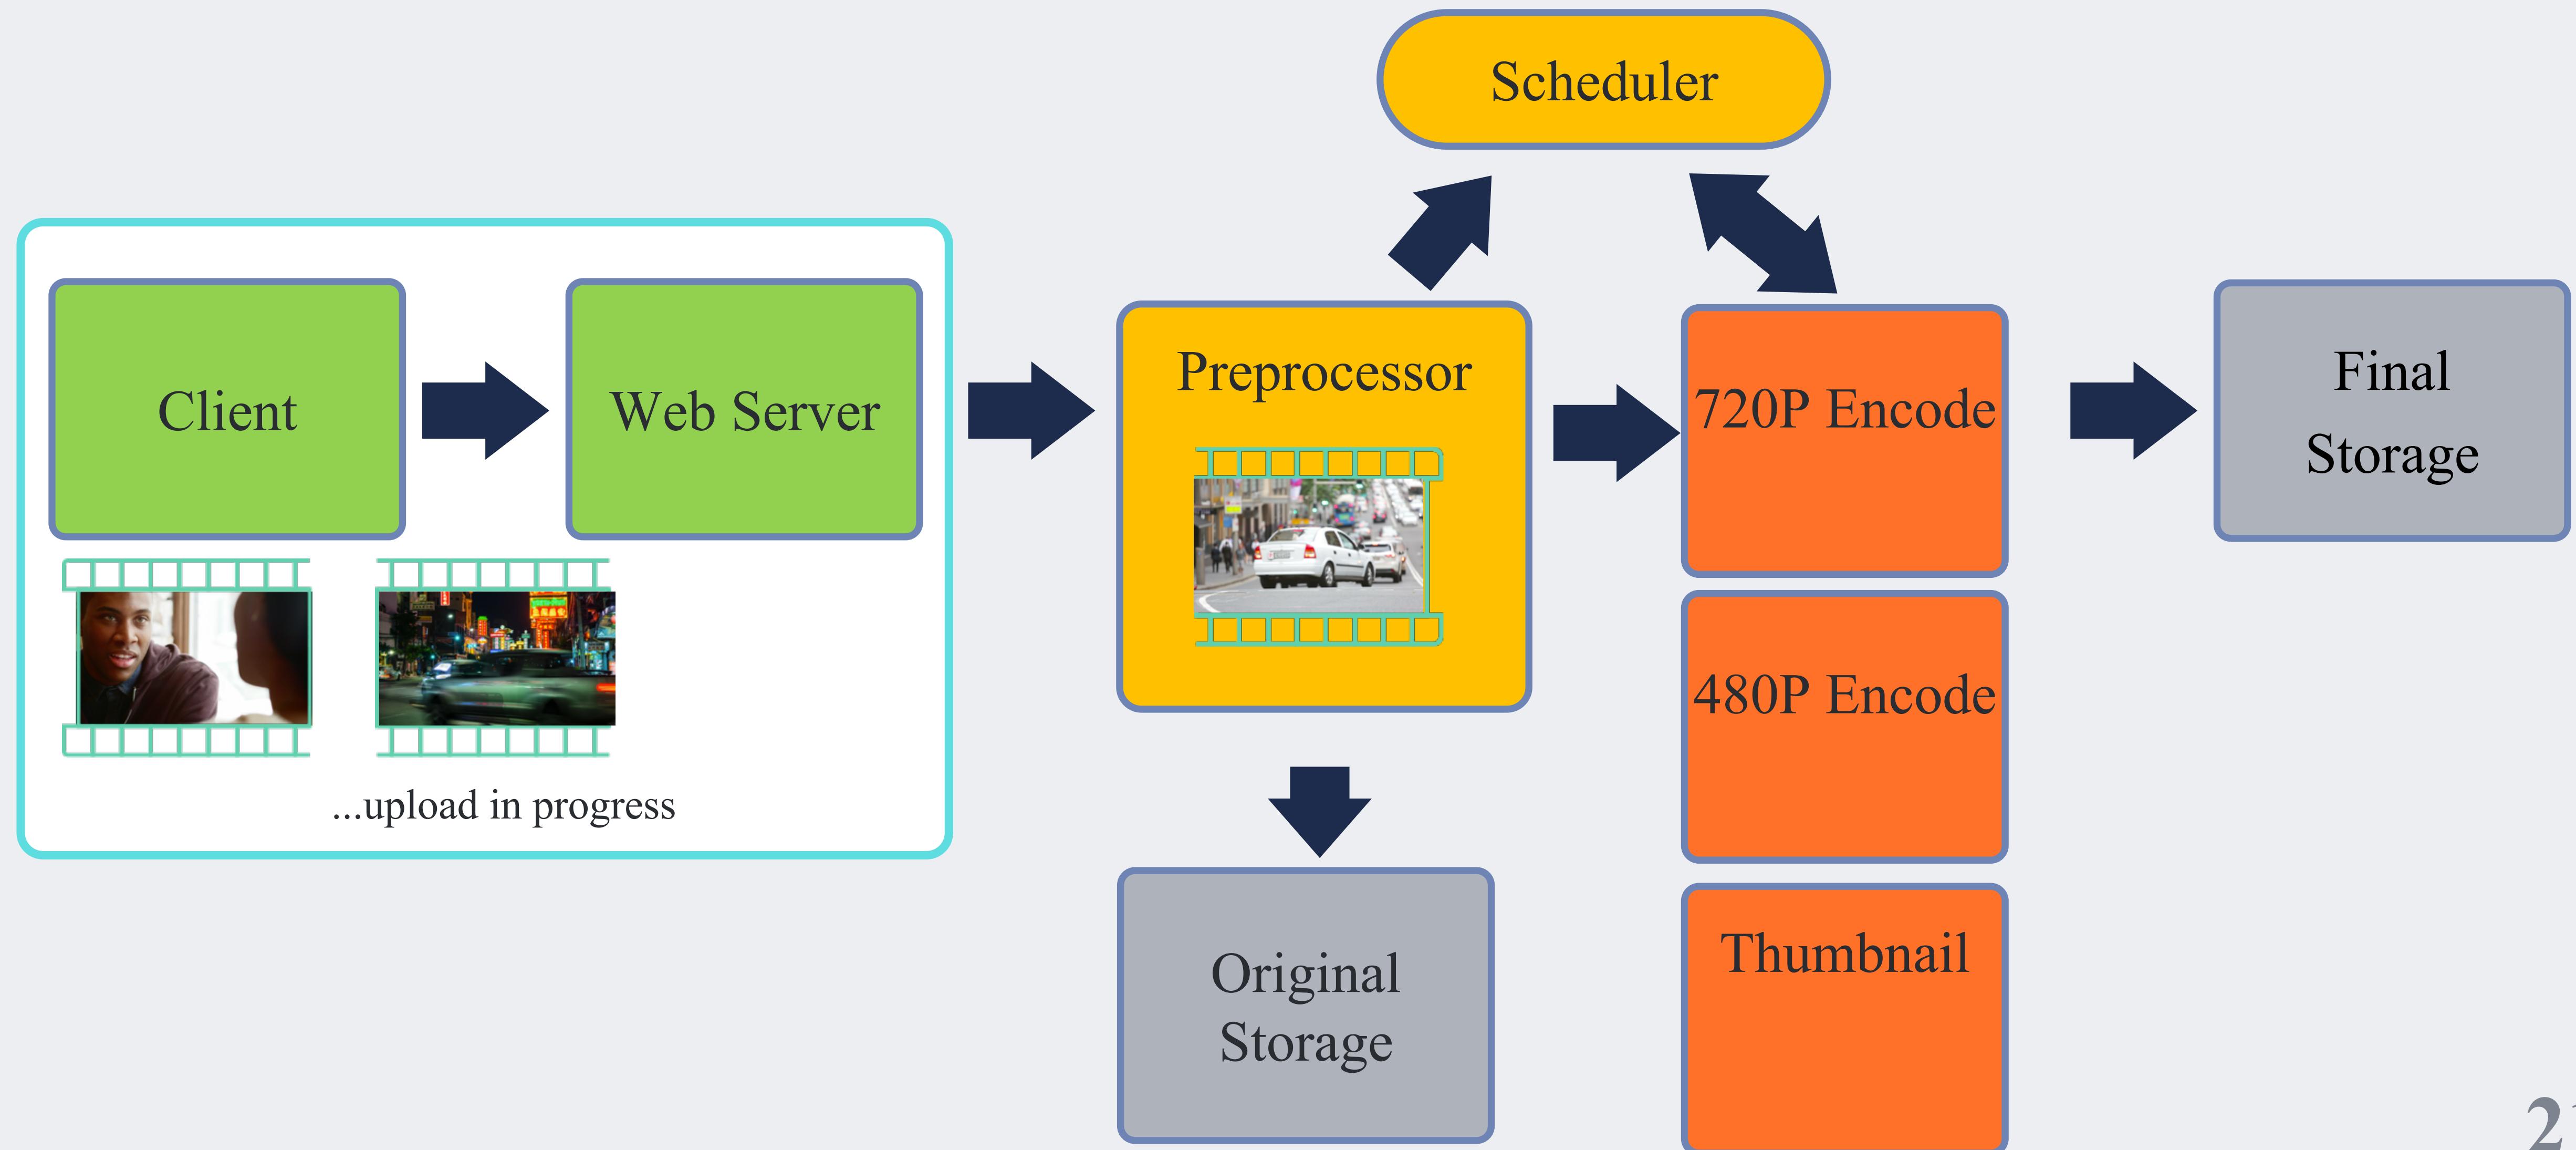
Overlap upload and processing



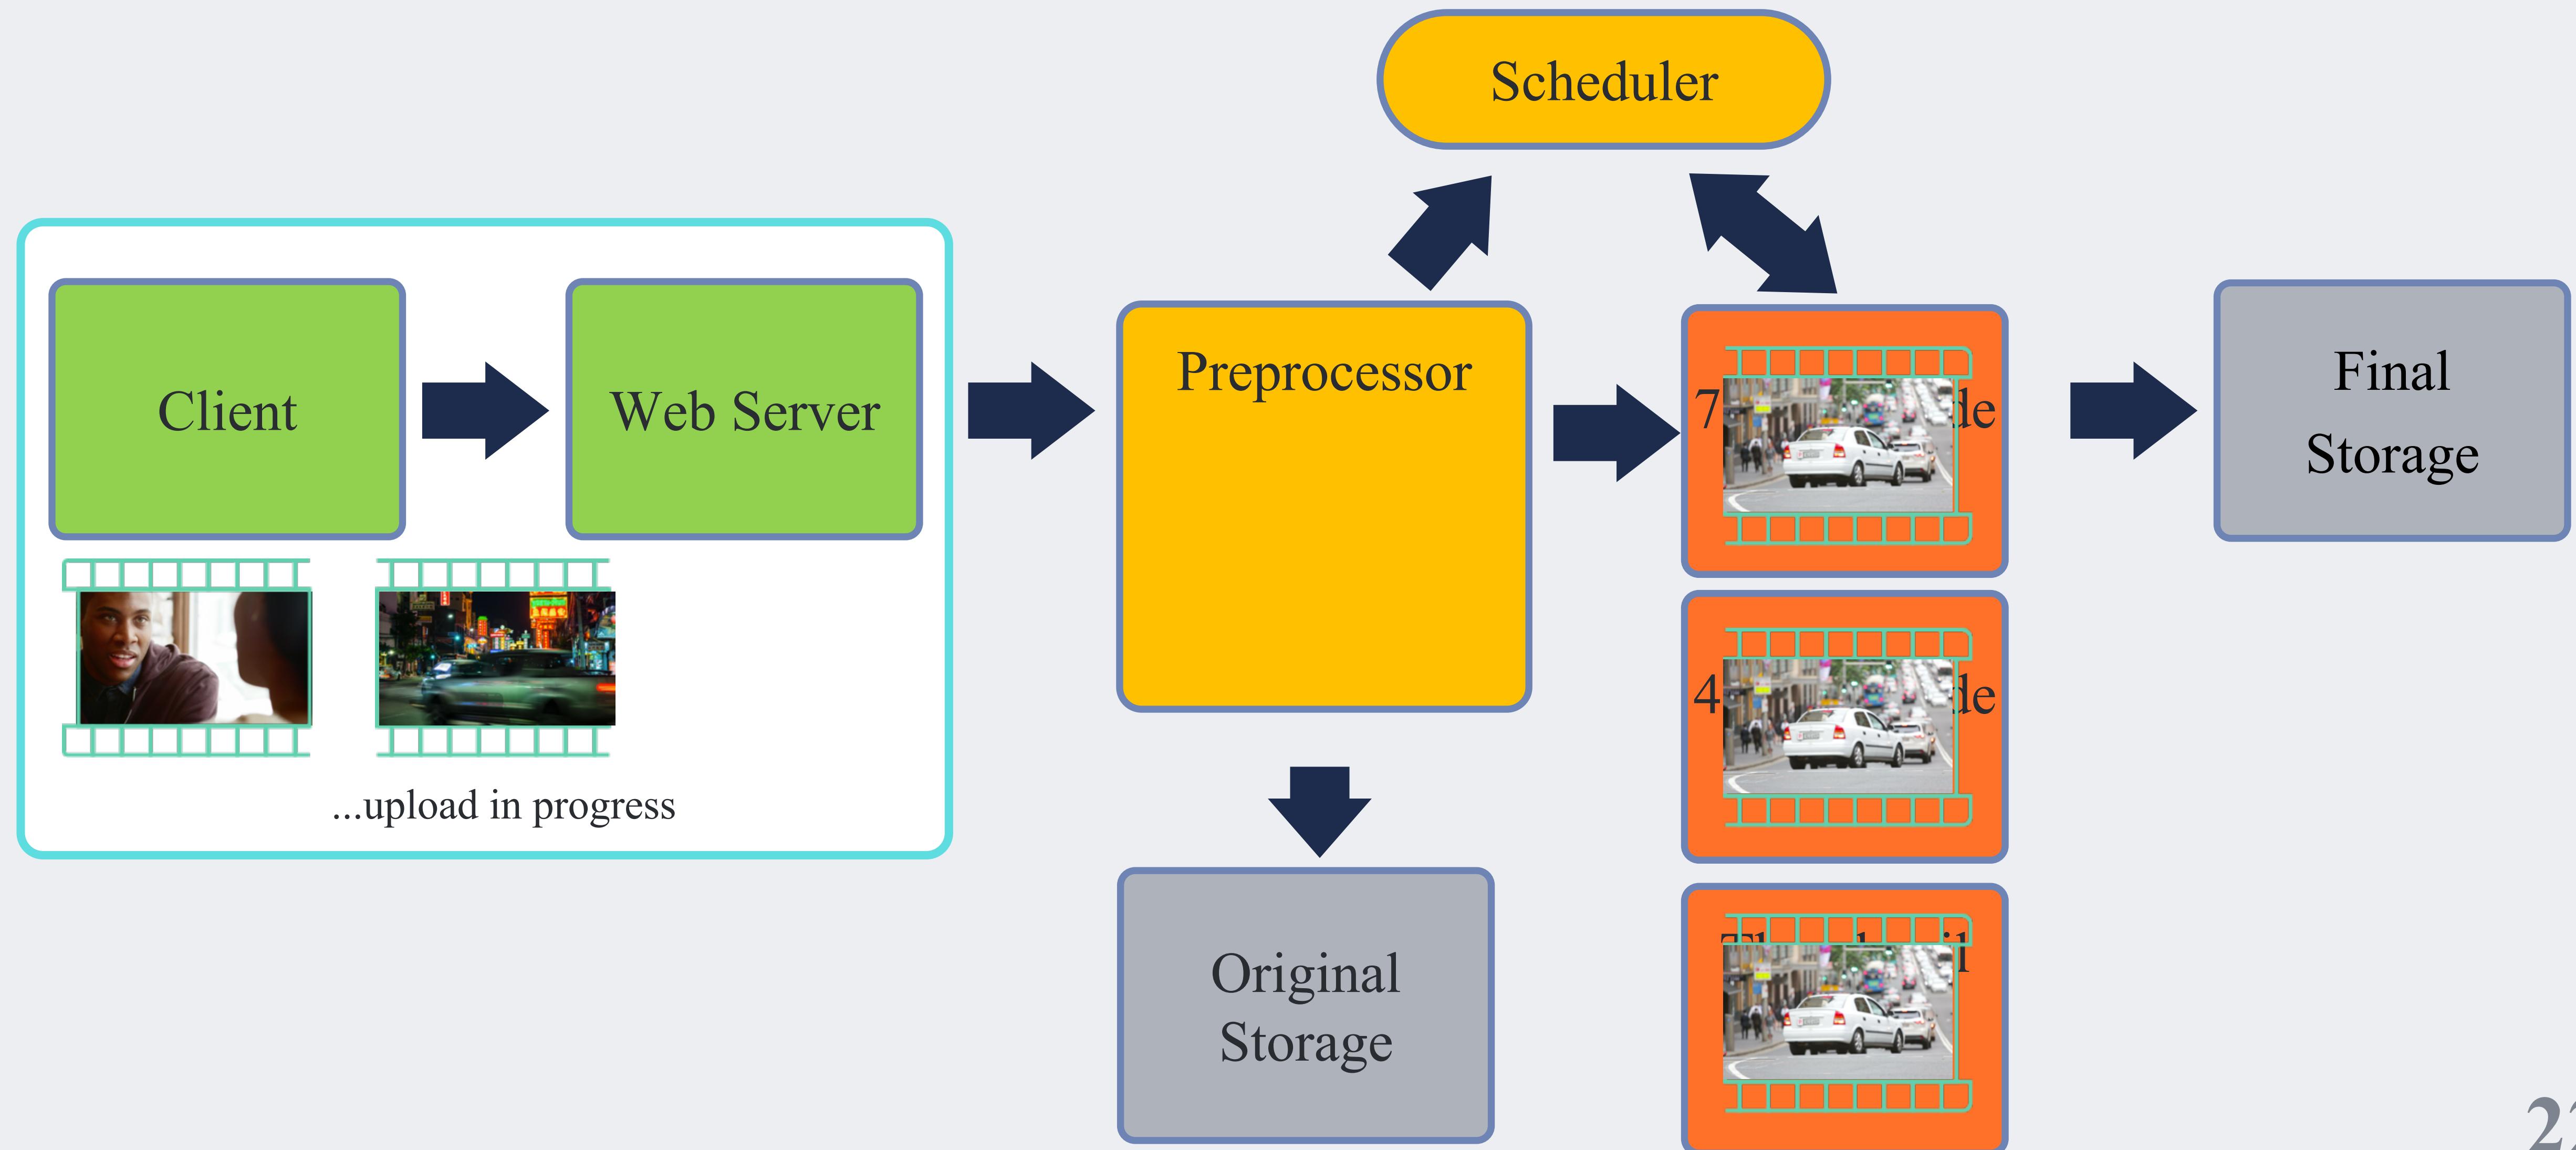
Parallel processing w/ many workers



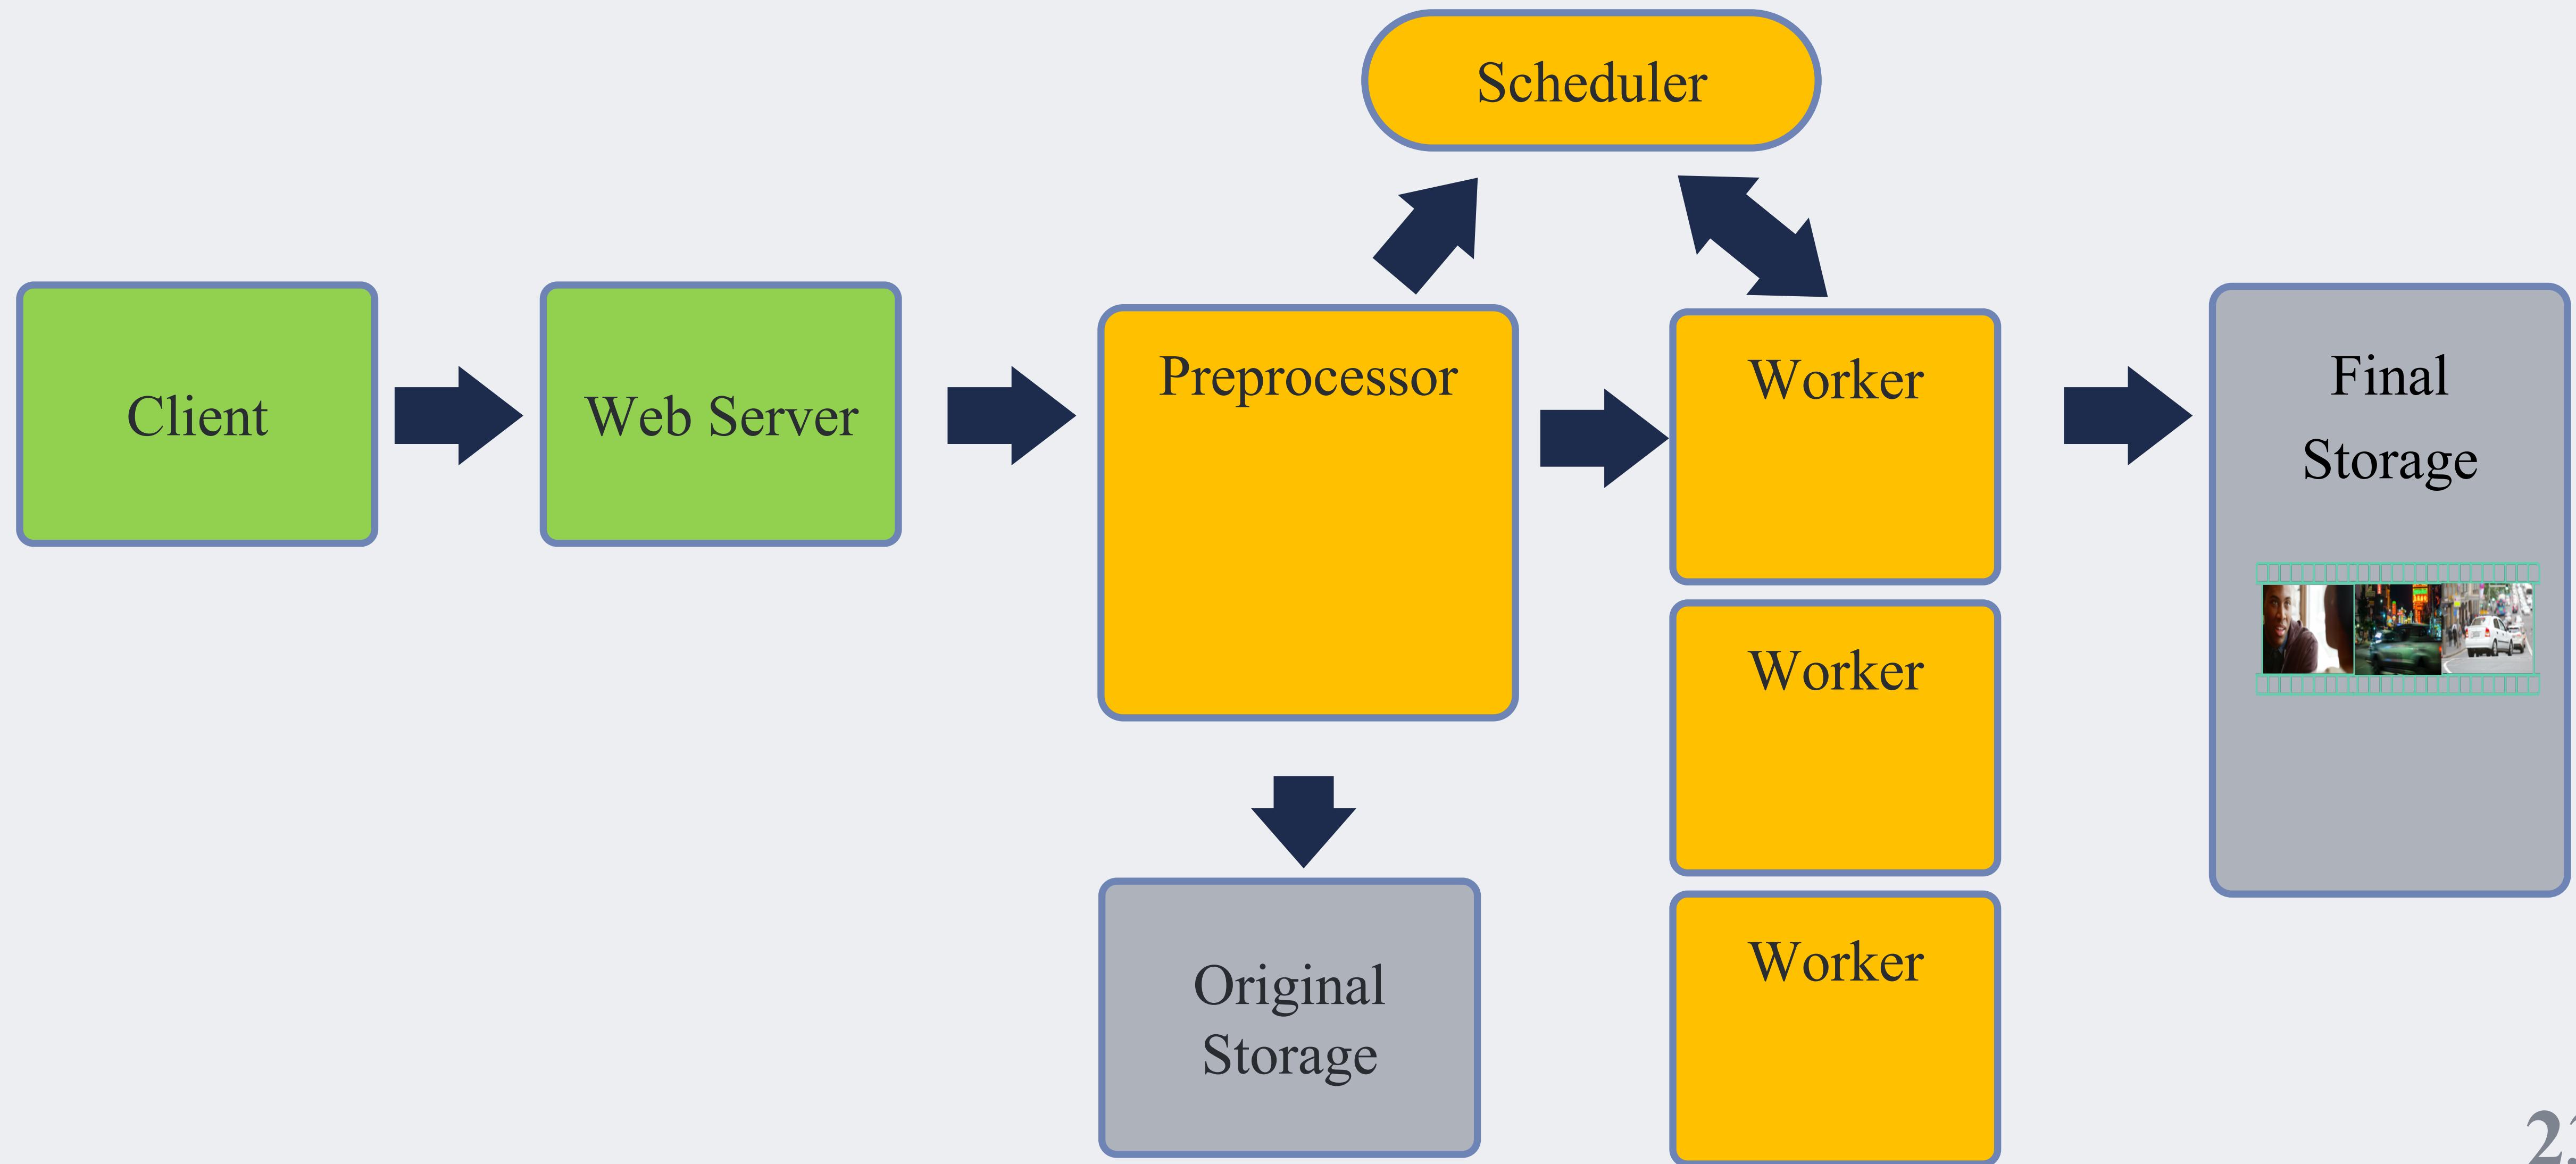
Parallel processing w/ many workers



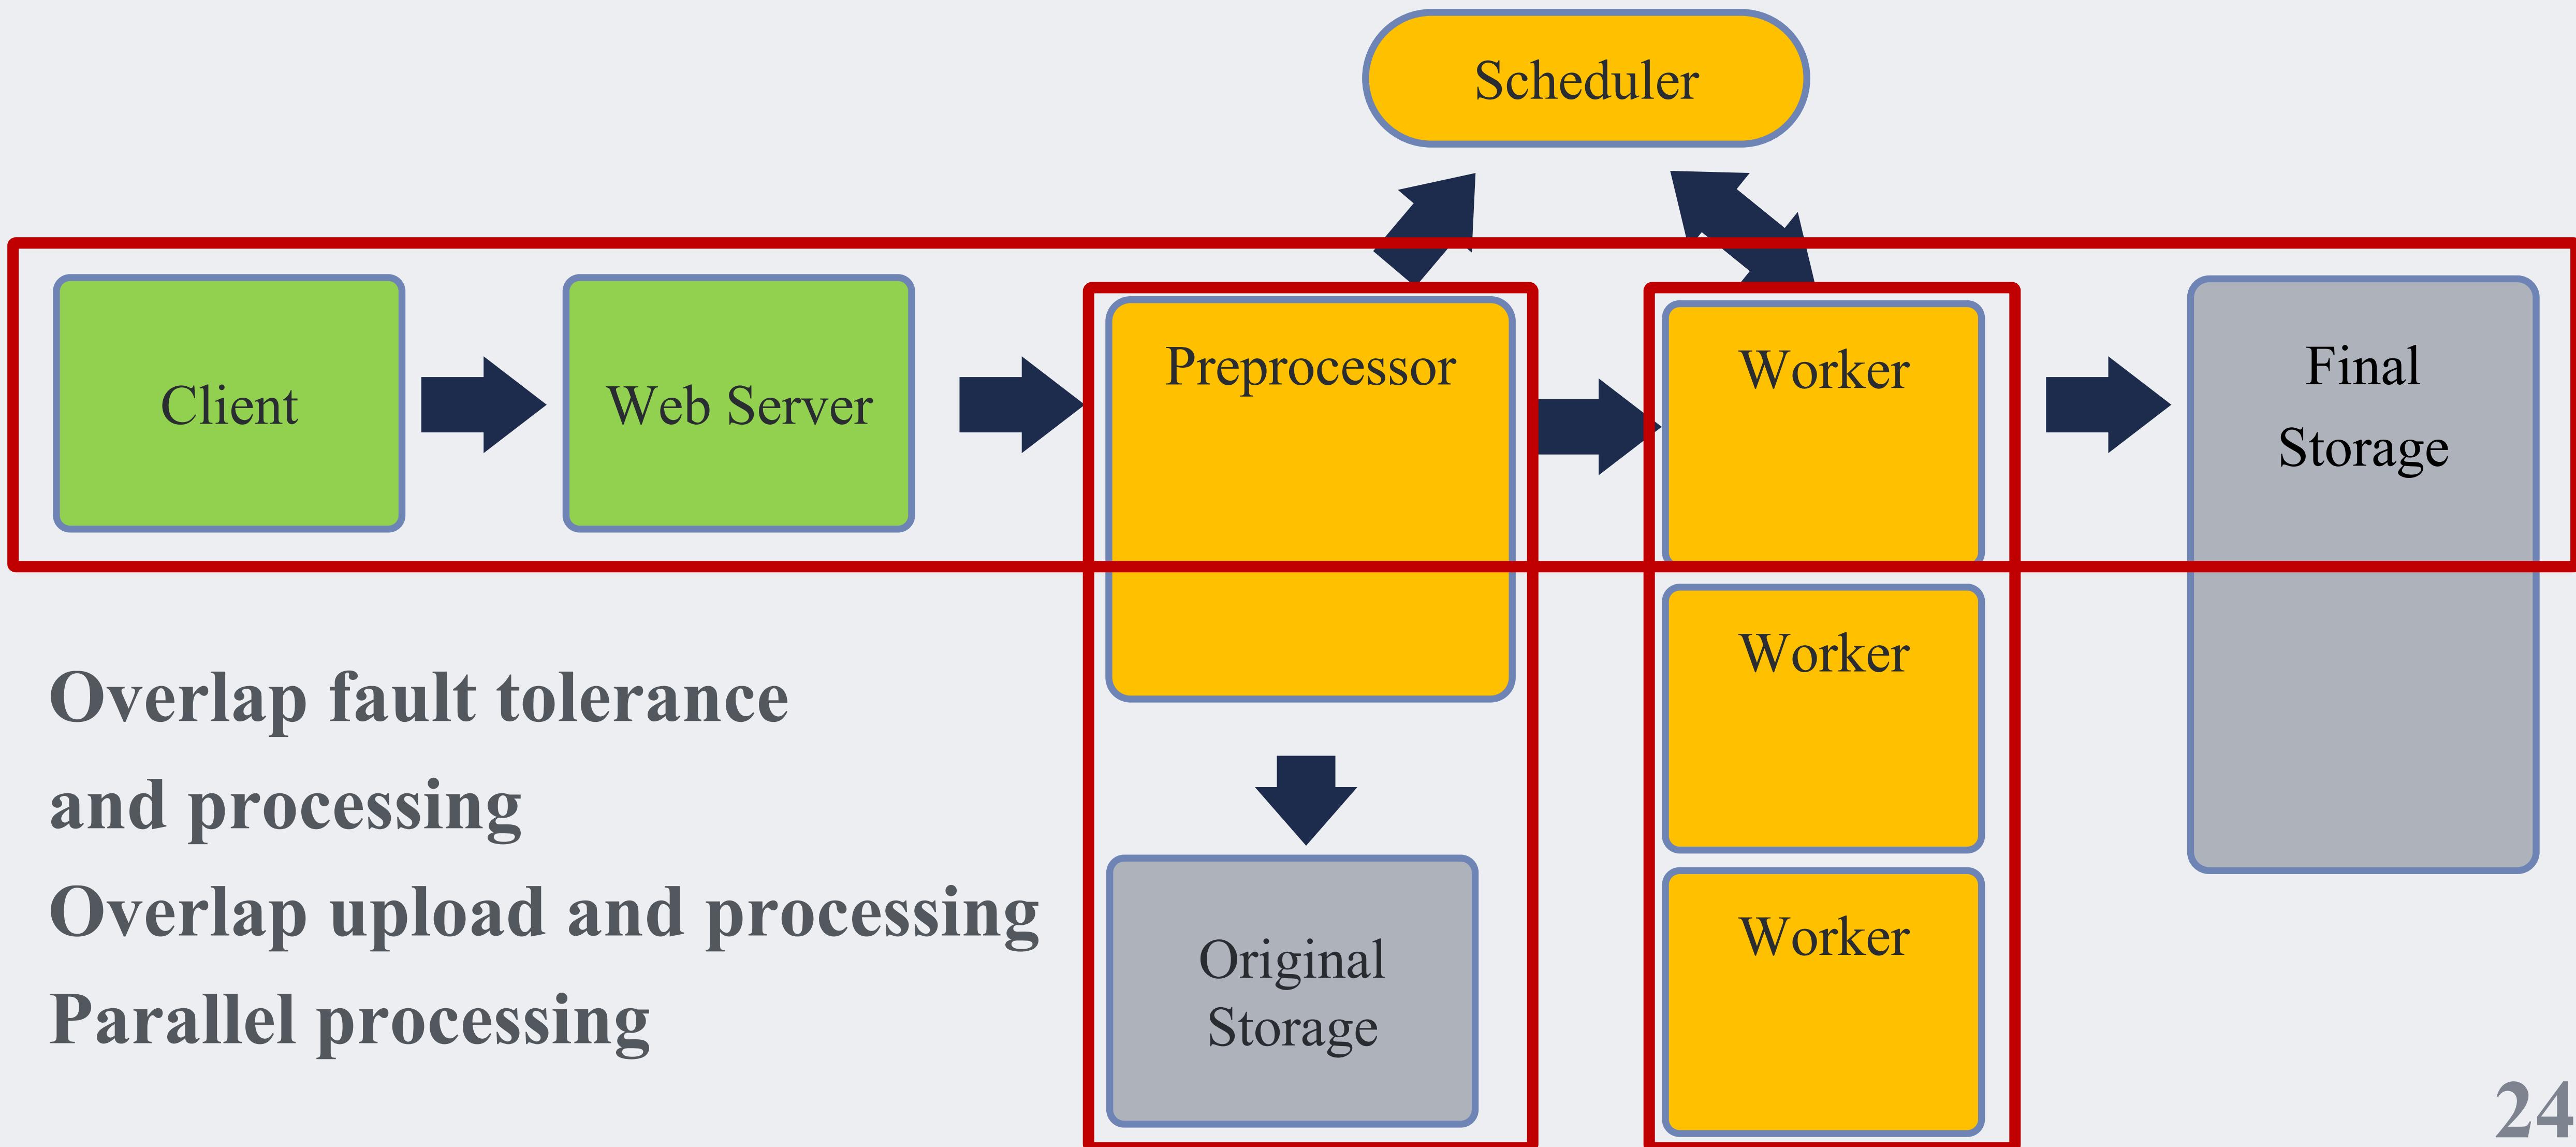
Parallel processing w/ many workers



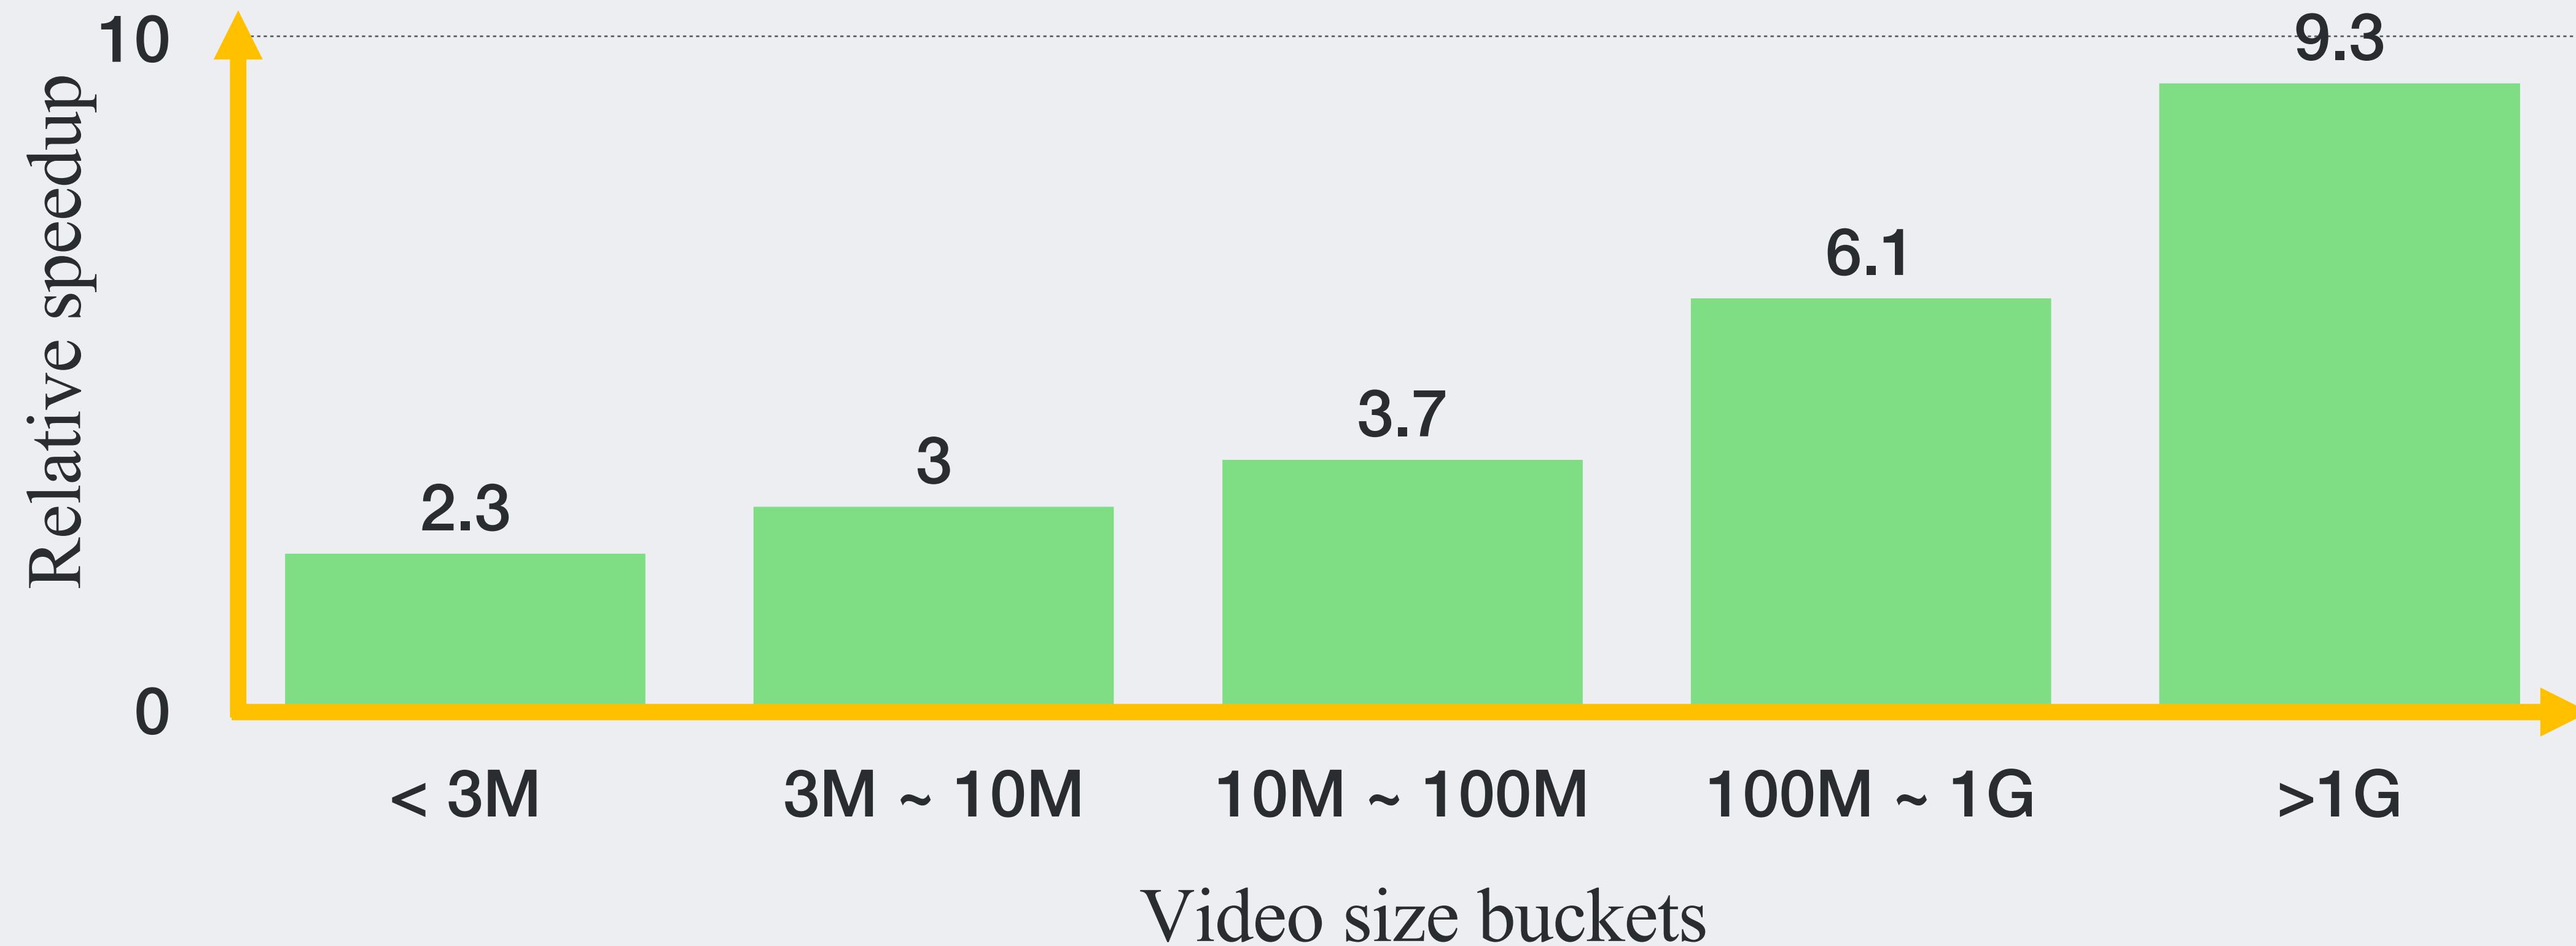
Parallel processing w/ many workers



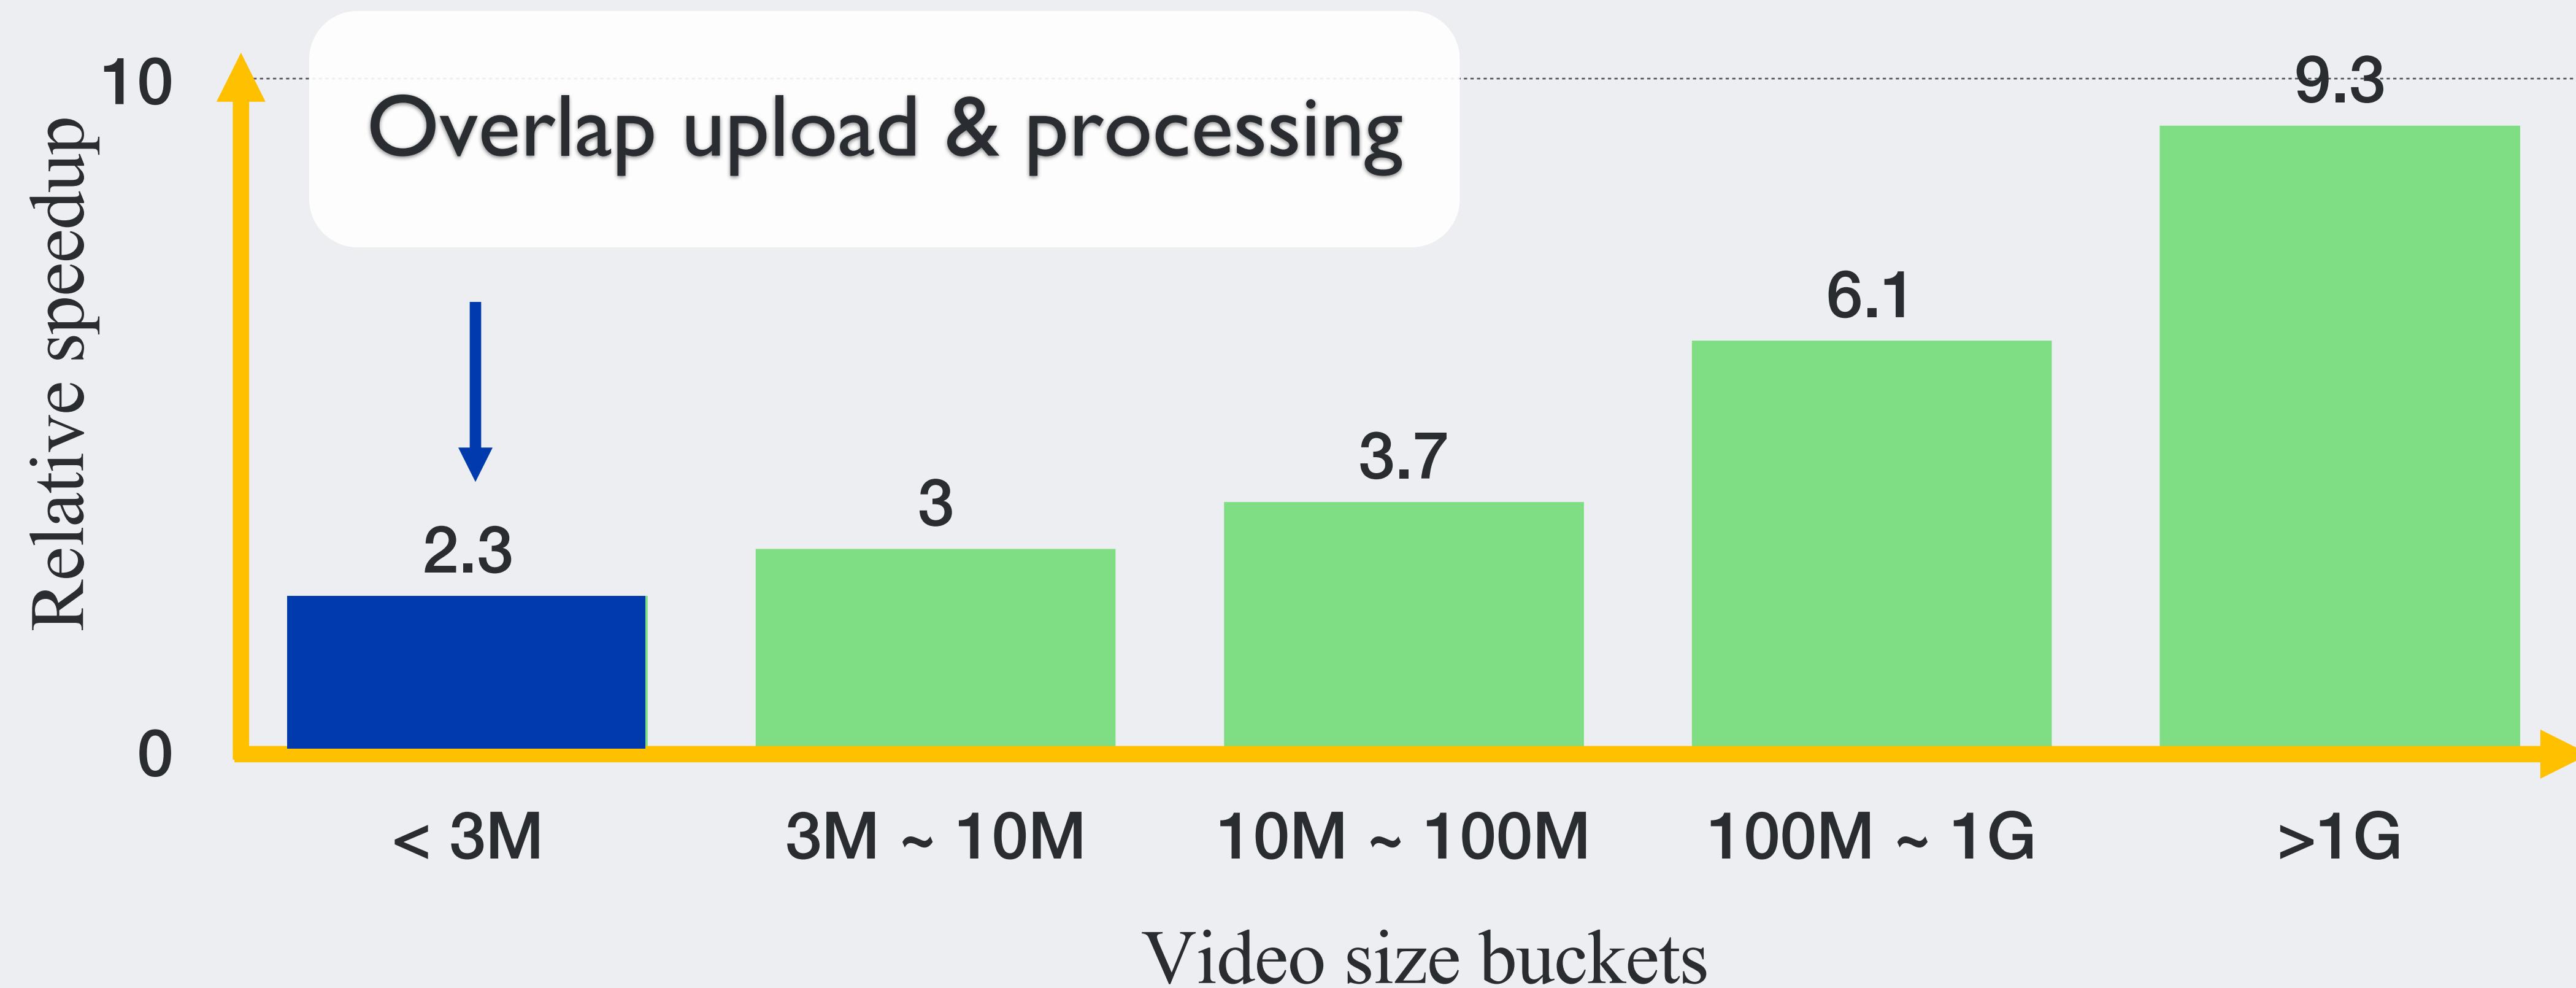
Three sources of parallelism



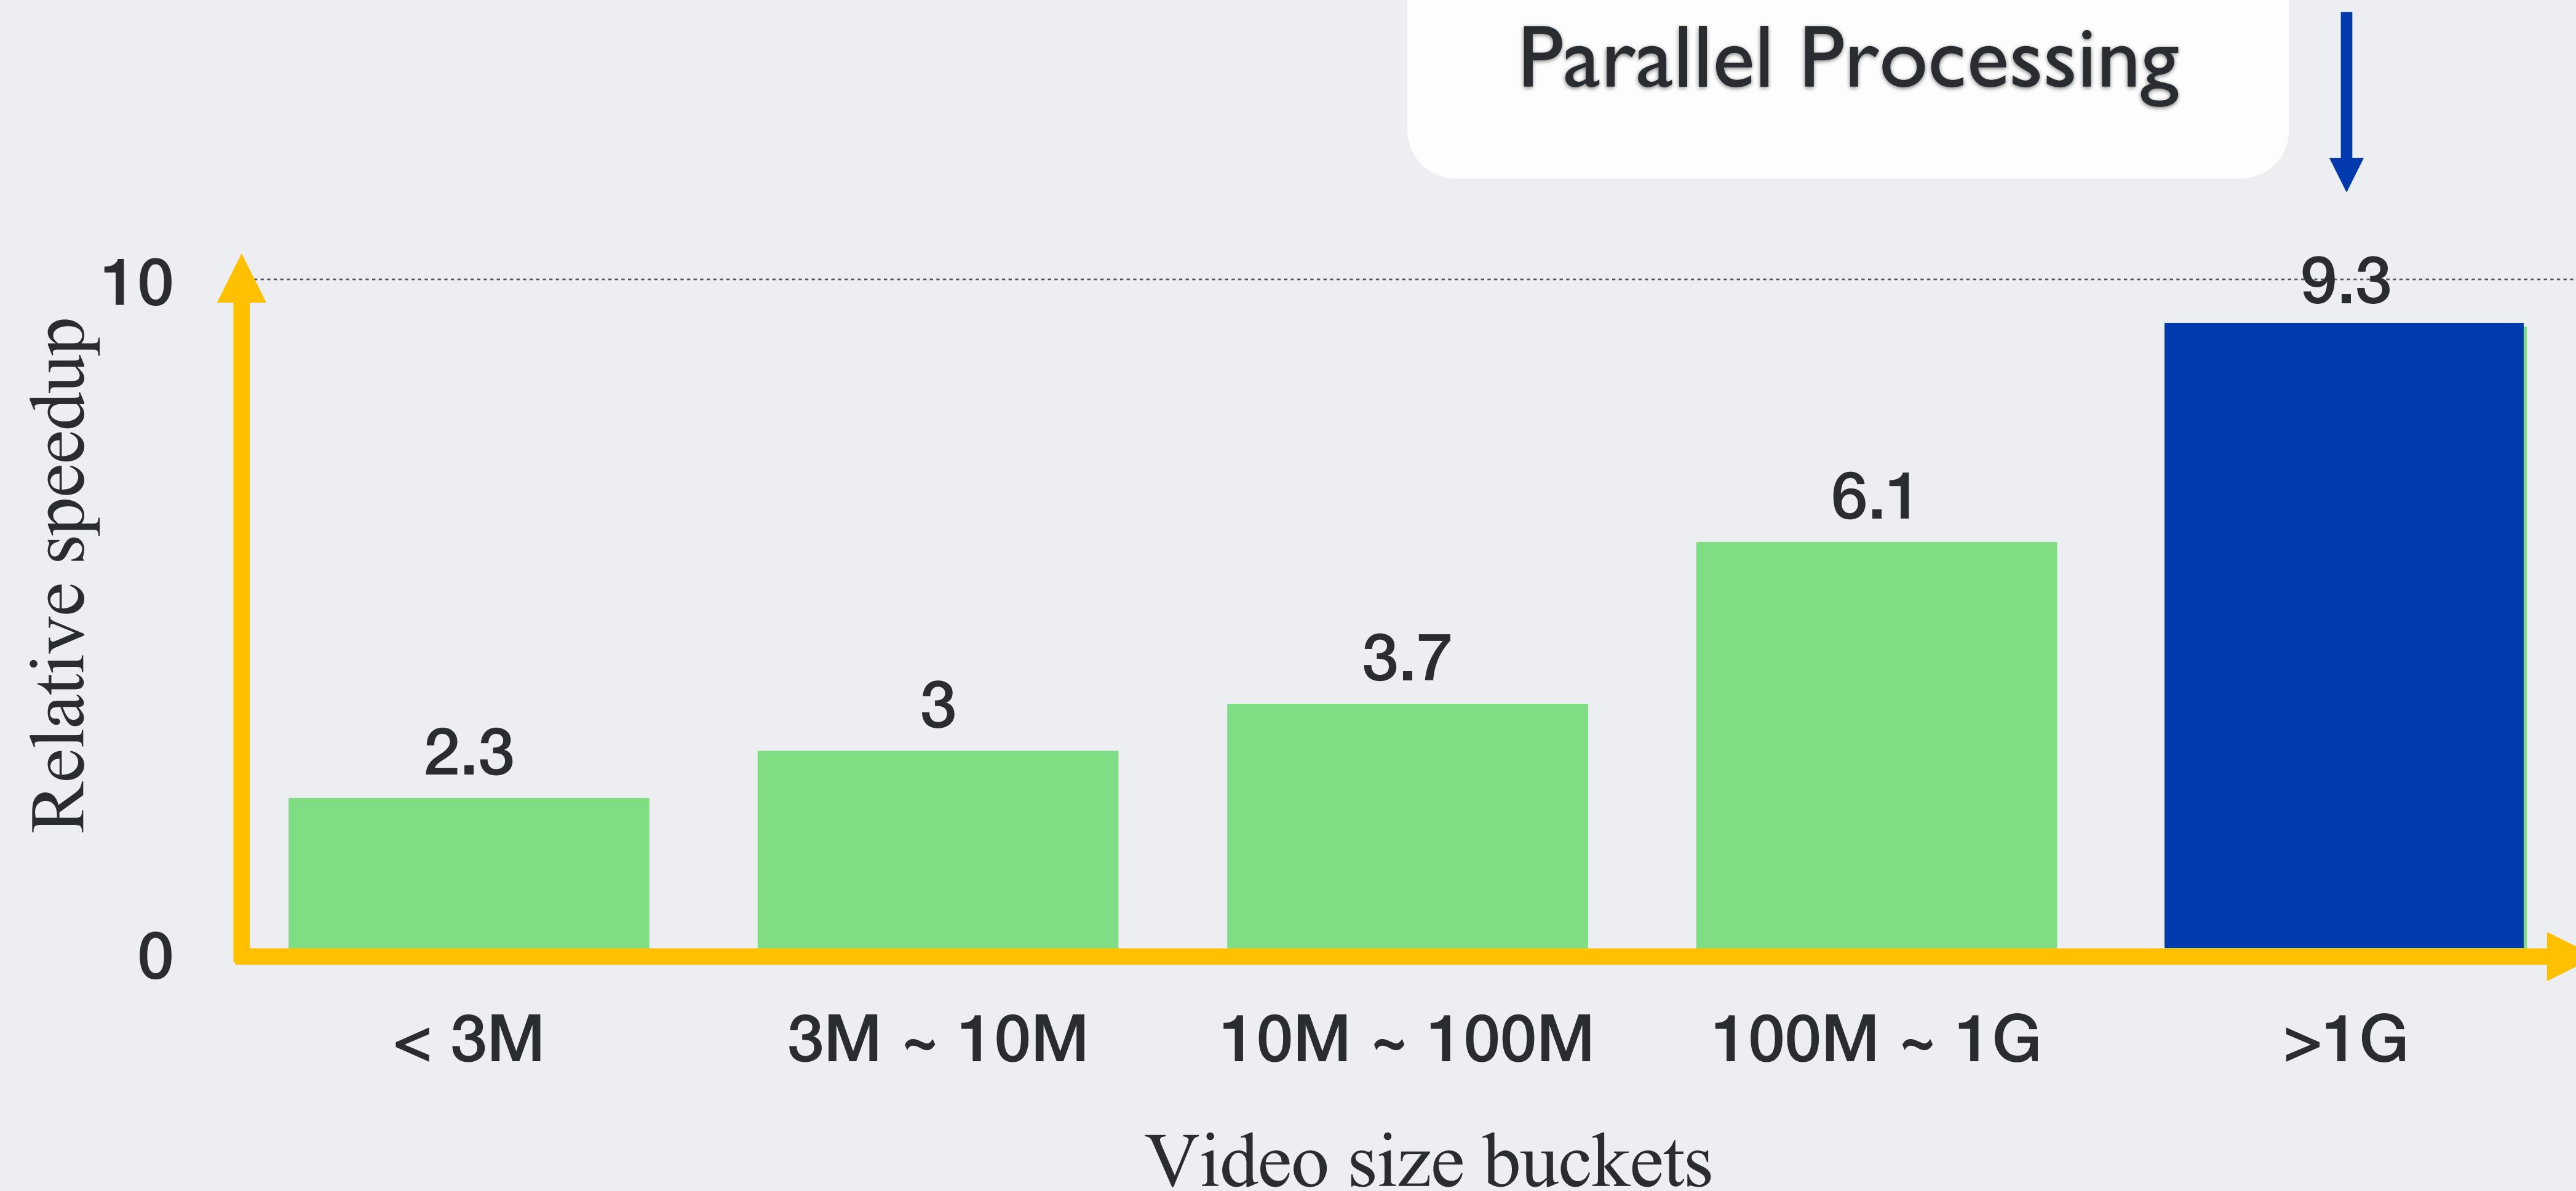
Results: 2.3x ~ 9.3x speedup



Results: 2.3x ~ 9.3x speedup



Results: 2.3x ~ 9.3x speedup



Challenges for video processing @ FB

Speedy

Users 2.3x share video speedup quickly

Flexible

Thousands of engineers can write pipelines for tens of apps

Robust

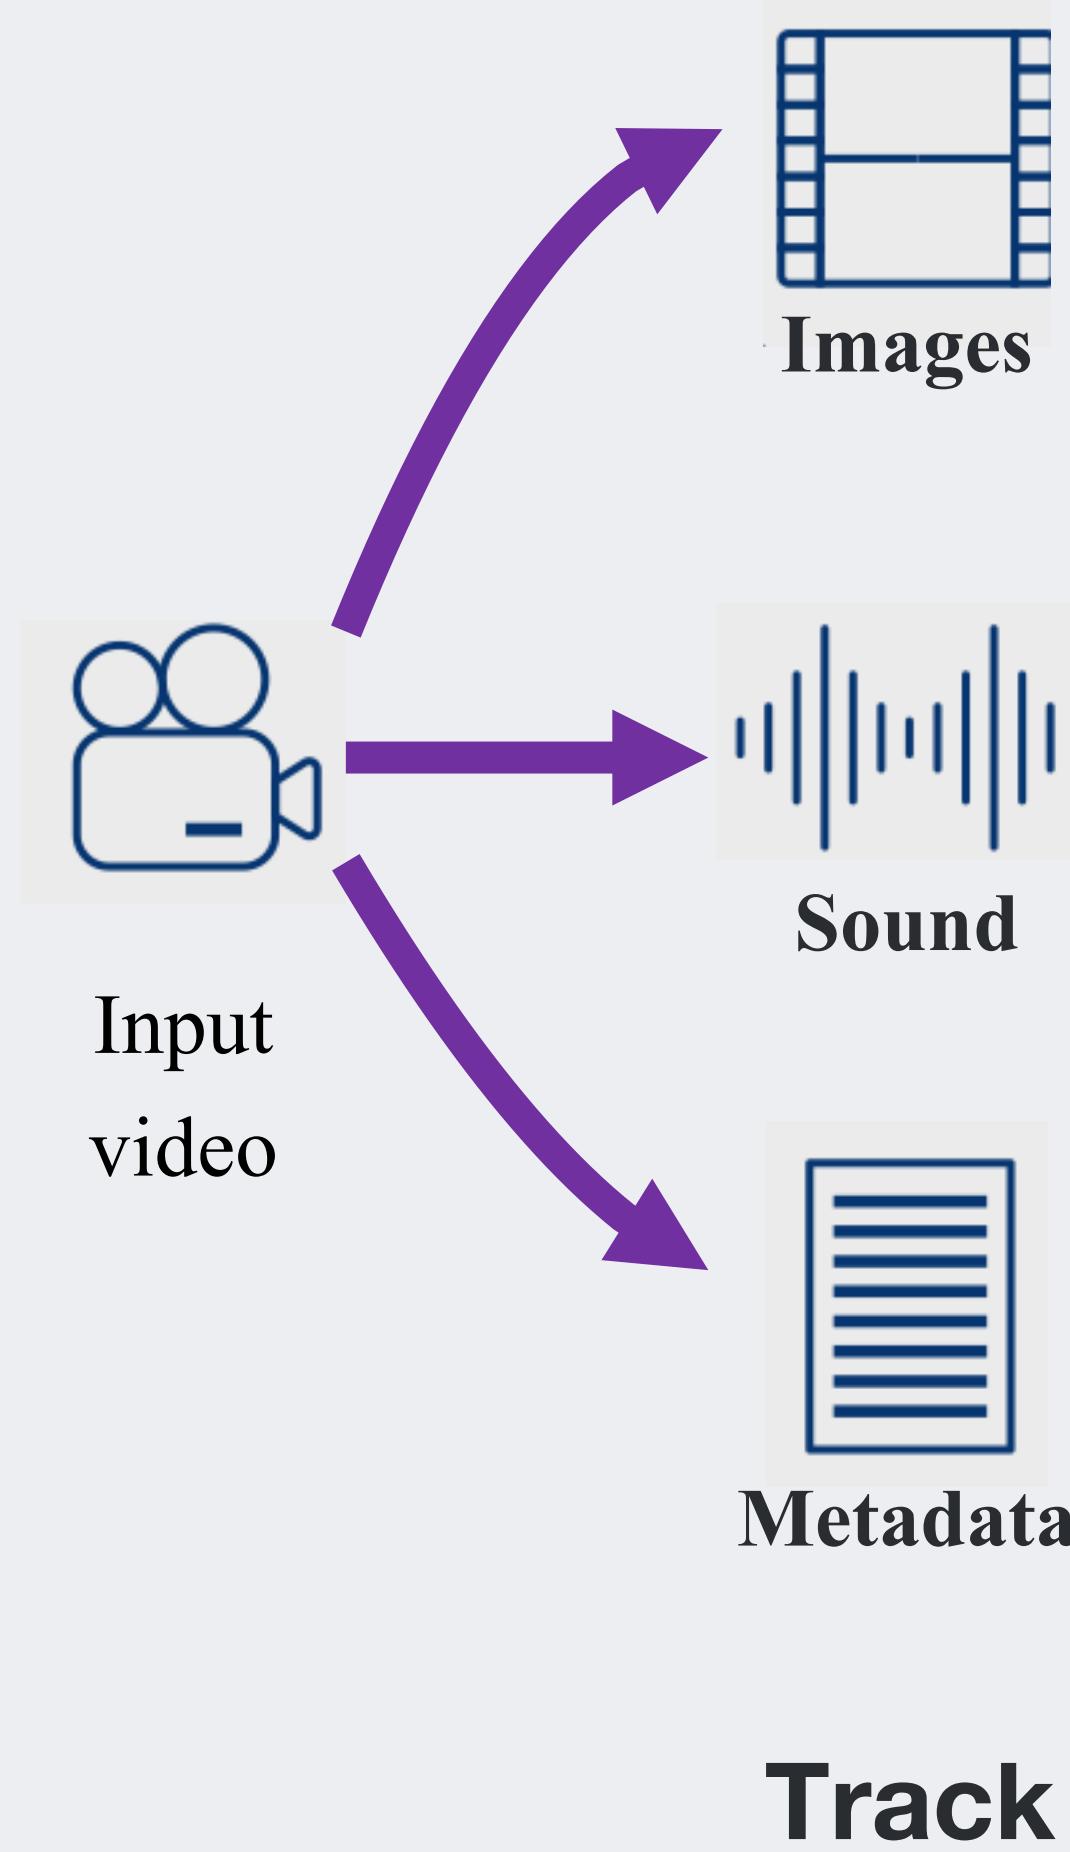
Handle faults and overload that is inevitable at scale

Flexible: build DAG framework

Thousands of engineers can write pipelines for tens of apps

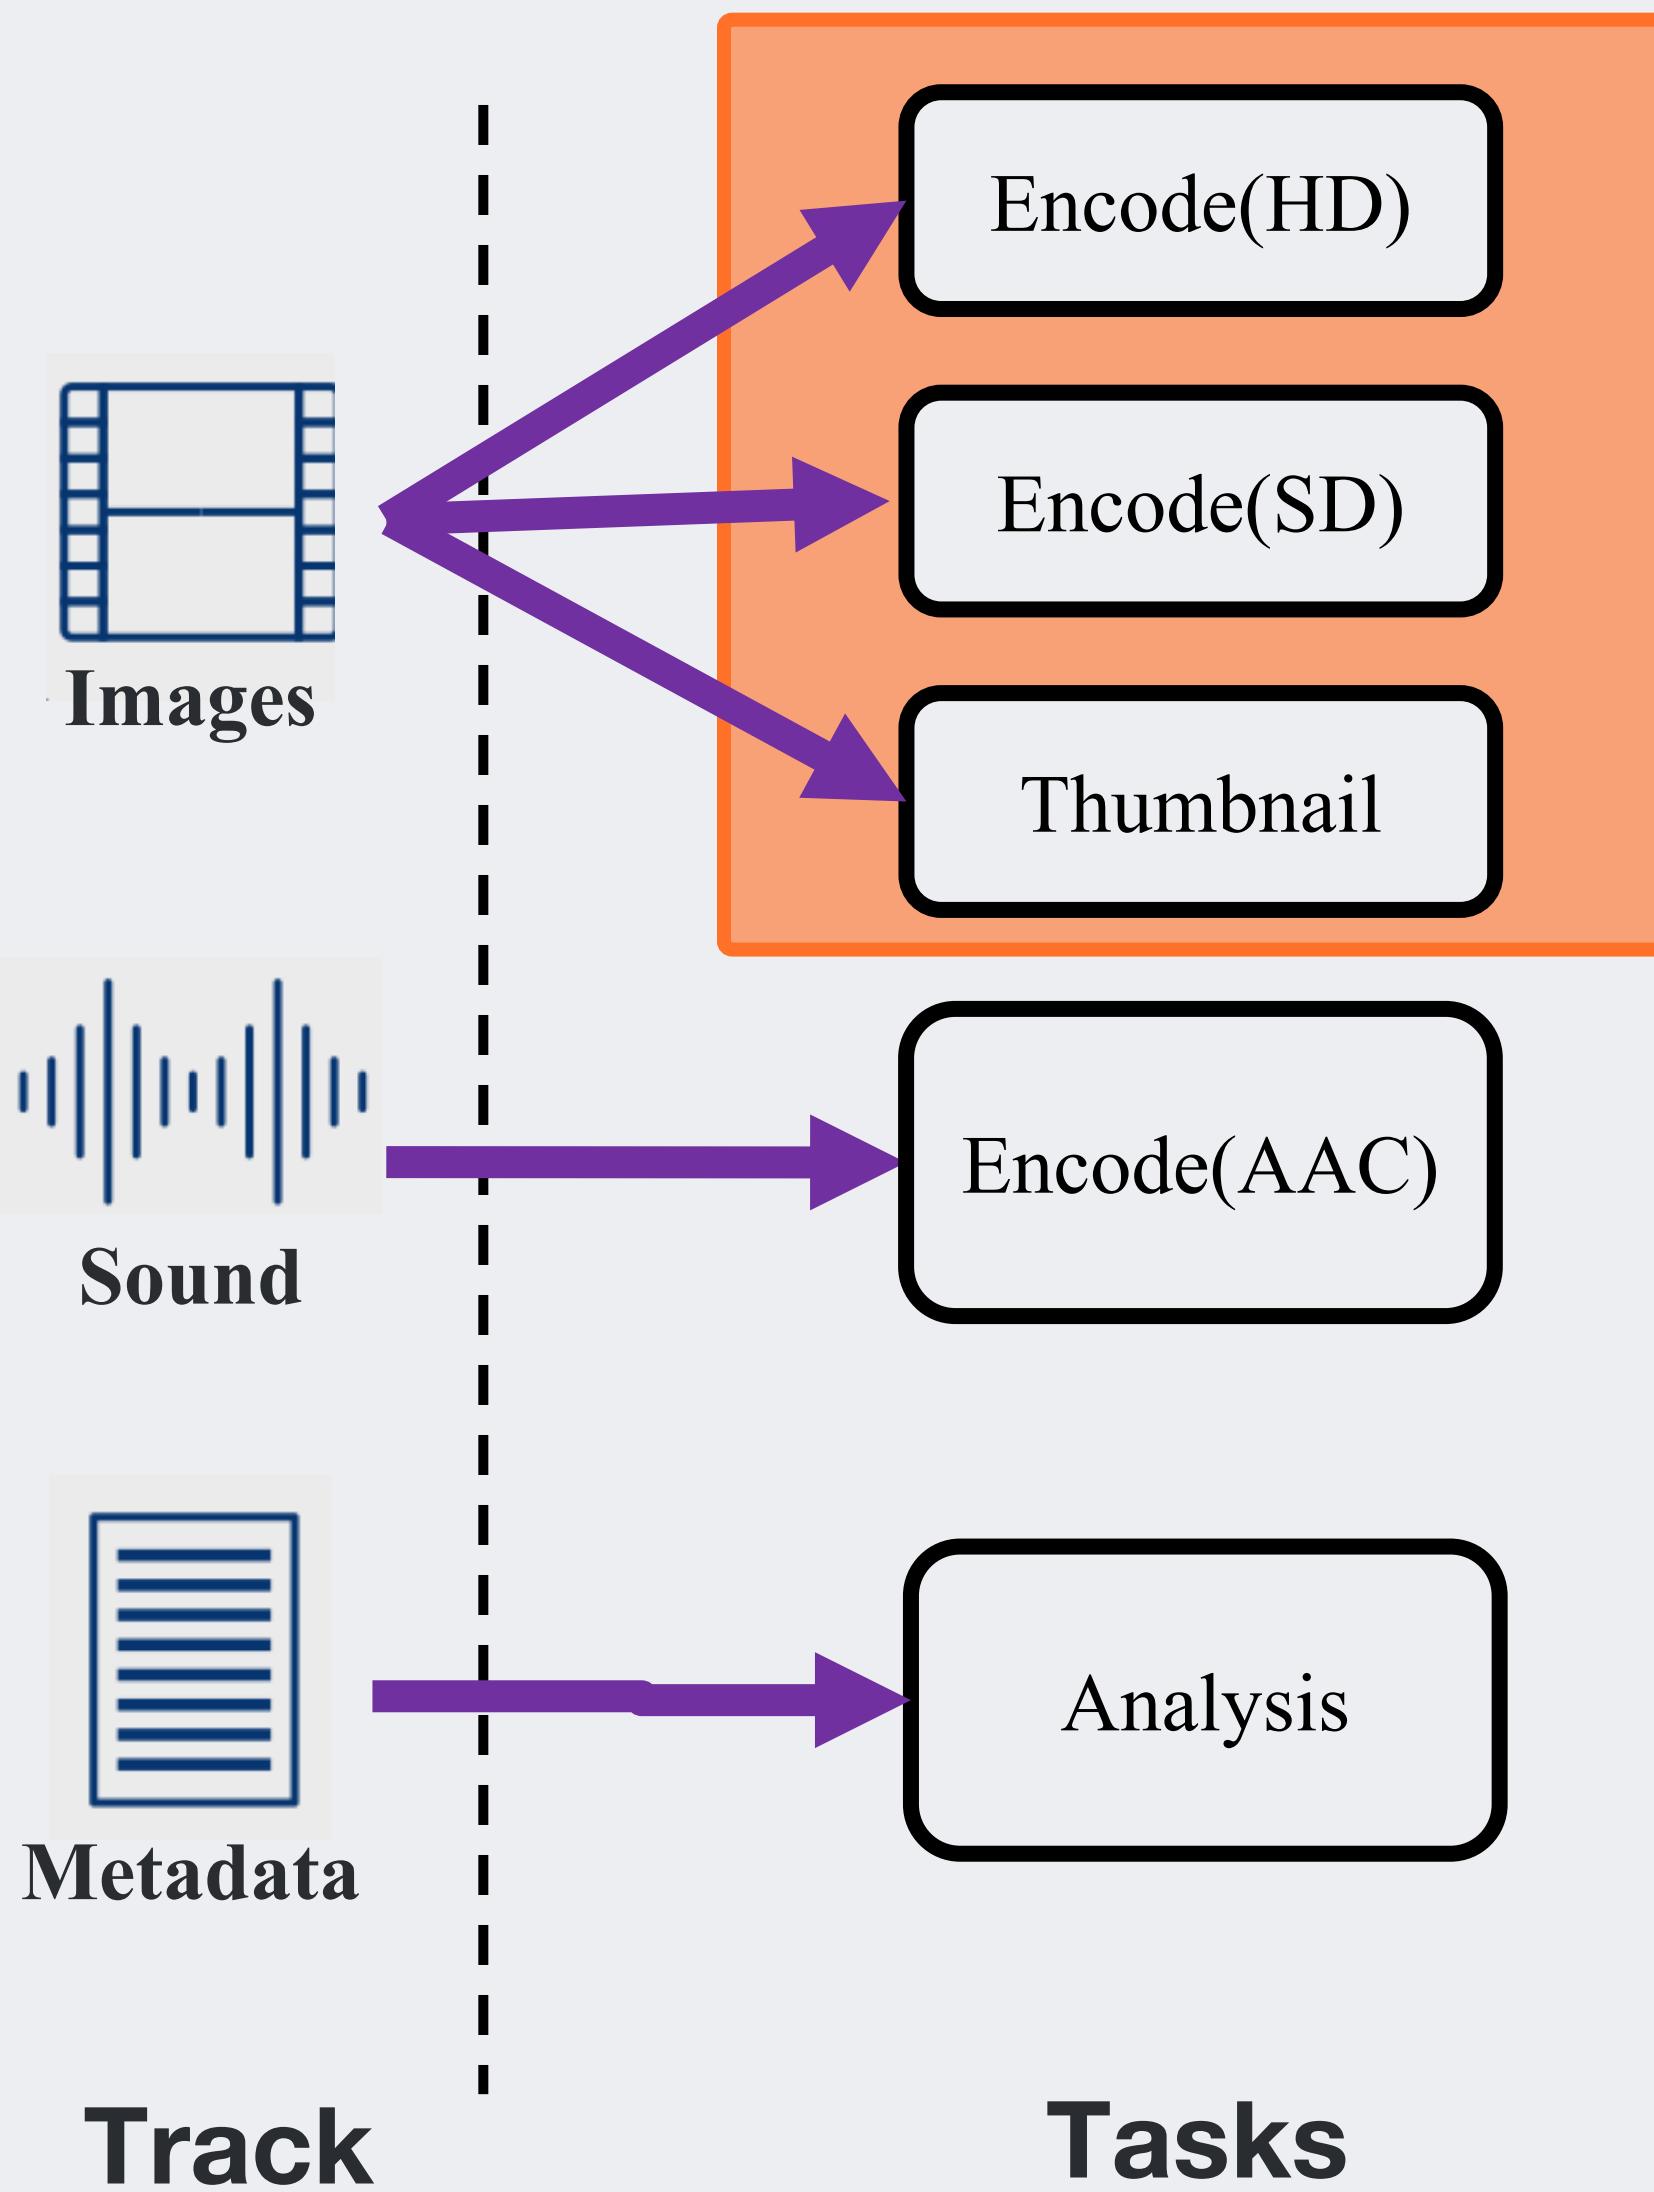
- DAG of computation on the stream-of-tracks abstraction
- Engineers write only sequential tasks in a familiar language
- Dynamic DAG generation per video

DAG on stream-of-tracks abstraction



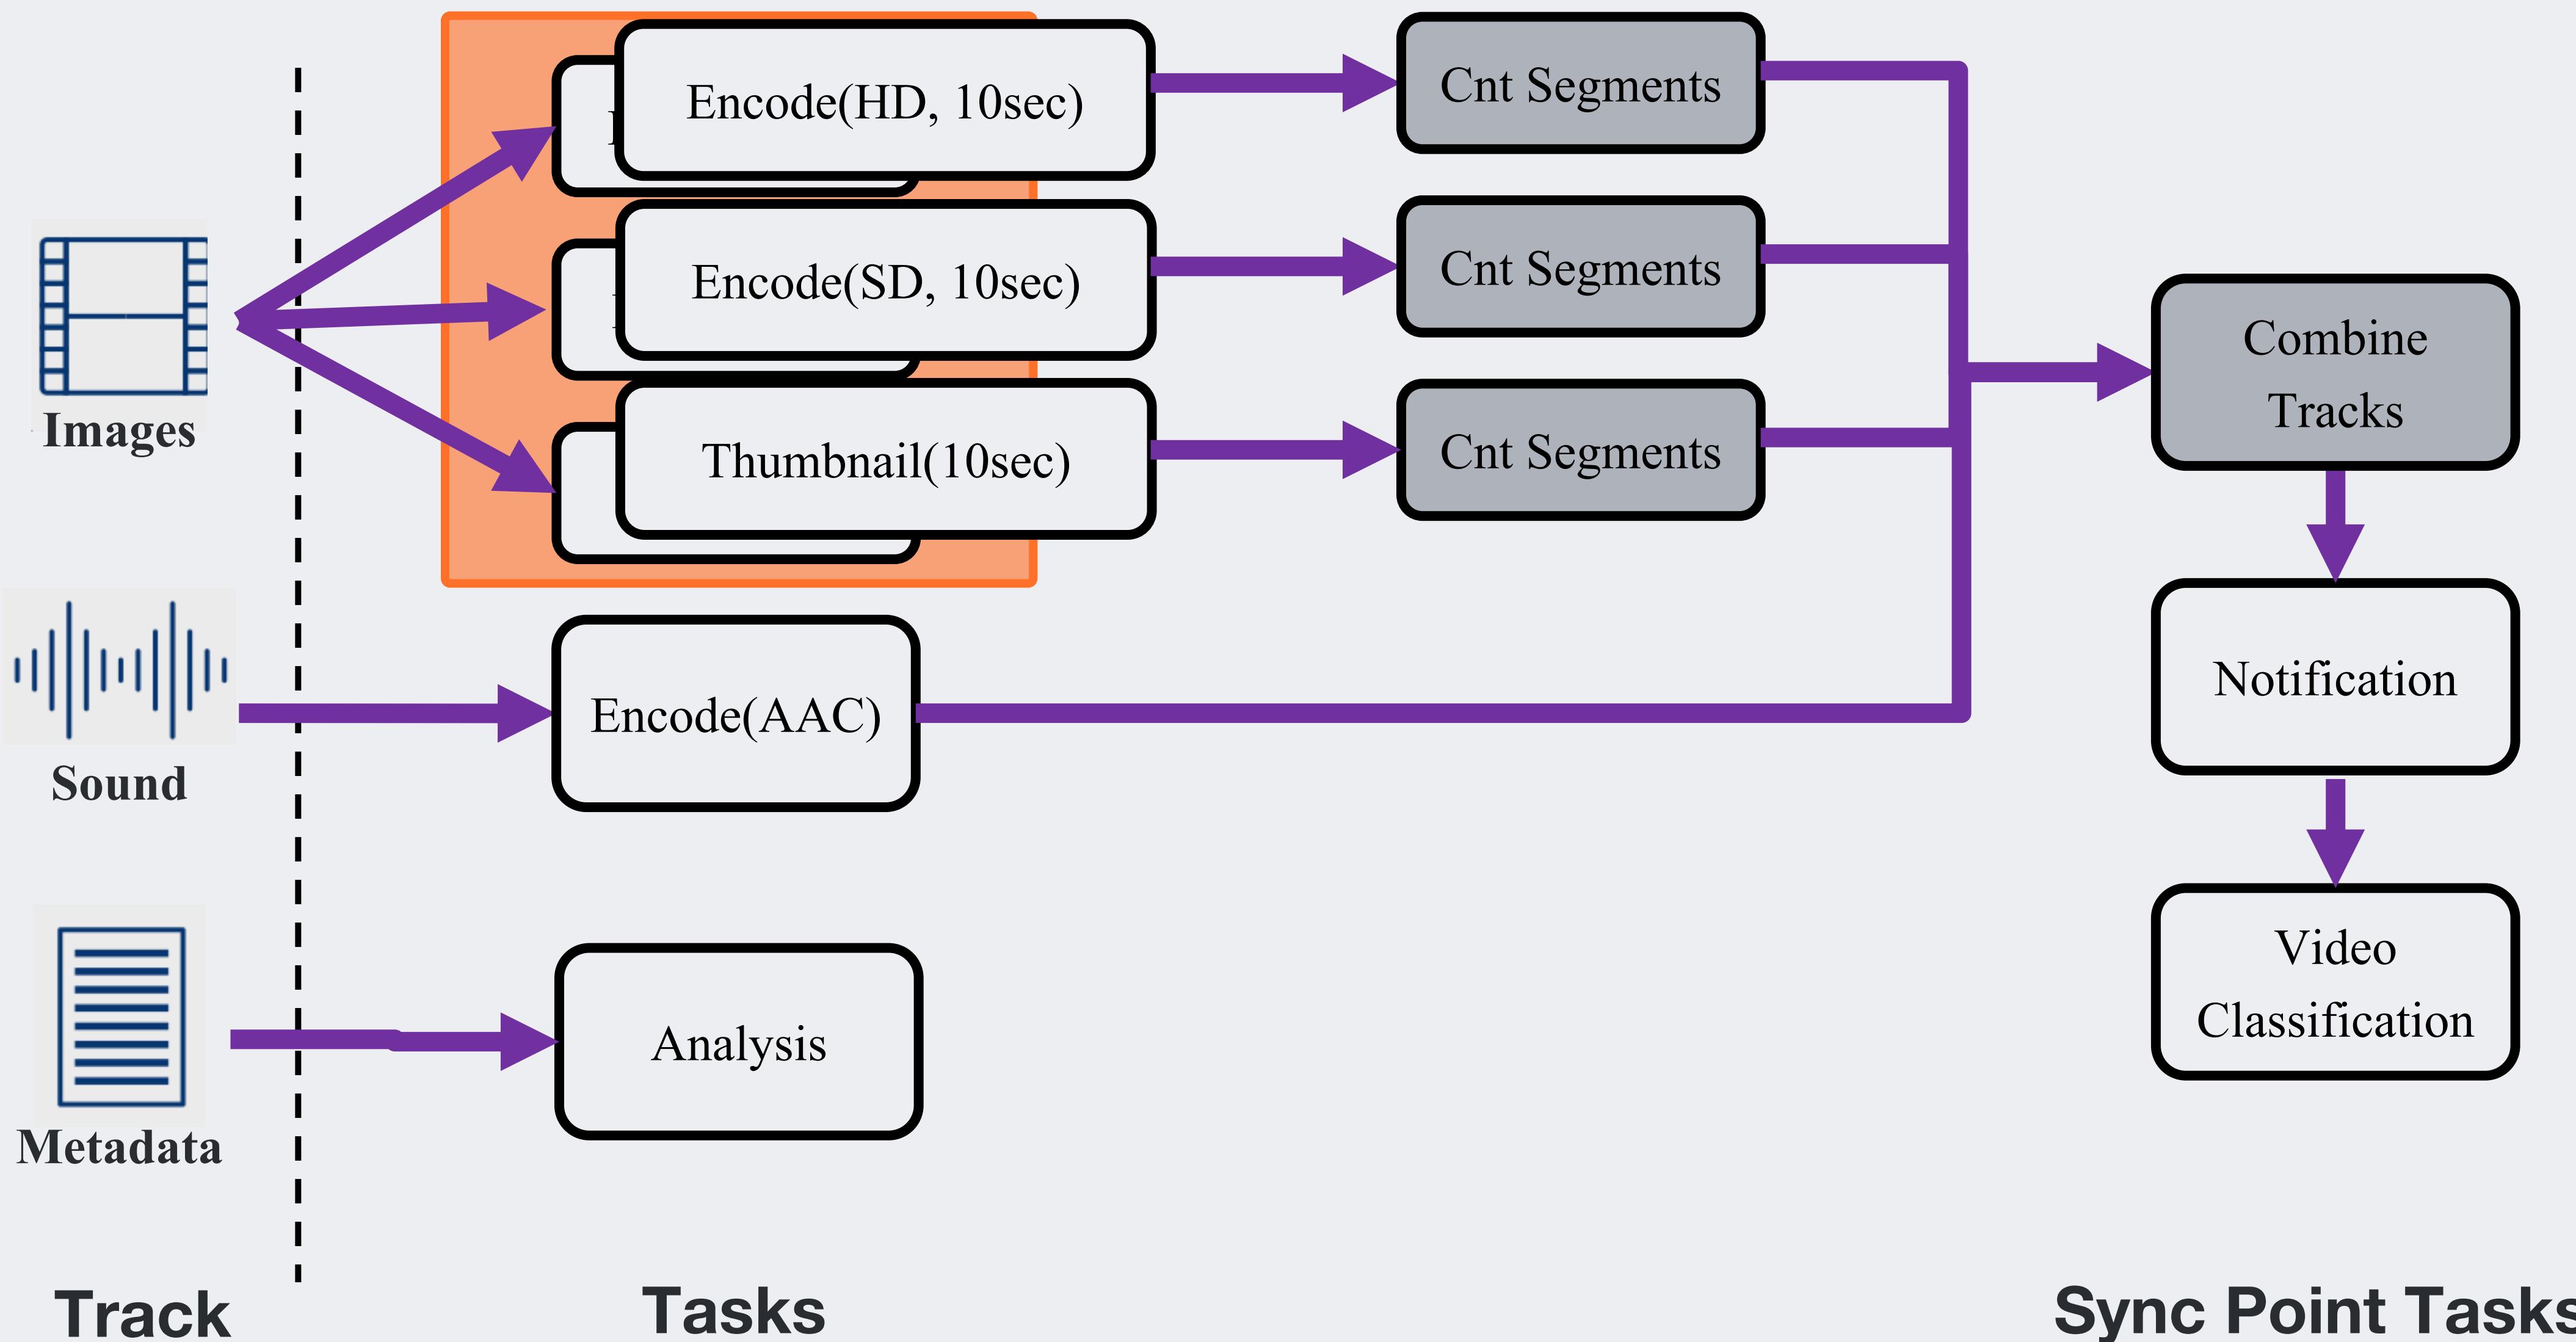
```
$pipeline = Pipeline.build()  
  
$video_track=$pipeline>addTrack(IMG_TYPE)  
->addTask()  
  
$audio_track=$pipeline>addTrack(AUD_TYPE)  
->addTask()  
  
$meta_track=$pipeline>addTrack(META_TYPE)  
->addTask()
```

DAG on stream-of-tracks abstraction

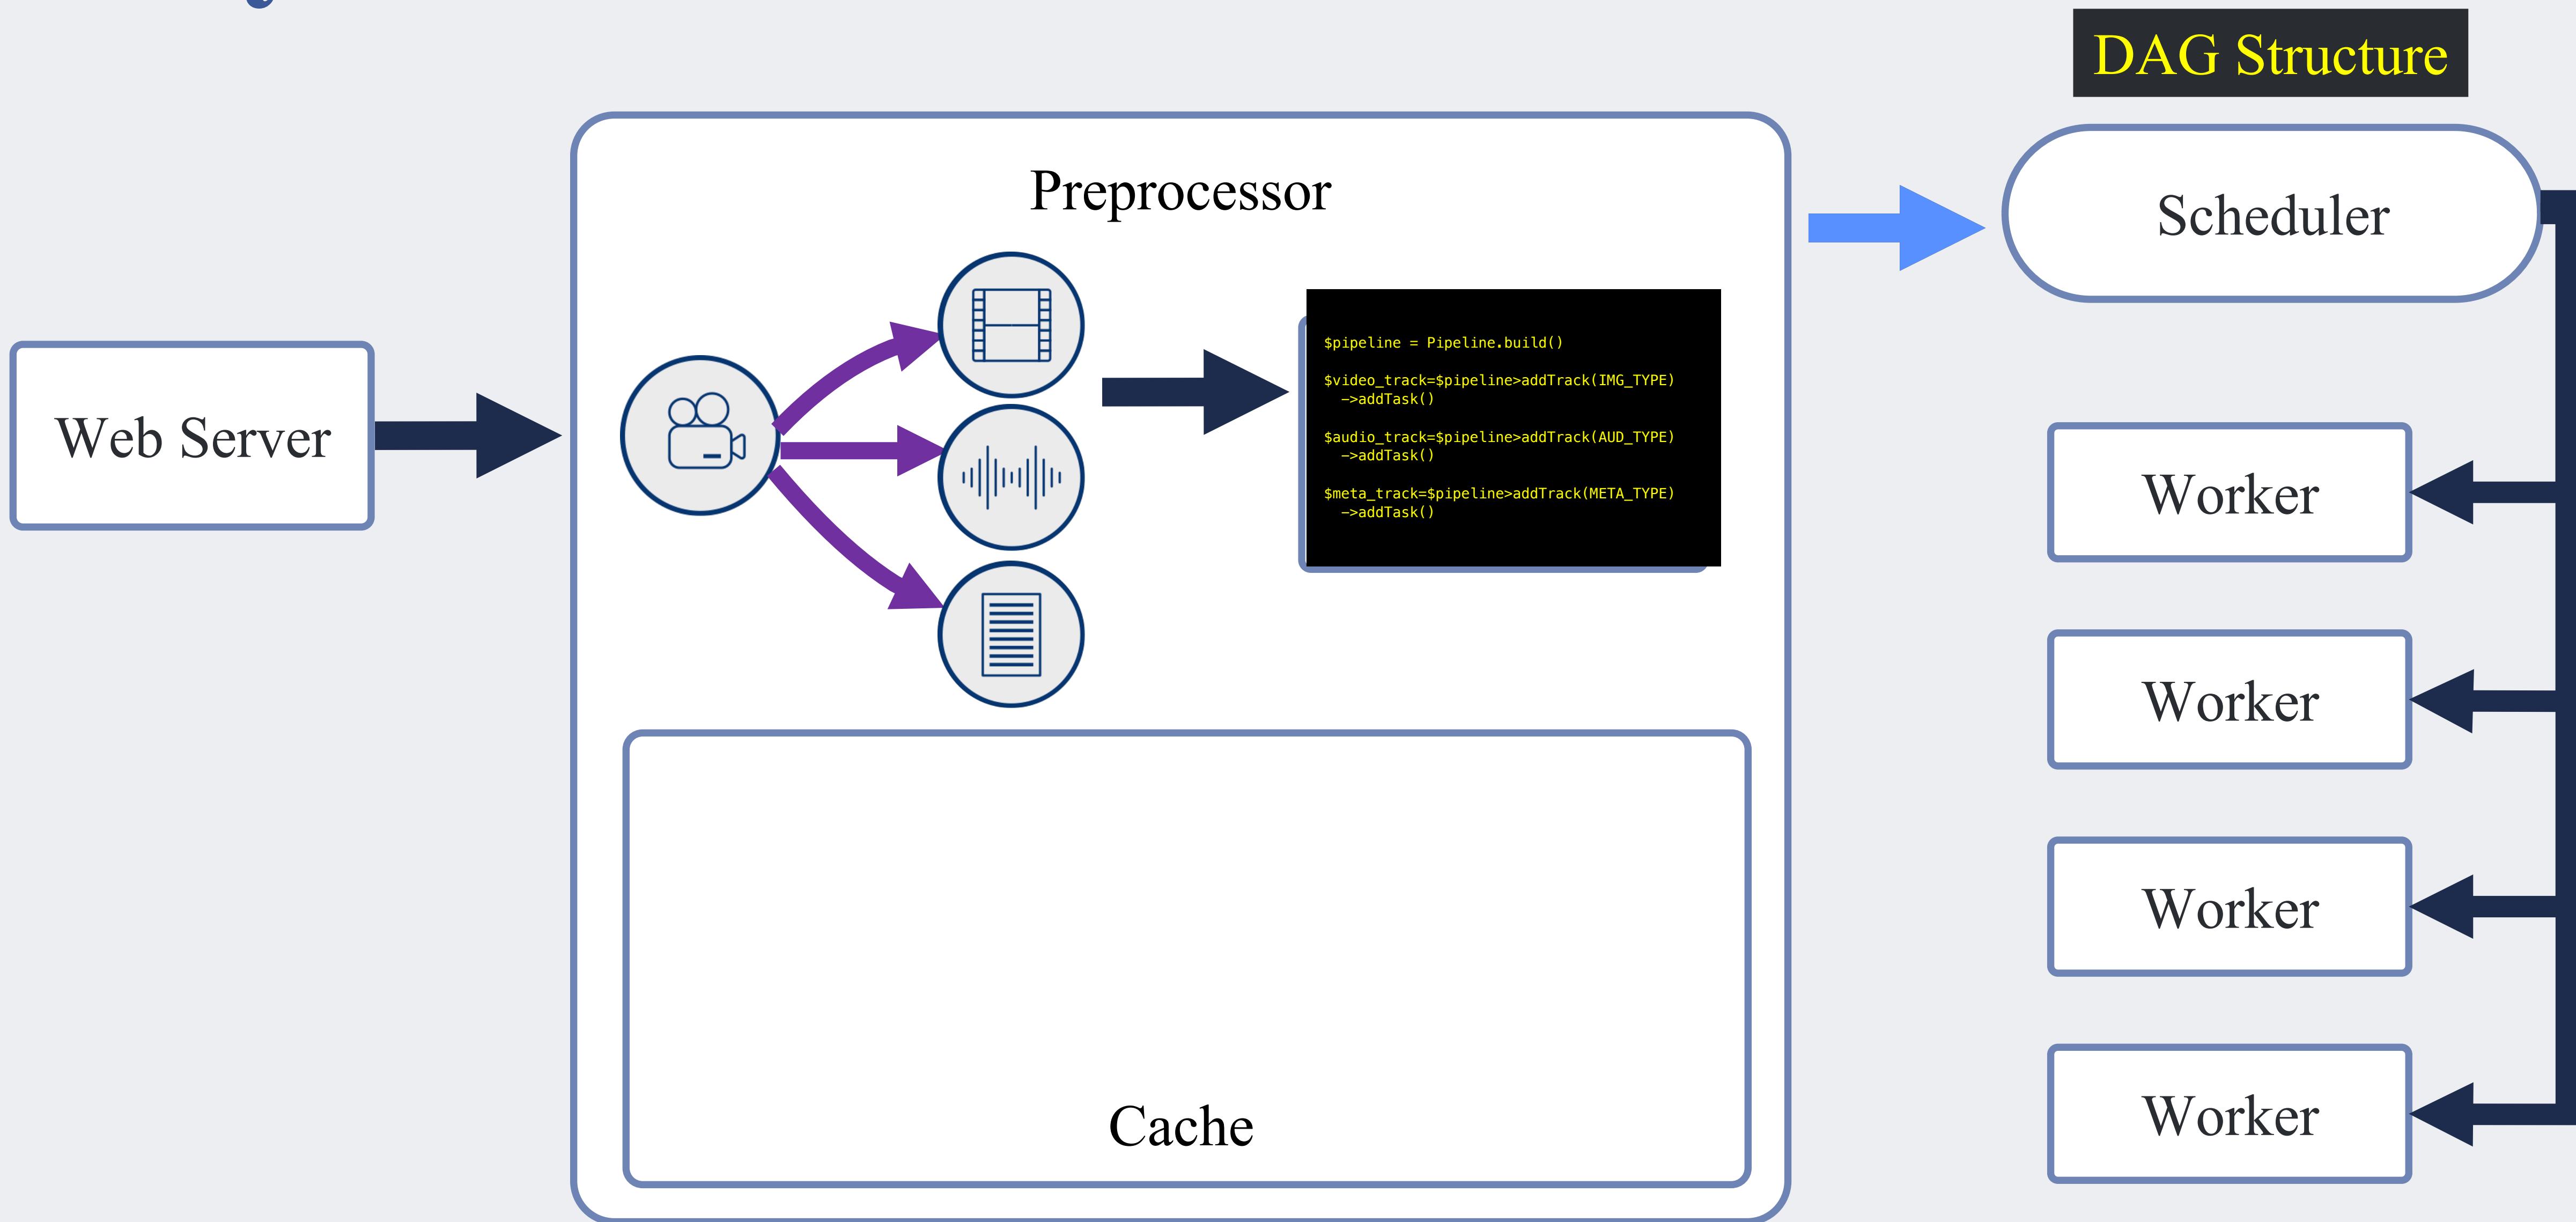


```
$pipeline = Pipeline.build()  
  
$video_track=$pipeline>>addTrack(IMG_TYPE)  
    ->addTask(Encode(HD), 10s), Encode(SD), Thumbnail  
        Encode(SD, 10s), Thumbnail(10s))  
  
$audio_track=$pipeline>>addTrack(AUD_TYPE)  
    ->addTask(Encode(AAC))  
  
$meta_track=$pipeline>>addTrack(META_TYPE)  
    ->addTask(Analysis)
```

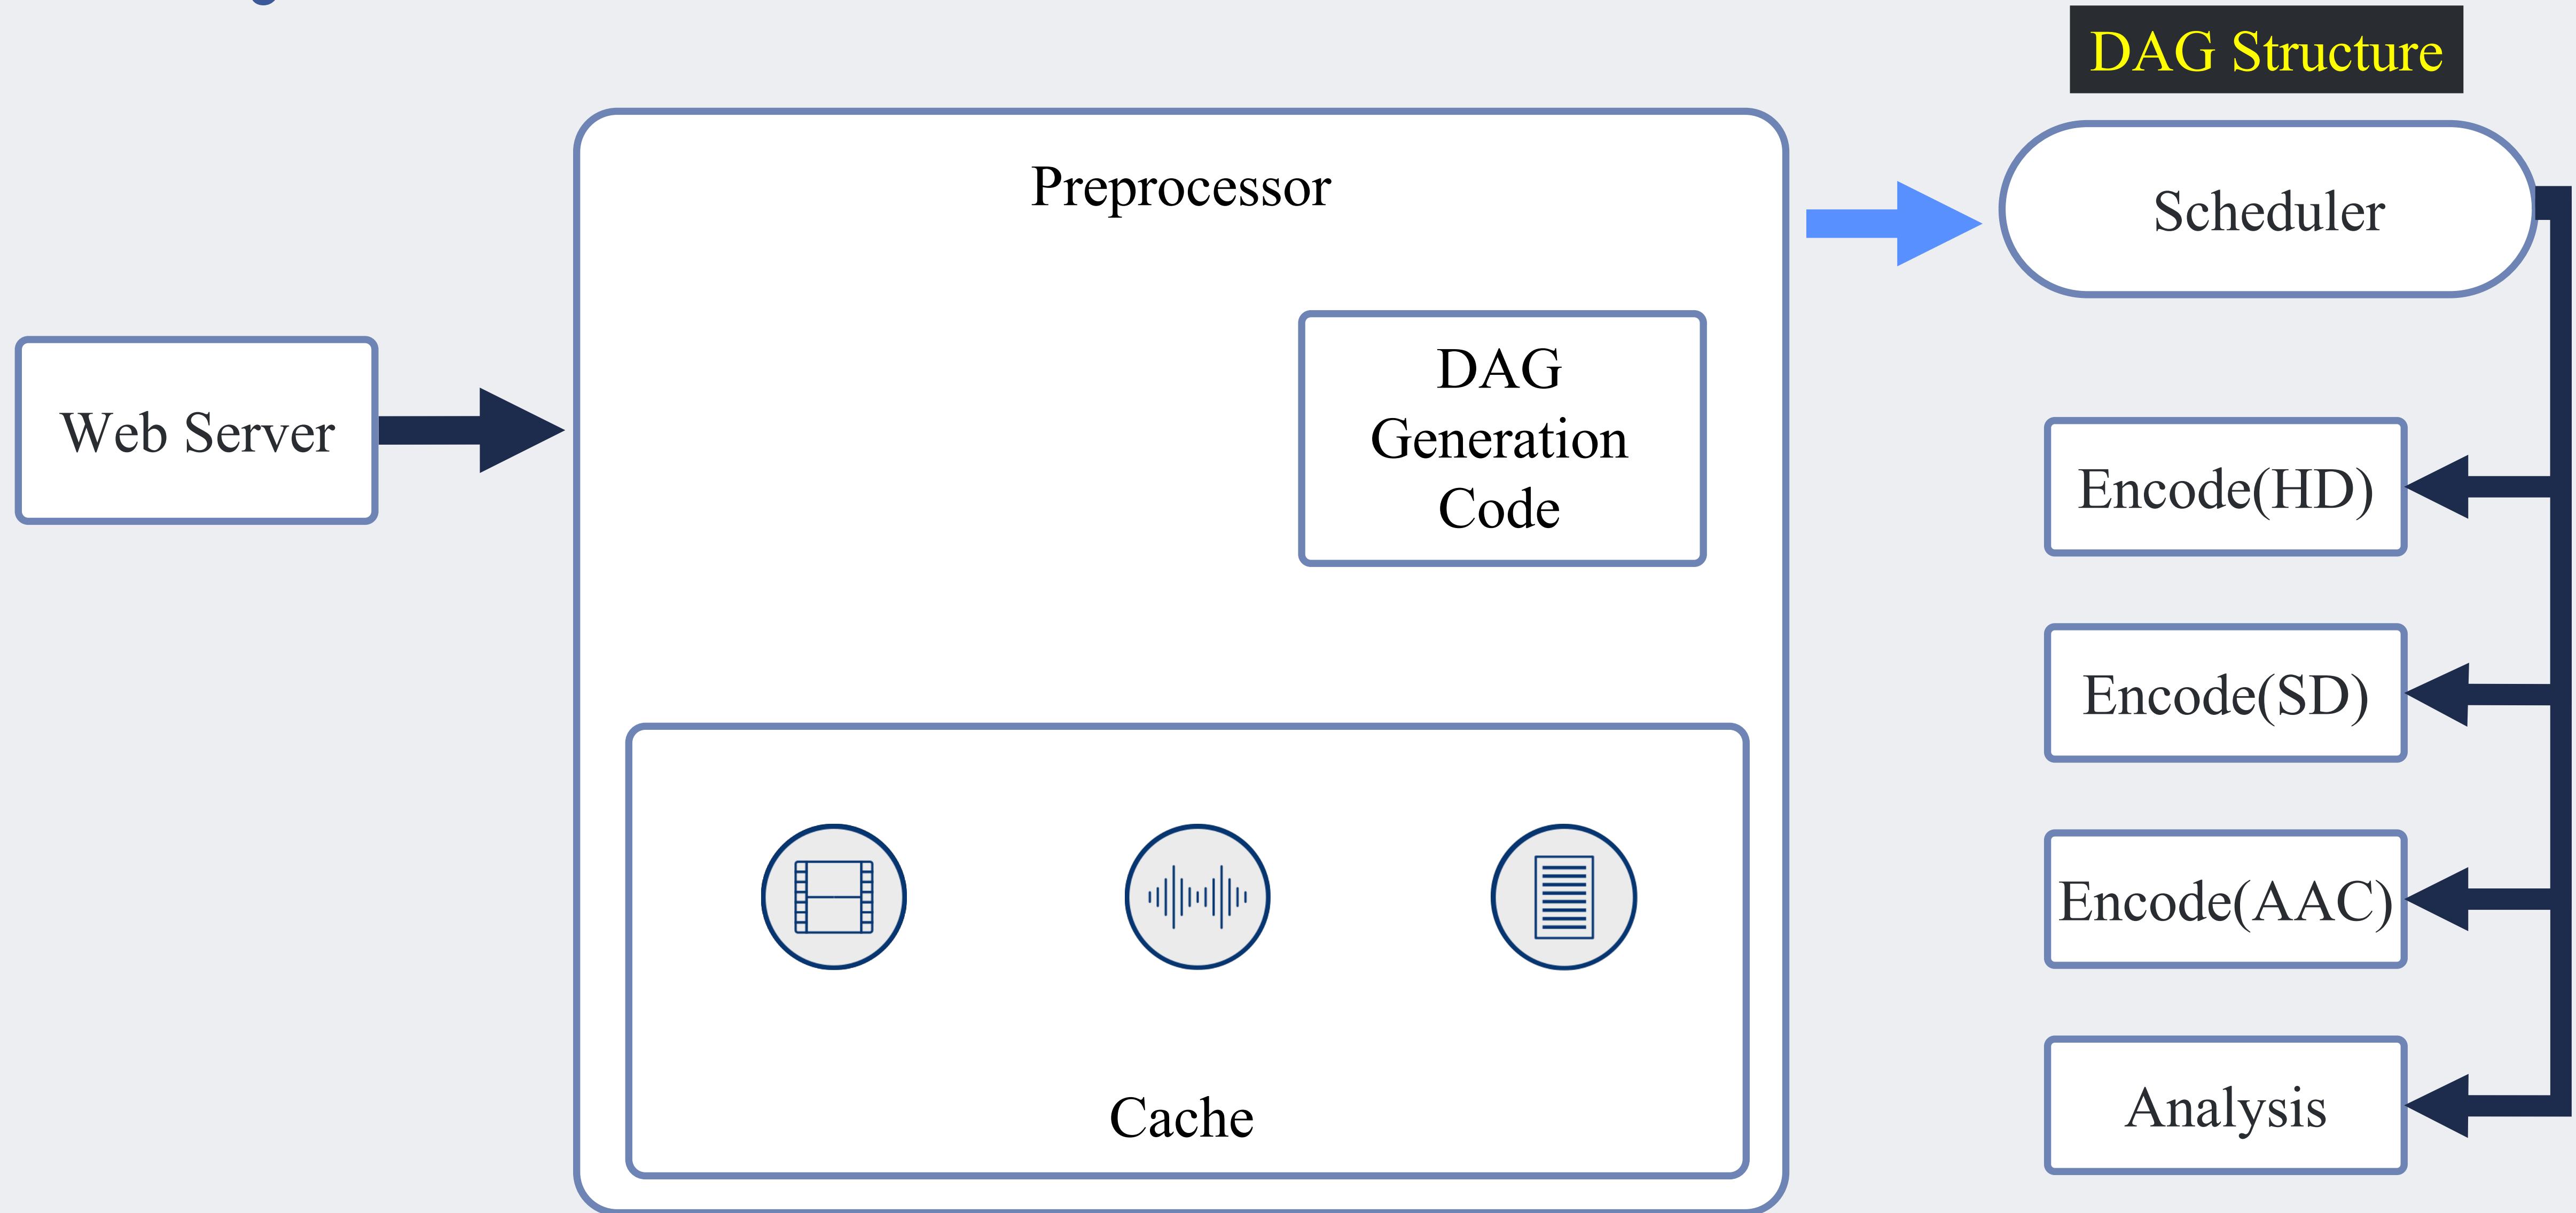
DAG on stream-of-tracks interface



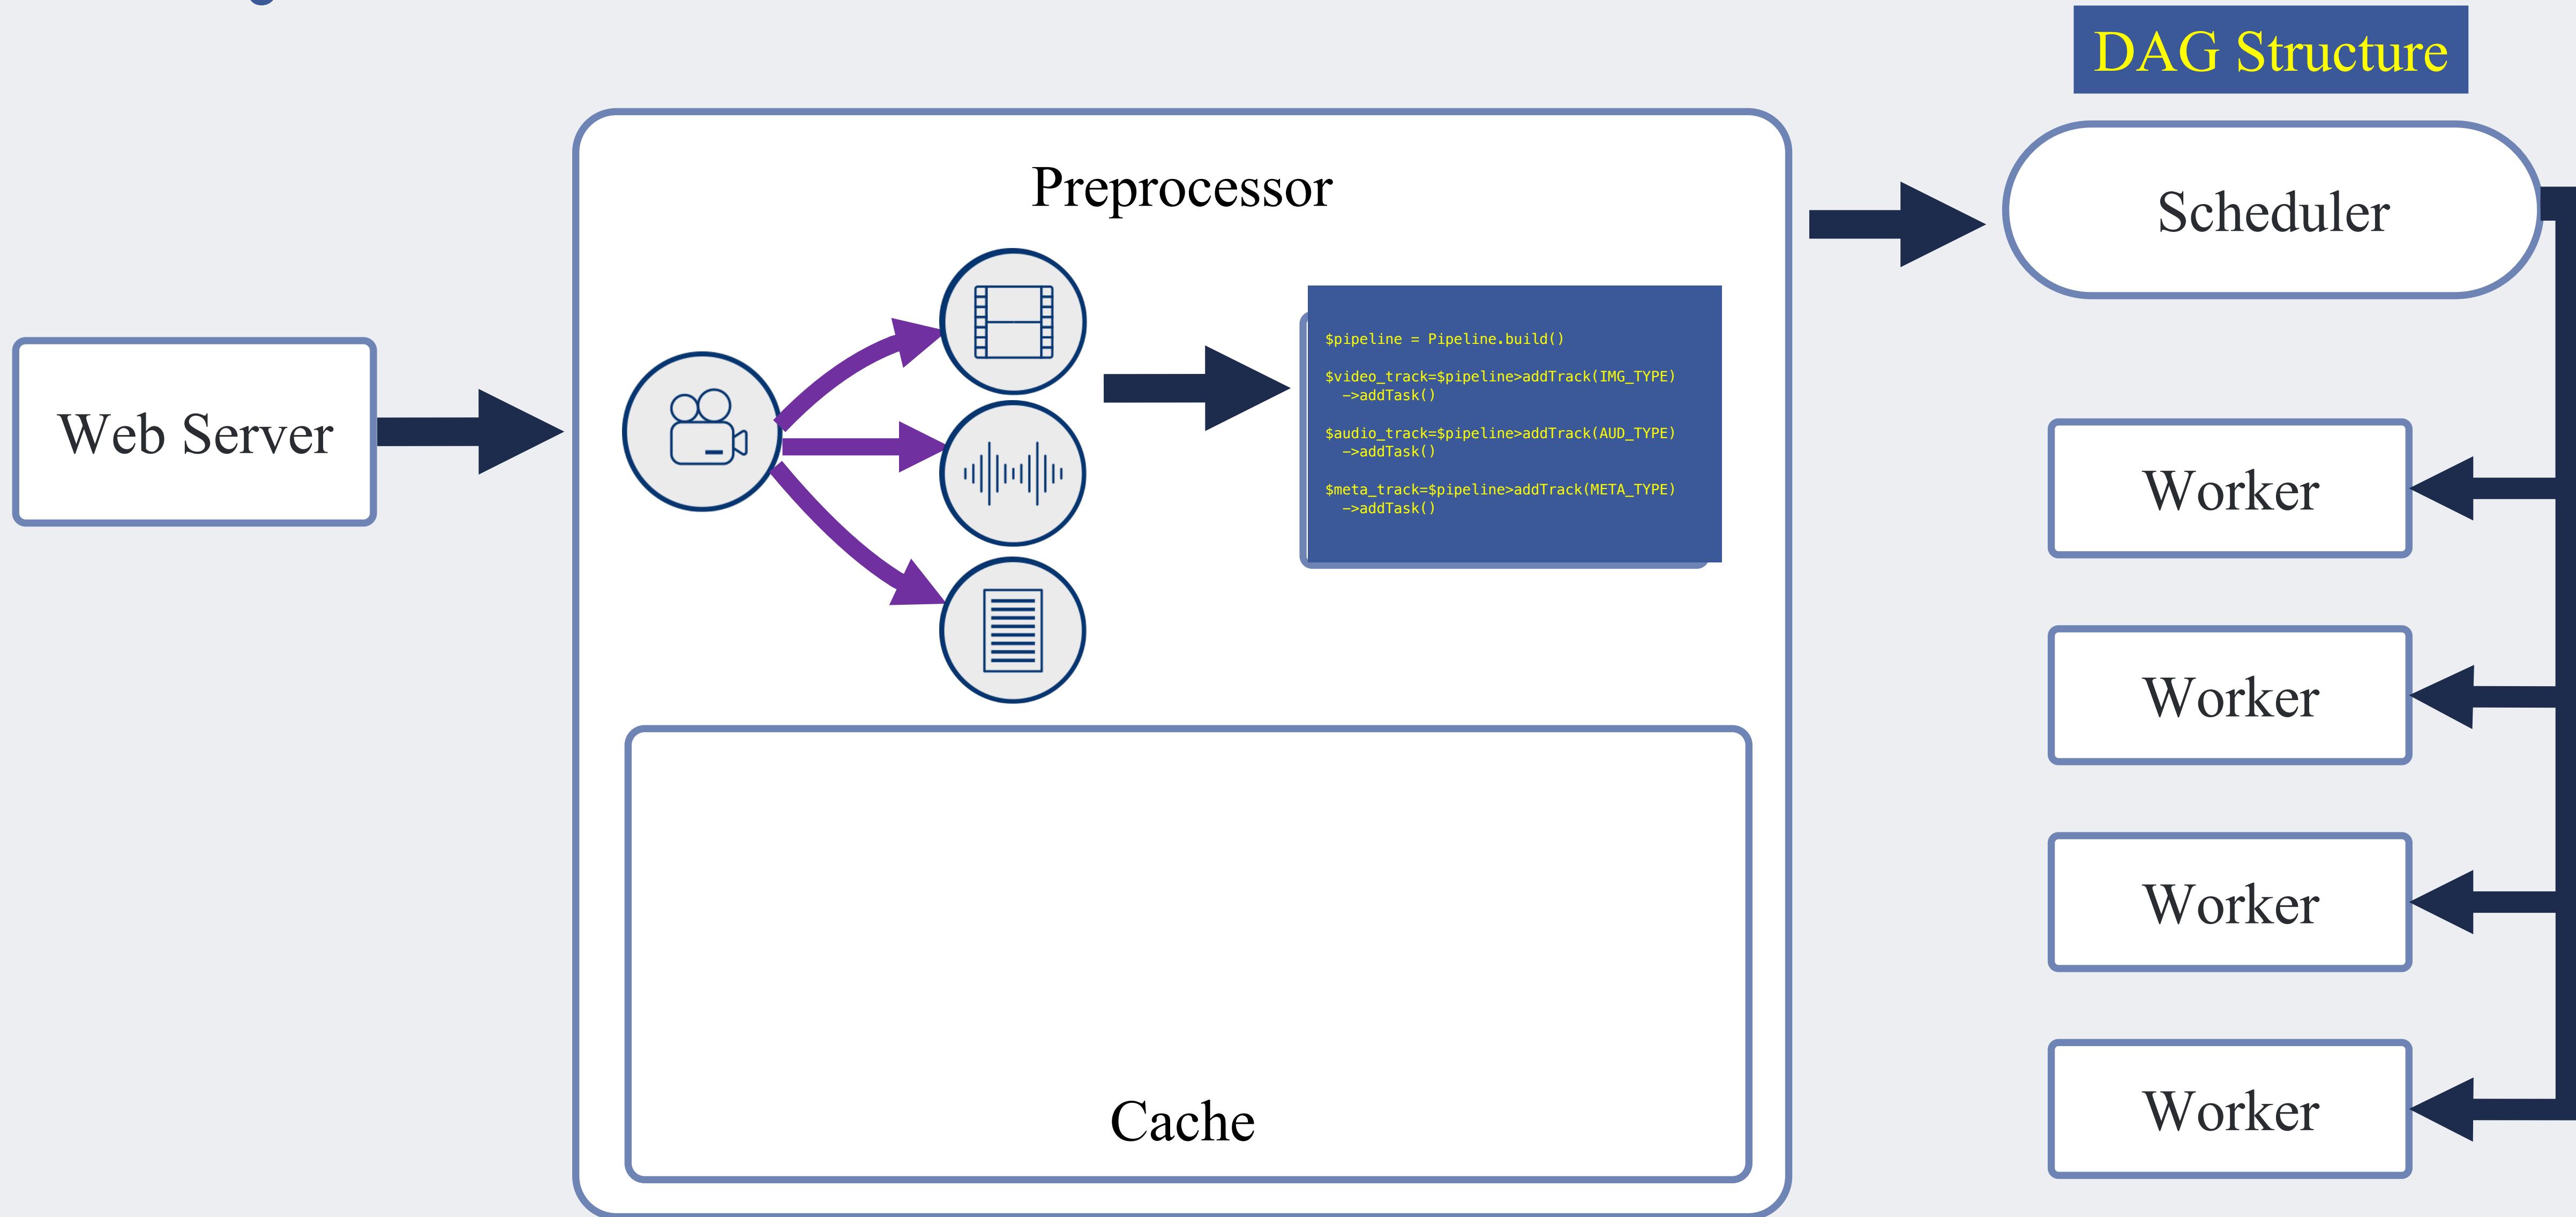
Dynamic DAG Generation



Dynamic DAG Generation



Dynamic DAG Generation



One system for 15+ applications

- Generate **billions** of tasks per day
- Varying DAG size
 - 360 video has **thousands** of tasks per upload
 - Newsfeed post averages at **153** tasks per upload
 - Instagram averages at **22** tasks per upload
 - Messenger averages at **18** tasks per upload

Challenges for video processing @ FB

Speedy

2.3x ~ 9.3x speedup

Flexible

Thousands of Engisystem for 15+ applications or tens of apps

Robust

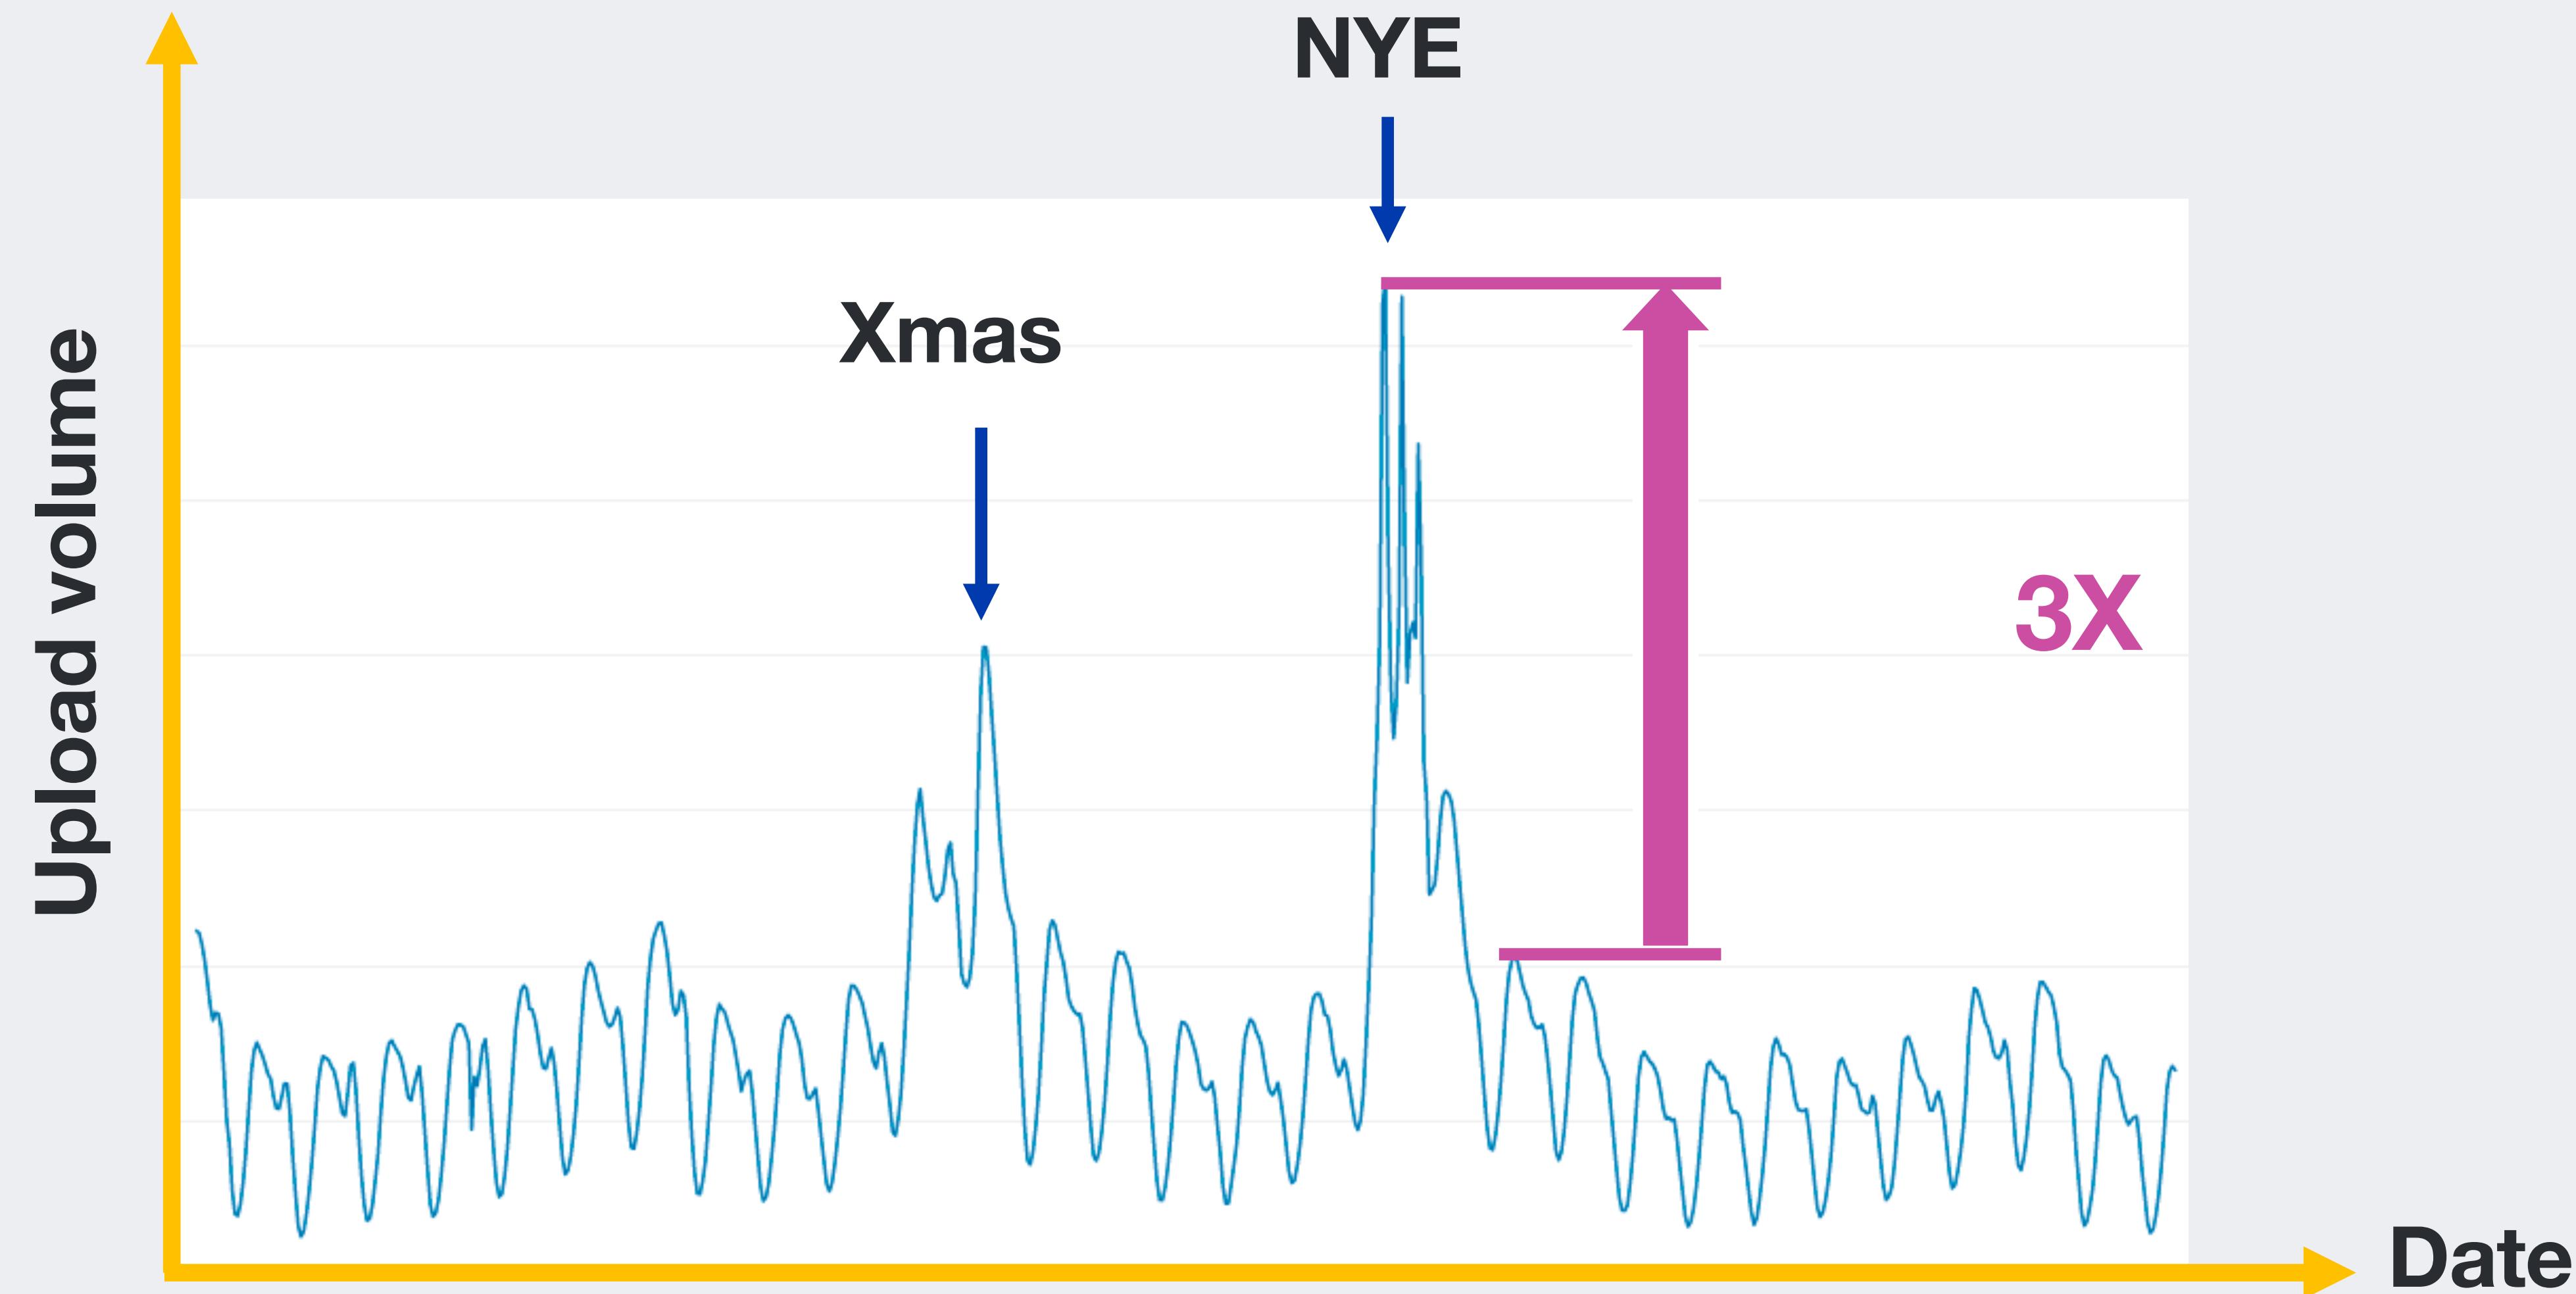
Handle faults and overload that is inevitable at scale

Robust: tolerate overload

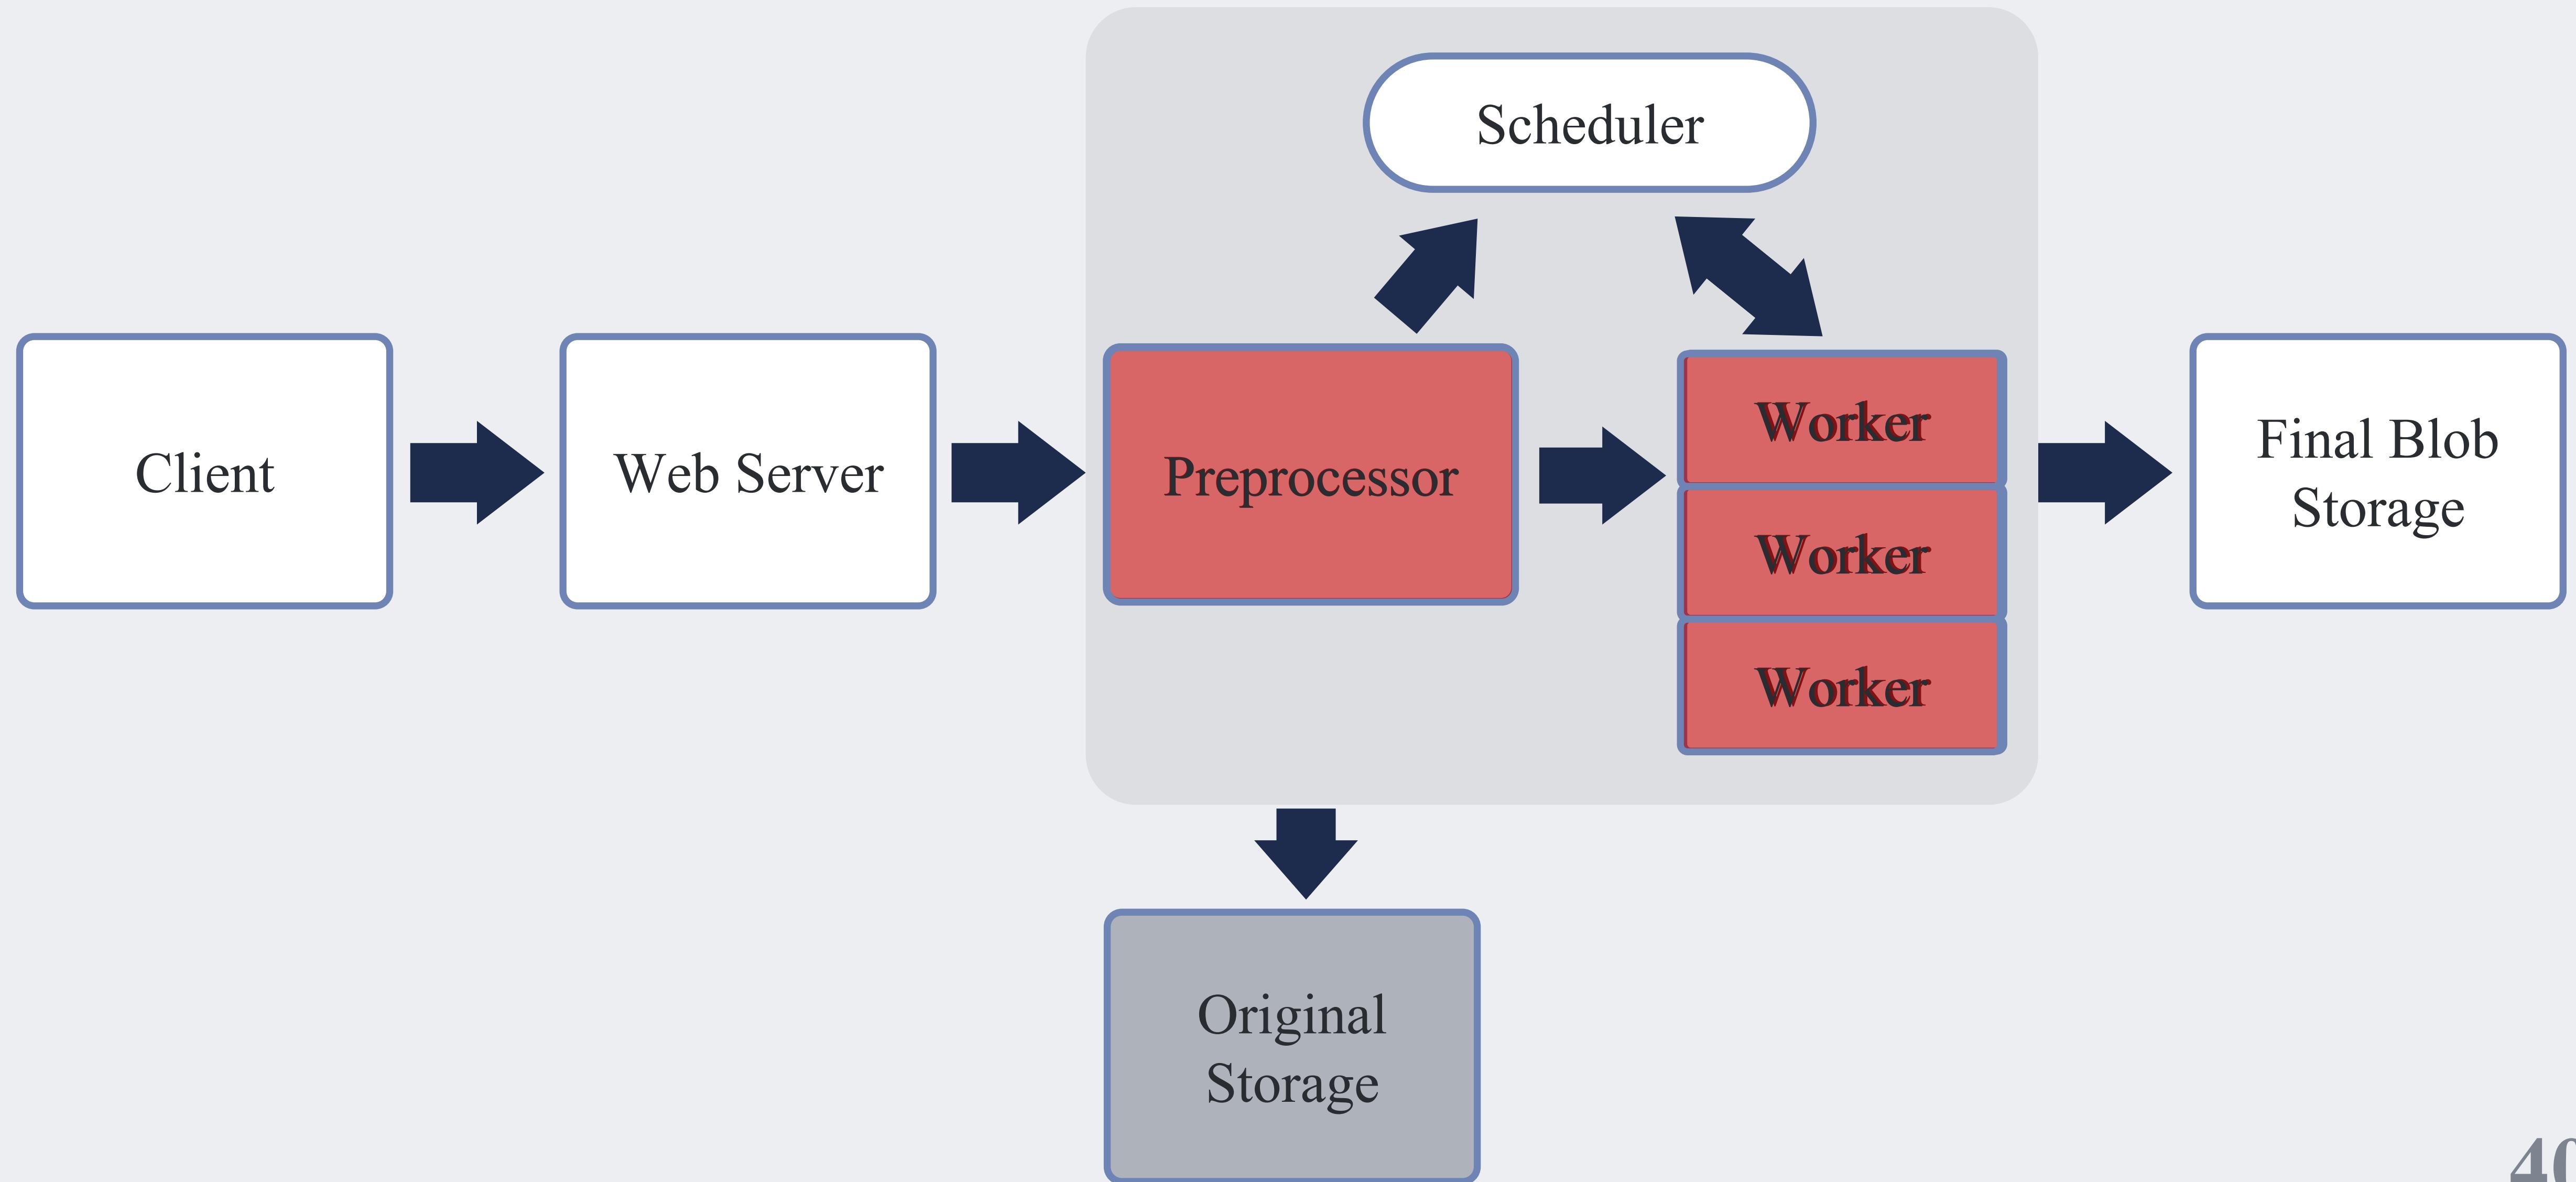
Handle faults and overload that is inevitable at scale

- Rely on priority to degrade non-latency-sensitive tasks
- Defer full video processing for some new uploads
- Load-shedding across global deployments

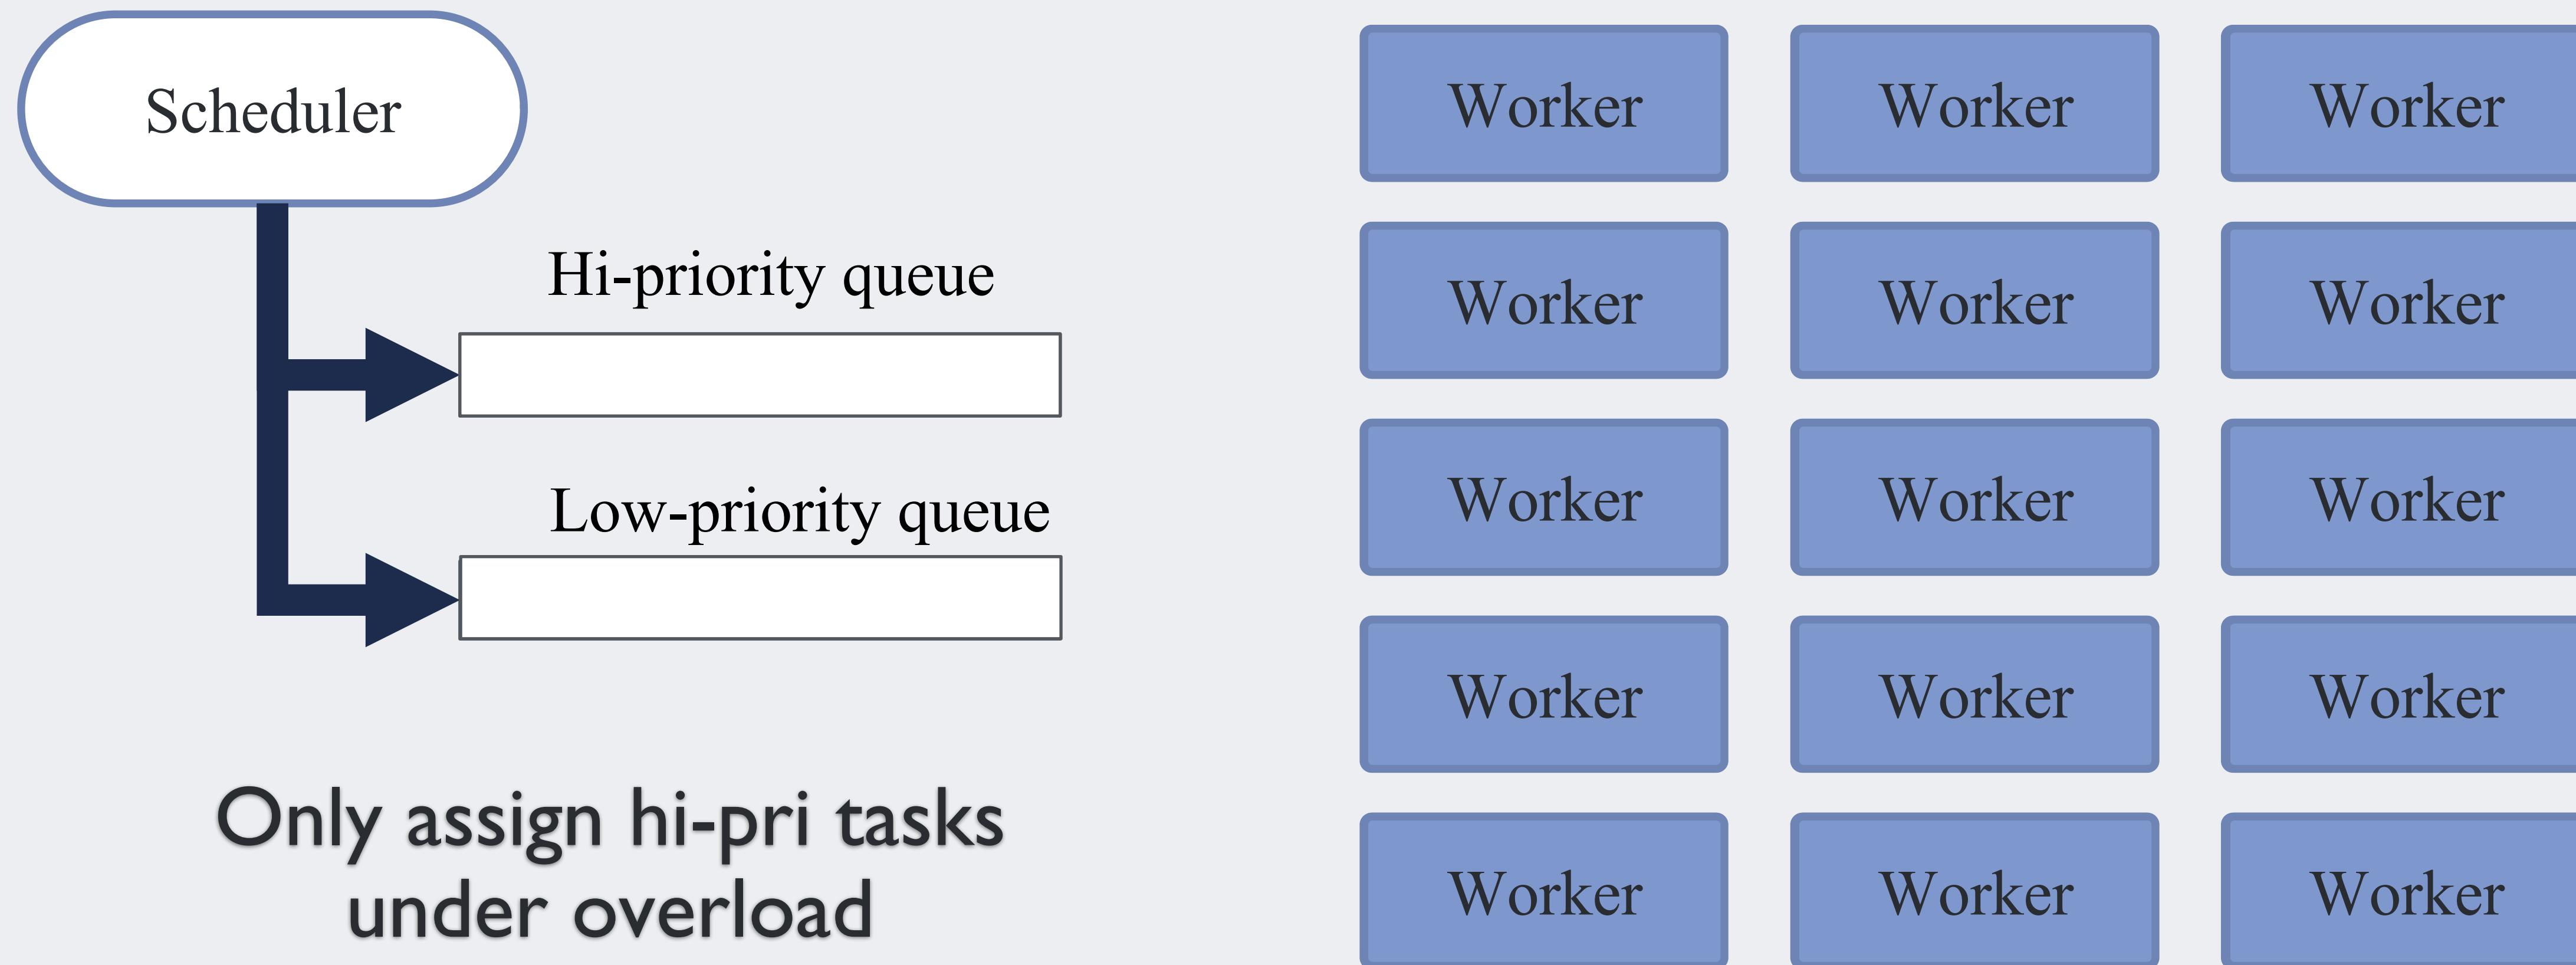
3X peak load during New Year Eve



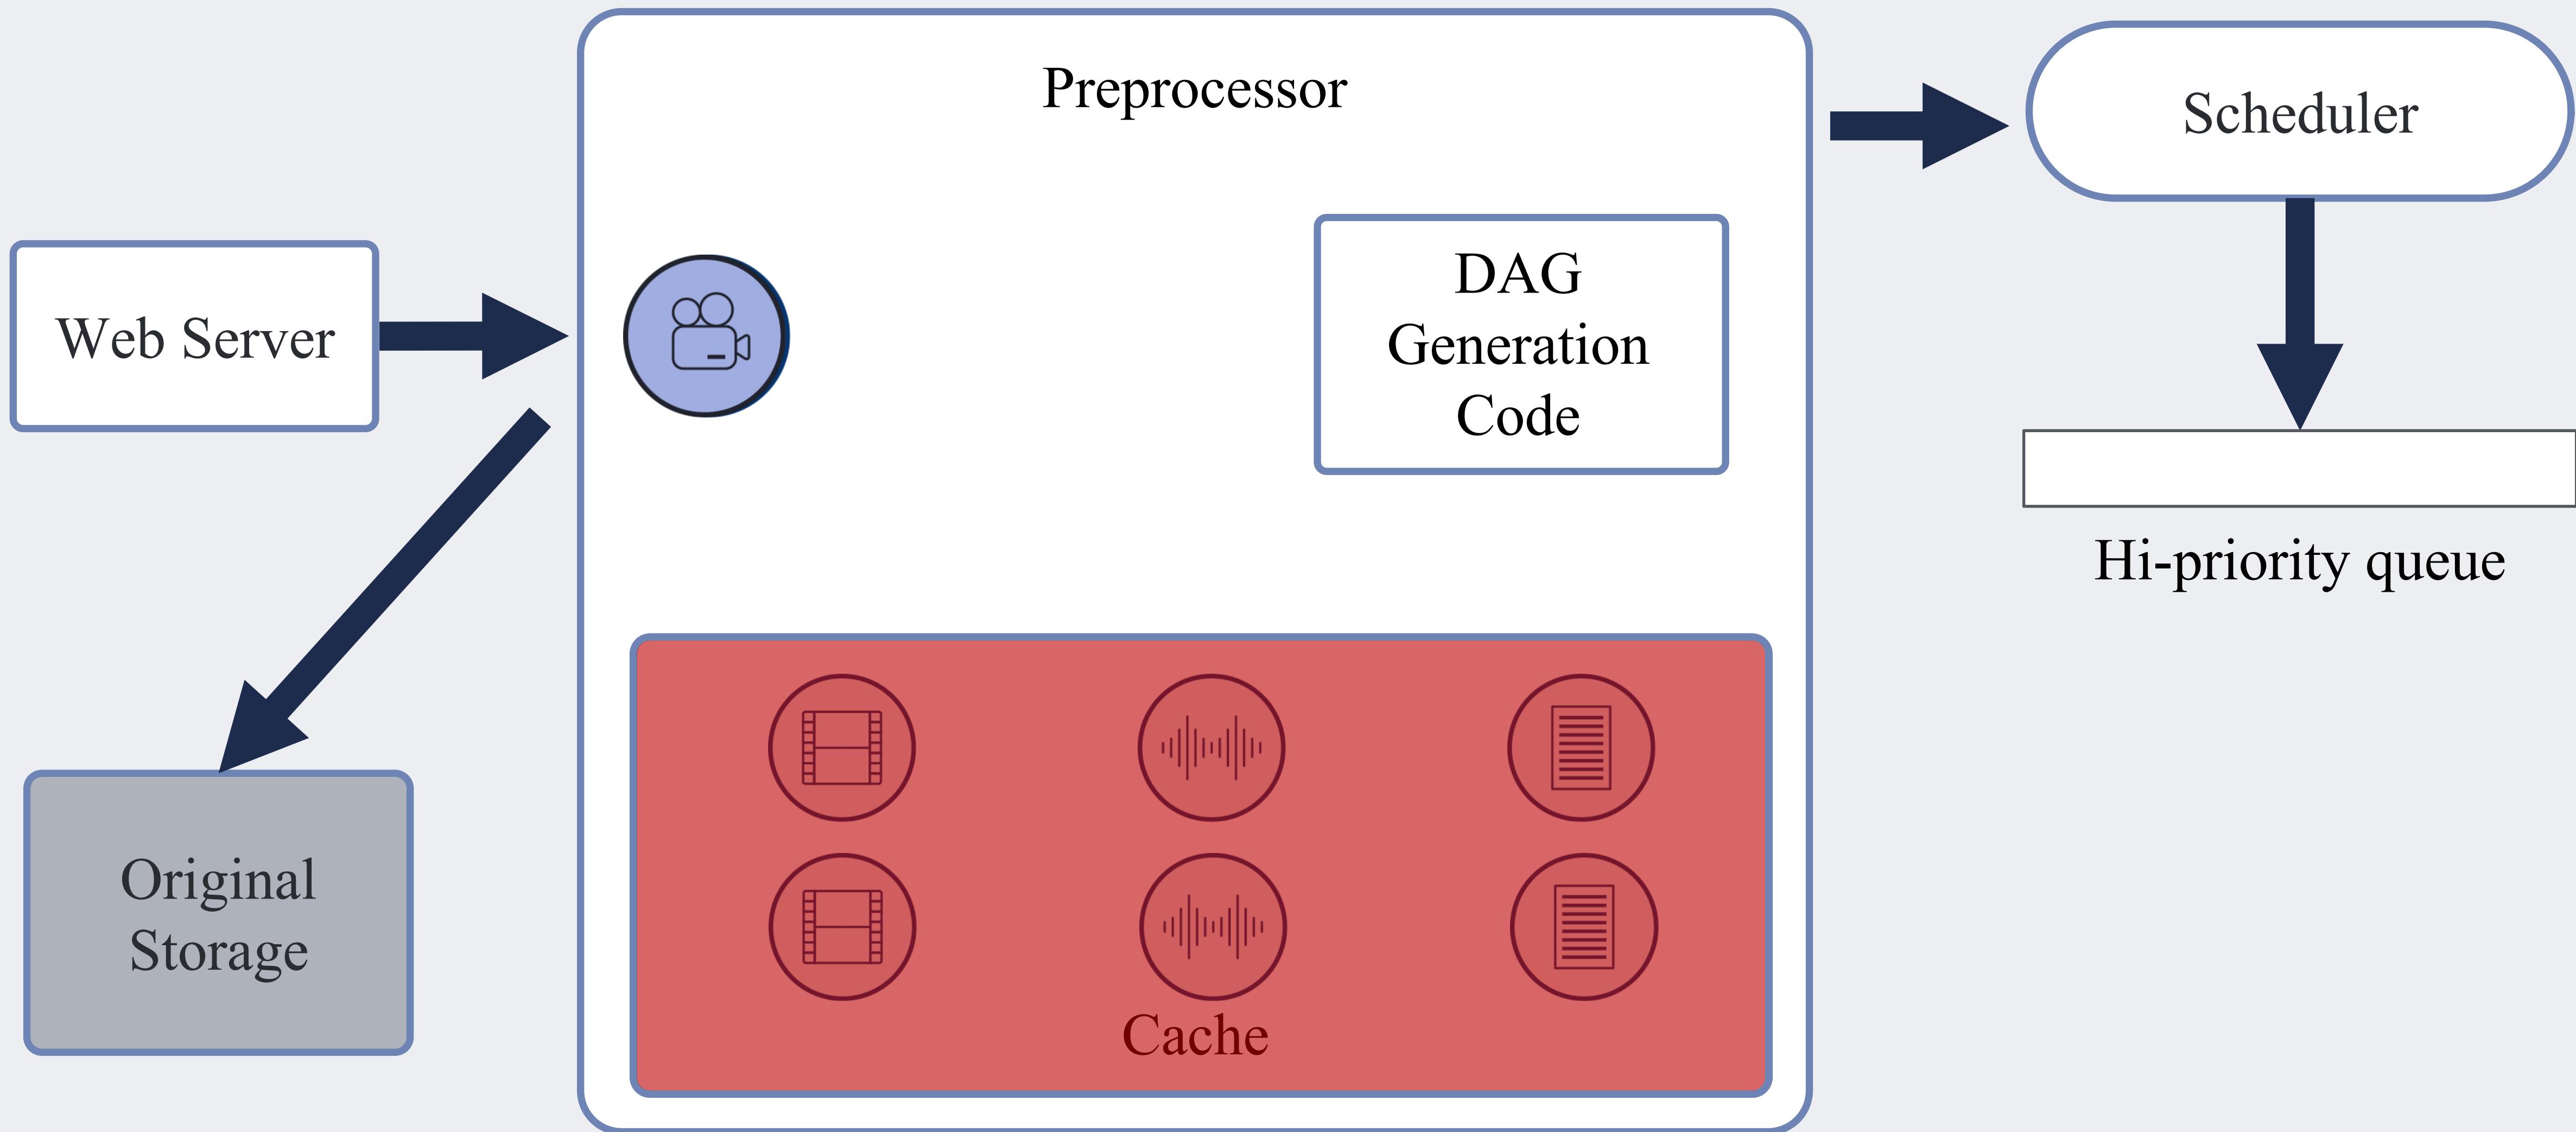
Prepare for overload



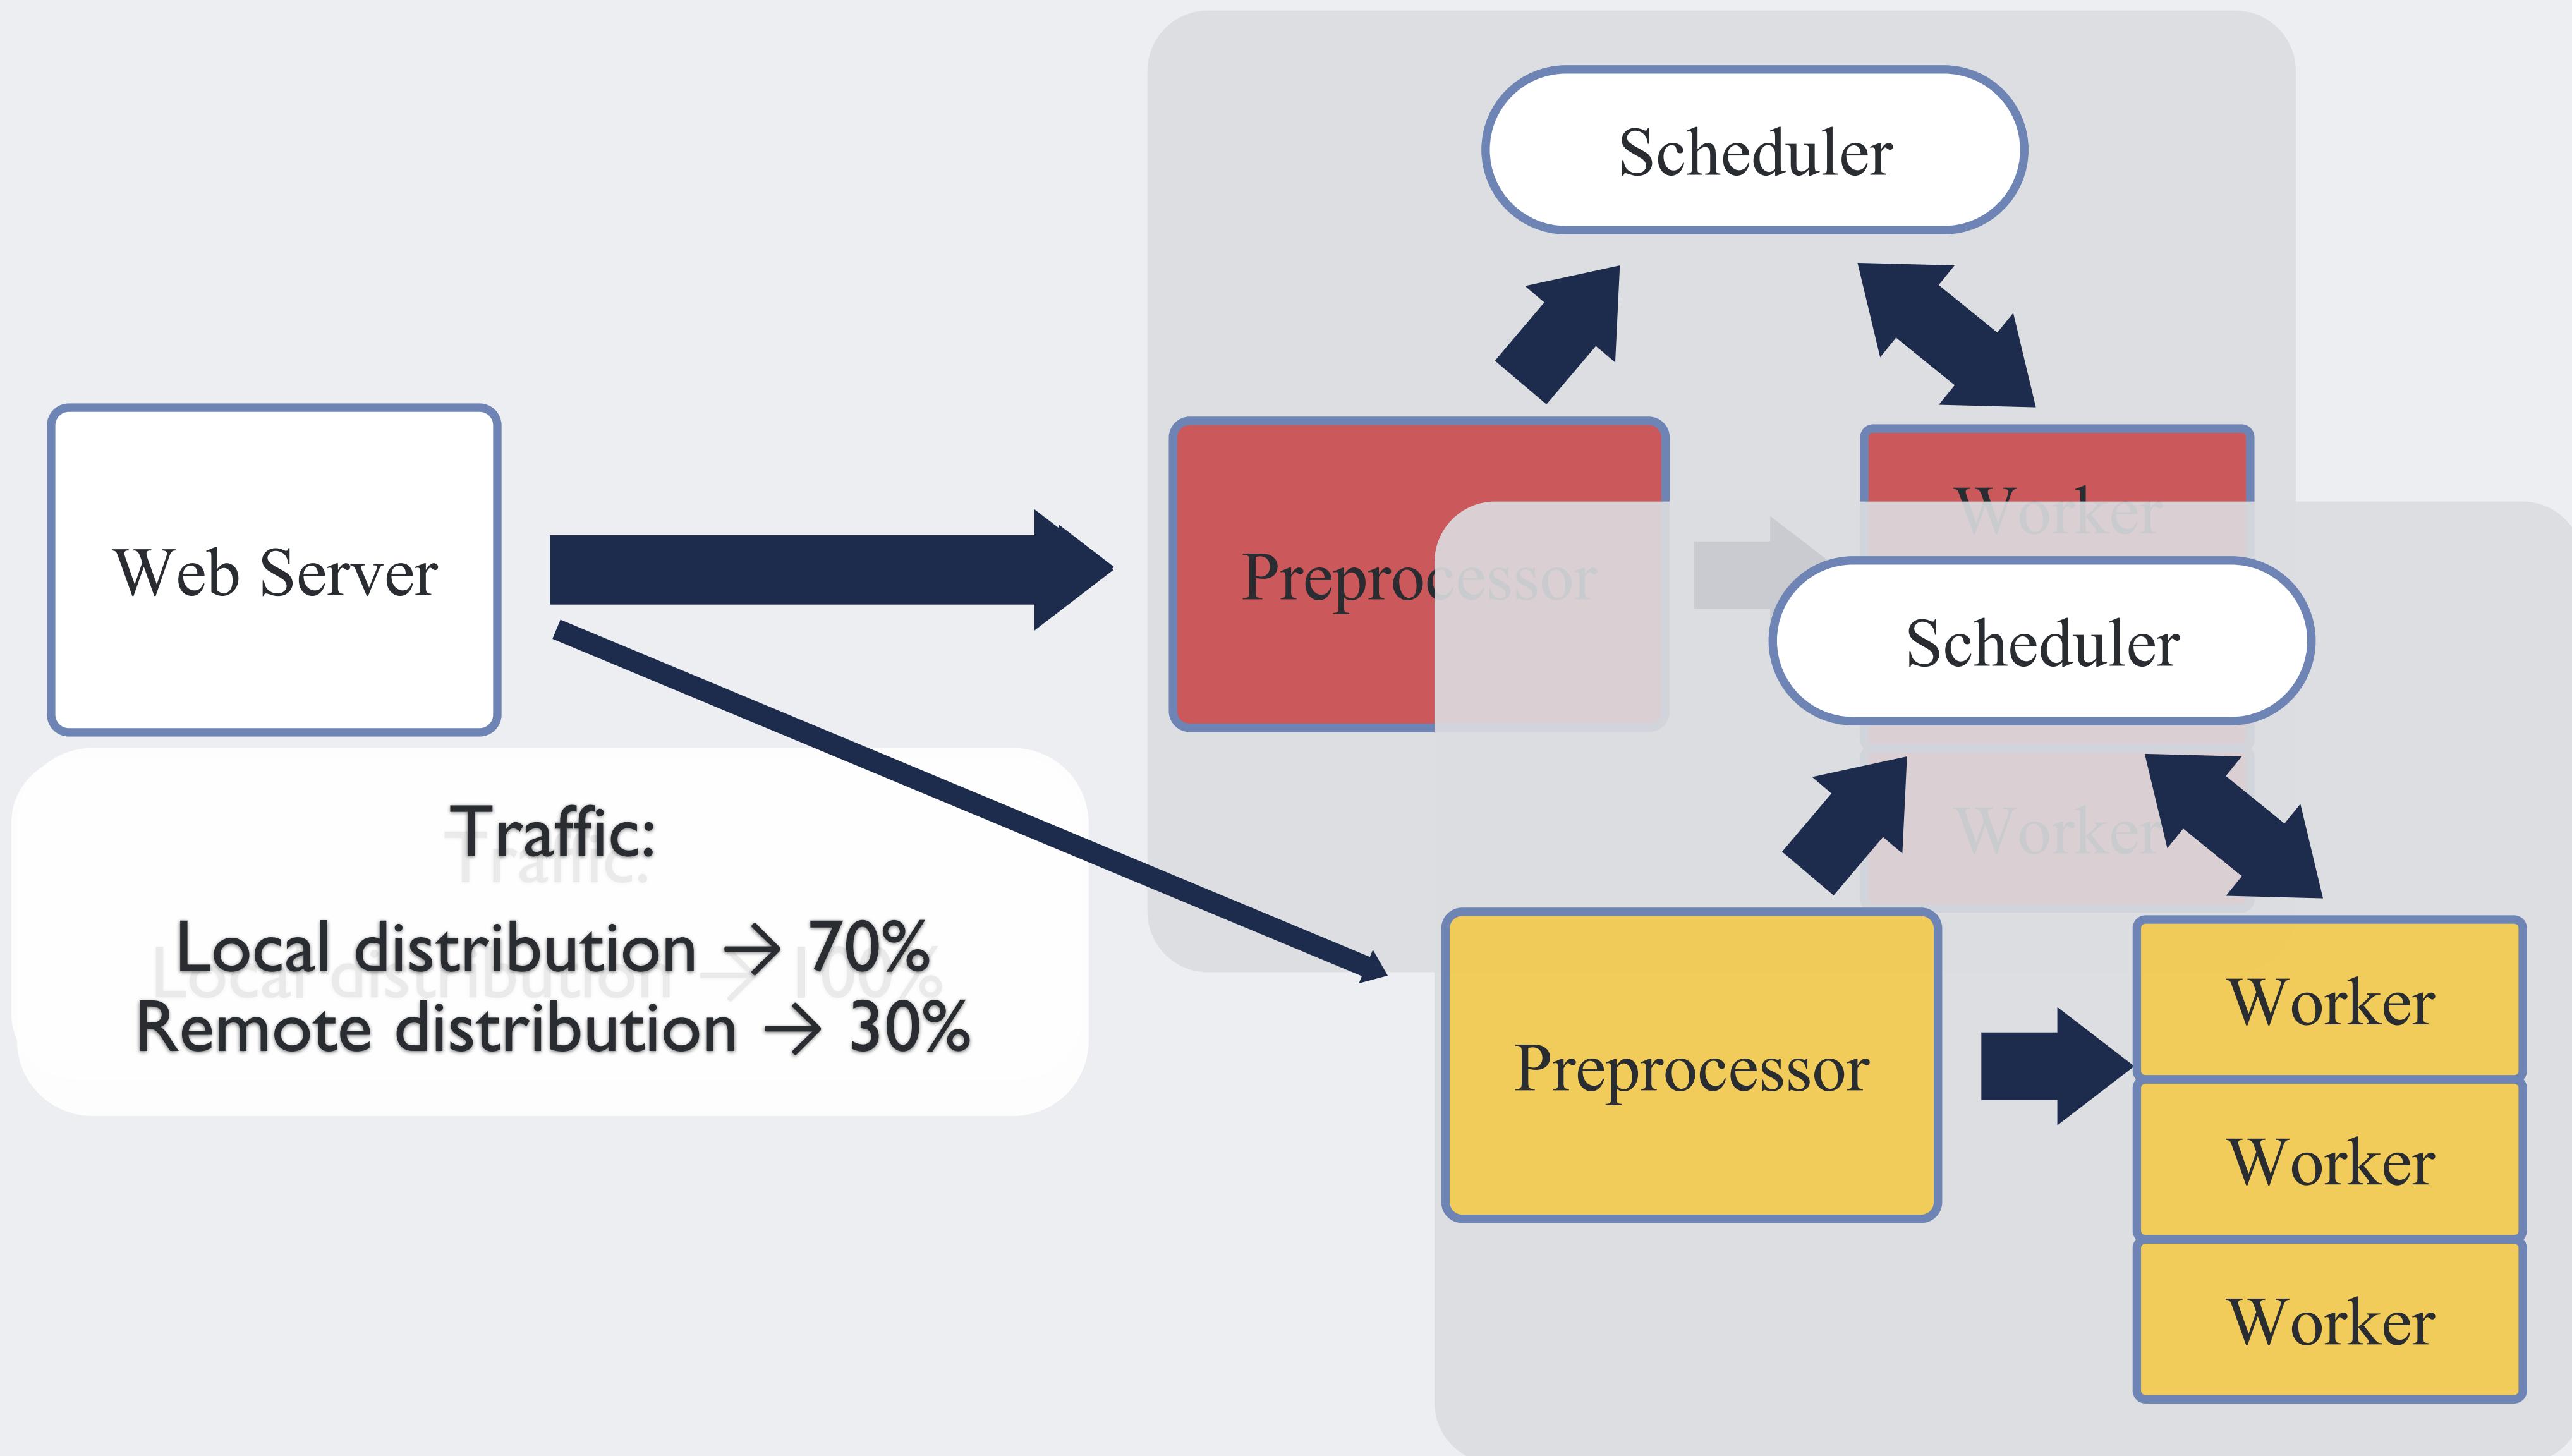
Use priority for worker overload



Defer full video processing



Regional redirection



Challenges for video processing @ FB

Speedy

2.3x ~ 9.3x speedup

Flexible

One system for 15+ applications

Robust

Handle faults Tolerate 3x traffic spikes inevitable at scale

More details in paper

- Advanced DAG control
 - Task group: batch multiple tasks for schedule
 - Priority control: annotate latency-sensitive task
 - Optional task: okay to fail or skip
 - Customizable error handling: early termination
- Failure monitoring and recovery
- Overload scenario caused by Kraken and system bugs
- Lessons learned

Failures from Global Inconsistencies

Lesson Learned

```
preprocessor = null

if (segment.is_first_segment()) {
    preprocessor = get_preprocessor()
    storage_write(video_id, "preprocessor", preprocessor)
} else {
    preprocessor = storage_read(video_id, "preprocessor")
}

forward_segment(preprocessor, segment)
```

One preprocessor handles all segments of one video

Mapping from video to preprocessor determined when upload starts

Storage system is eventually consistent, what could go wrong?

Related work

- Batch processing

SVE overlaps data ingestion and processing

- Stream processing

SVE offers dynamic DAG generation per input

- Video processing at scale

SVE support many production apps

Streaming Video Engine

- Deployed in production for 2 years
- Speedy to enable users to share videos quickly
 - Harness parallelism in upload, processing, and storage
- Flexible to support 15 app with tens of millions of uploads/day
 - Dynamic DAG generation on the stream-of-tracks abstraction
- Robust to tolerate faults and overload at scale
 - Prioritize processing and then shed load to other DCs or the future